



# Chapter 3:

# the Fetch and Decode

# Stages

# Background Required to Understand this Chapter




Basic Boolean Logic

Need for Branch Prediction

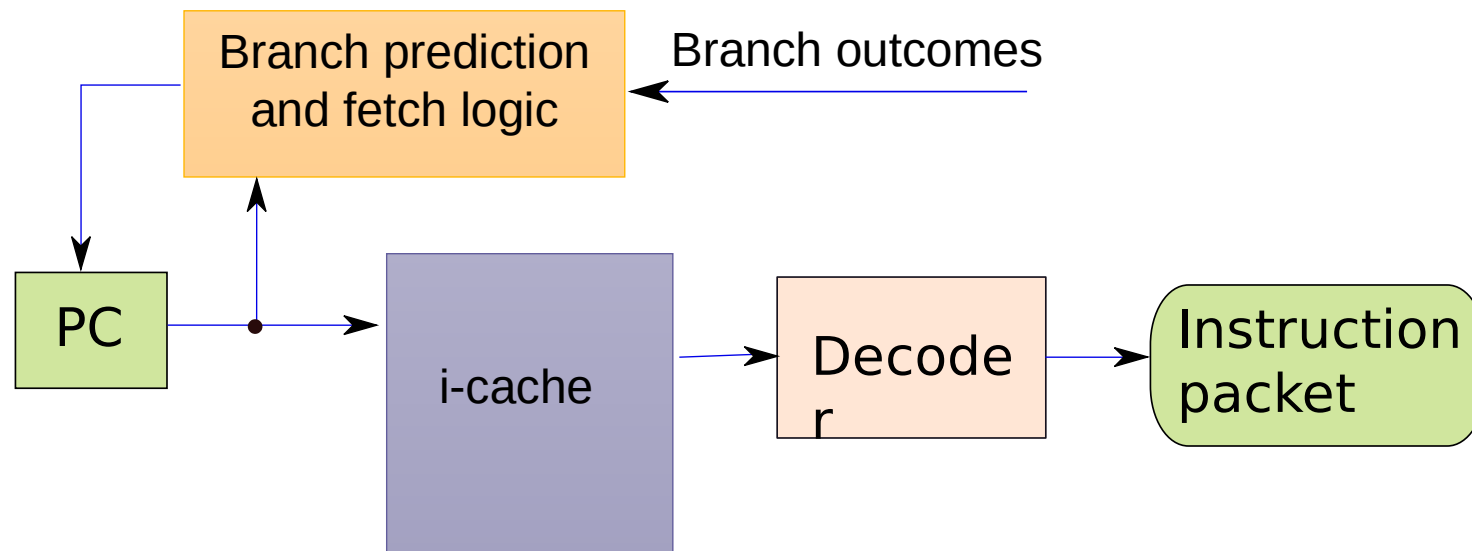
RISC and CISC ISAs

<http://www.cse.iitd.ac.in/~srsarangi/archbooksoft.html>




# Background

# The Fetch Stage




- The **branch prediction** and **fetch logic** computes the addresses of the instructions to be fetched in the next cycle
- The corresponding instructions are fetched from the i-cache
- They are sent down the pipeline

# Reading 4 instructions in one go from the i-cache

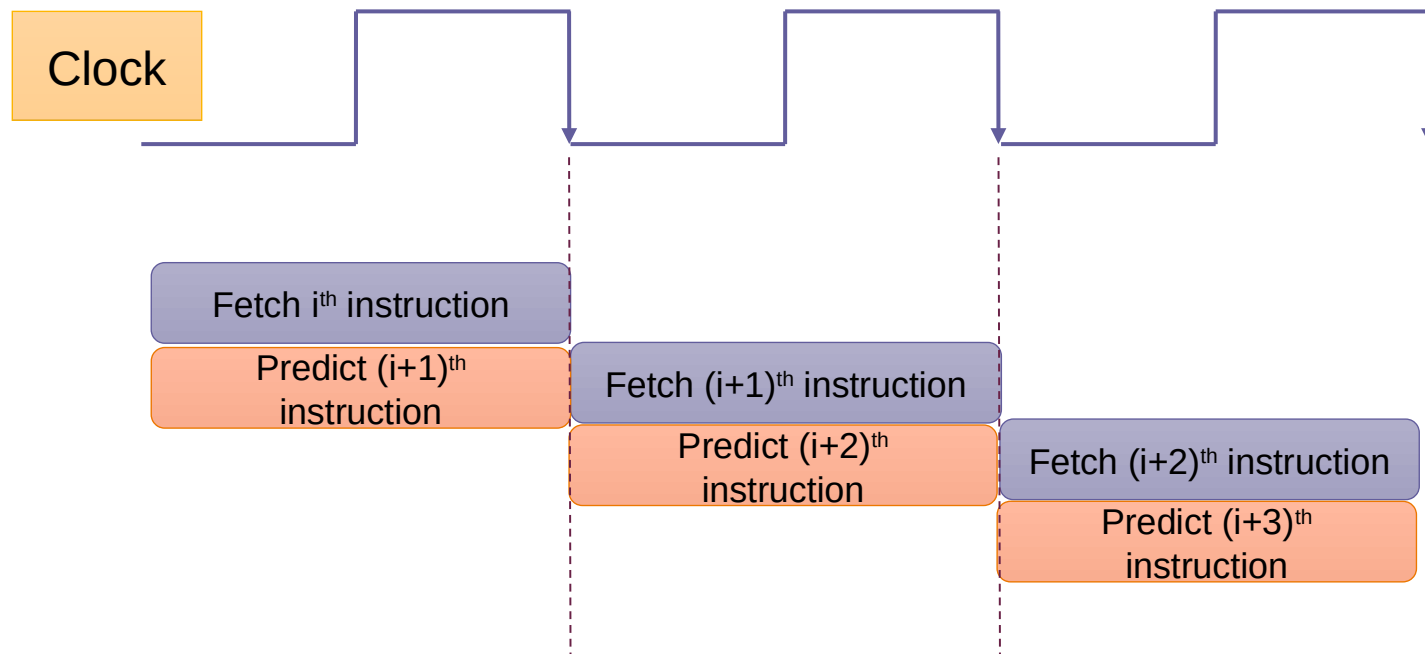


Assume no  
branches



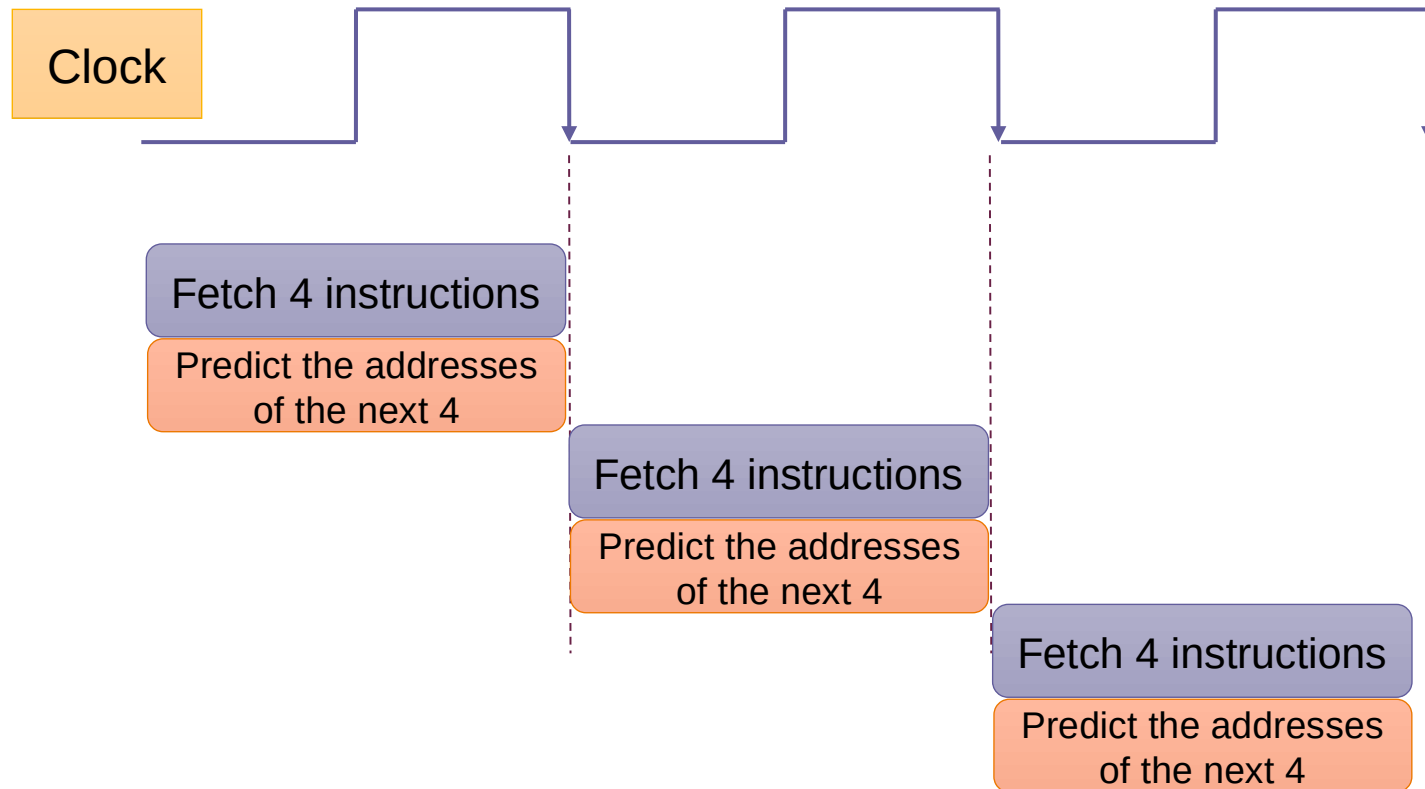
They might be in different cache lines

# Timing: Fetching a single instruction per cycle




There is no time to look at an instruction.  
We just know its PC.

# Timing: Fetching 4 instructions simultaneously



This is very hard to do. Why?

# Branches in a Bundle of 4 Instructions



- If we have branches in a bundle of 4 instructions, the instructions will not be contiguous
- To predict the second branch, we need to first predict the first branch




# Branch Prediction: Three Problems


**Predict if an instruction is a branch or not.**

**If it is a branch, predict its direction.**

**Predict its target.**



# Contents

- 
- 1. Prediction of the Instruction Type**
  - 2. Prediction of the Branch Outcome**
  - 3. Prediction of the Branch Target**
  - 4. Decode Stage**

# Is an instruction a branch or not?

## Insights

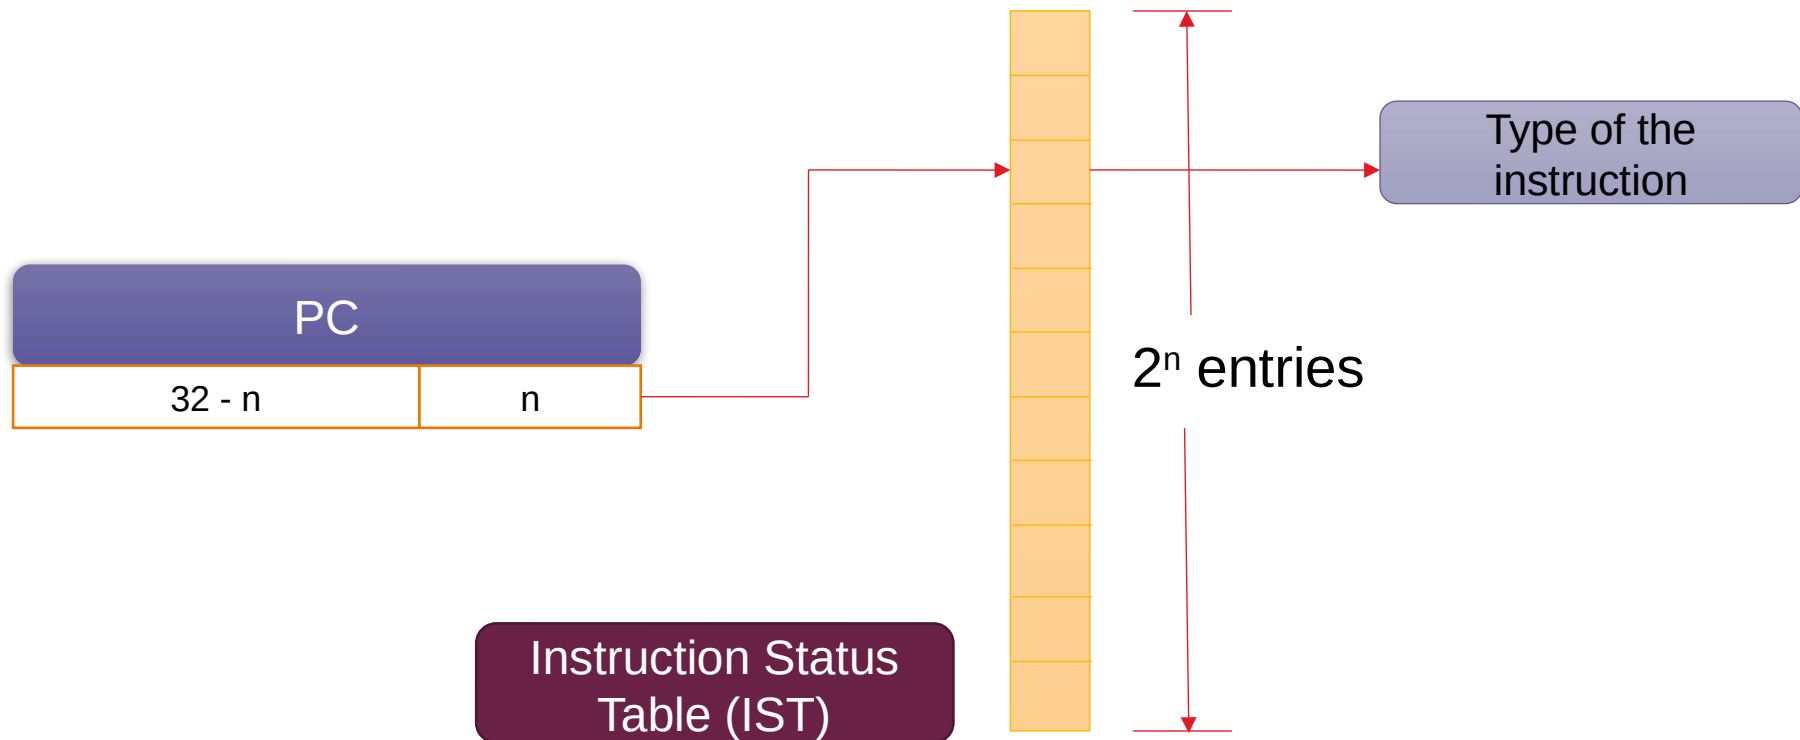
- Given a **PC**, the status of the instruction (branch or not) does not change.
- Can we use this information?

## Approach

- The last time that we had seen a branch **remember** its PC
- Next time we see a PC, **check** if we have seen it before
- Also **remember** the type of the branch:
  - Unconditional branch
  - Conditional branch  $\Rightarrow$   
depends on the result of a previous compare instruction
  - Function call
  - Return



# Make a structure in hardware to remember ...



**Assumption**

Assume a 32-bit PC address


# Basic Method of Operation

When we **see** a branch instruction

- **Access** its corresponding entry in the table
- **Record** the branch type

When we **see** an instruction:

- **Check** the corresponding entry of the table
- **Read** the type of the instruction



Why do we choose the  $n$  least significant bits?

# Instruction Status Table (IST)

If we choose the least significant 10 bits of the PC address, we will have 1024 entries in the IST

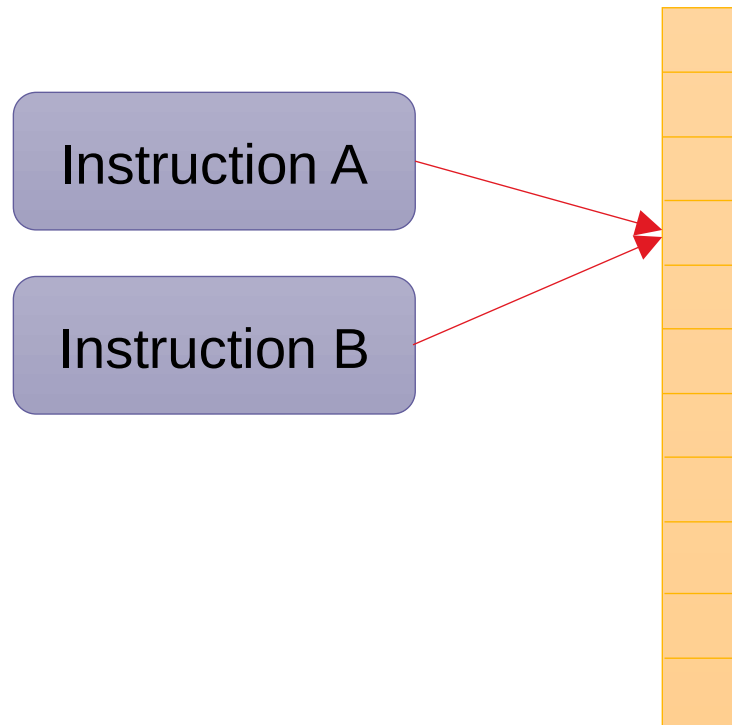
Why 10 bits?

For a 32-bit PC, we cannot have a  $2^{32}$  entry **table**.

- Too big
- Too slow
- Too much of power
- Too much of area

We need to manage with a **smaller** table

# Destructive Interference



Two instructions can **map** to the same entry; same **last**  $n$  bits.

Defined as **destructive interference**


# Branch Aliasing

- Possible for a **branch** and a **non-branch** instruction to map to the same entry
- Possible for two **branches** to map to the same entry
- Possible for two **non-branches** to map to the same entry
- Need for **disambiguation**
  - **Augment** each entry of the IST



# Disambiguation Between Addresses

Add a  $(32-n)$  bit tag




# Disambiguation

- We can keep the status of only **branches** in the IST
- However, for every **instruction**, we need to check the IST
- If there is no entry, then we need to predict that it is not a branch
- Can be **wrong**
- What to do  $\dashv$

Why does an IST work?

- Mainly because most pieces of code exhibit **temporal locality**
- In a given window of time **instructions** tend to repeat themselves

# Contents

- 
1. Prediction of the Instruction Type
  2. Prediction of the Branch Outcome
  3. Prediction of the Branch Target
  4. Decode Stage

# Predict the Direction of the Branch (Outcome)



- If a branch is **unconditional**  $\Rightarrow$  no need to predict (**taken**)
- If a branch is a **call/return**  $\Rightarrow$  no need to predict (**taken**)
- If a branch is **conditional**  $\Rightarrow$  need to **predict**

# Example Code

C

```
void foo() {  
    for (i=0; i < 5; i++) {  
        ...  
    }  
}
```

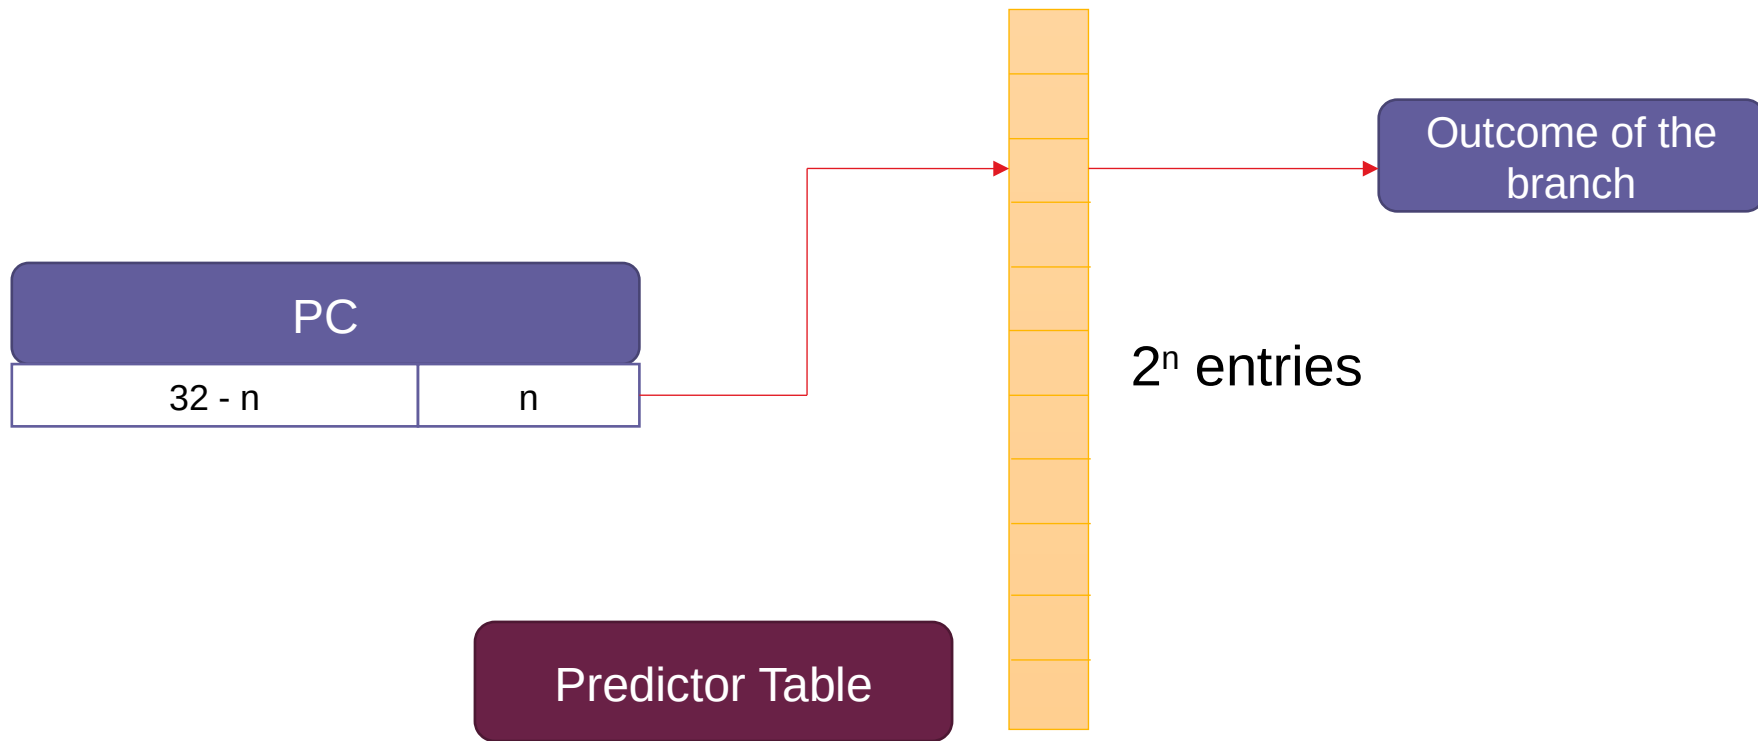
Assembly

```
.foo:  
    mov    r0, 0  
.loop:  cmp    r0, 5  
        beq    .exit  
        add    r0, r0, 1  
        b     .loop
```

Predict

This **branch** has roughly similar behavior most of the time. Predominantly **not taken**.

# Simple Bimodal Predictor



Remember the **last outcome**.  
**Current prediction** = **last outcome**

# Bimodal Predictor - II

Each entry saves the last recorded outcome of the branch

- taken or not-taken

What is the problem:

- The *beq* instruction is evaluated 6 times
- Assume the default is not-taken
- First 5 times  $\Rightarrow$  correctly predicted not-taken
- 6<sup>th</sup> and last time  $\Rightarrow$  incorrectly predicted to be taken

Now assume we call the function foo() once again

- The first time  $\Rightarrow$  we will predict taken. This is wrong
- Again the last time will be wrong
- Misprediction rate: 2/6

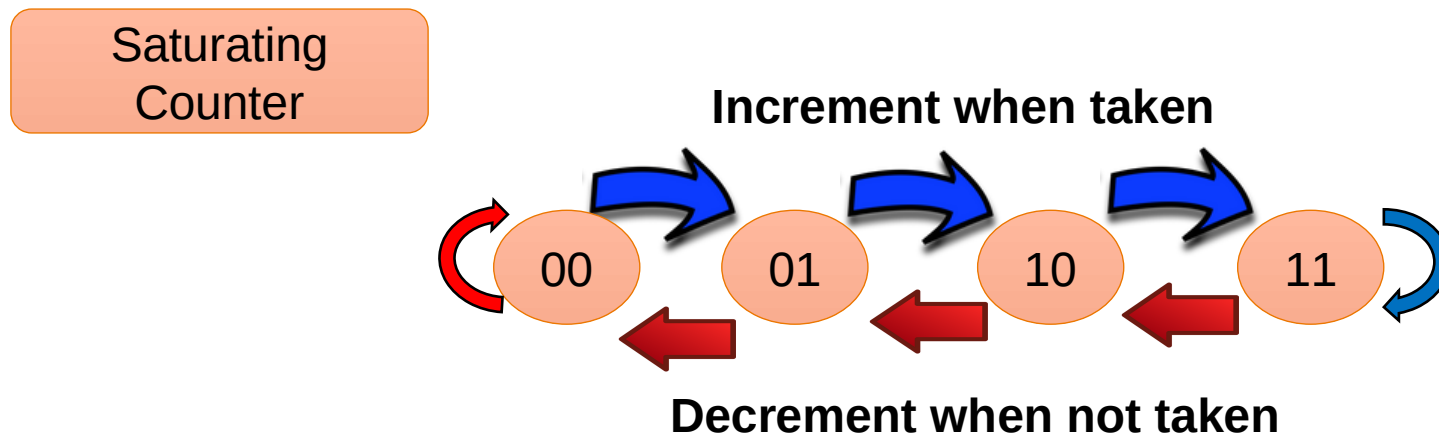
# Increasing Accuracy

What is the **problem**?

- The **beq** instruction is fundamentally biased towards **not-taken**
- One exception at the end of the **for** loop cannot change its inherent behavior


Let us add some hysteresis

- Instead of having a single bit in each entry, have a 2-bit counter





# Saturating Counters



**Algorithm** for updates:

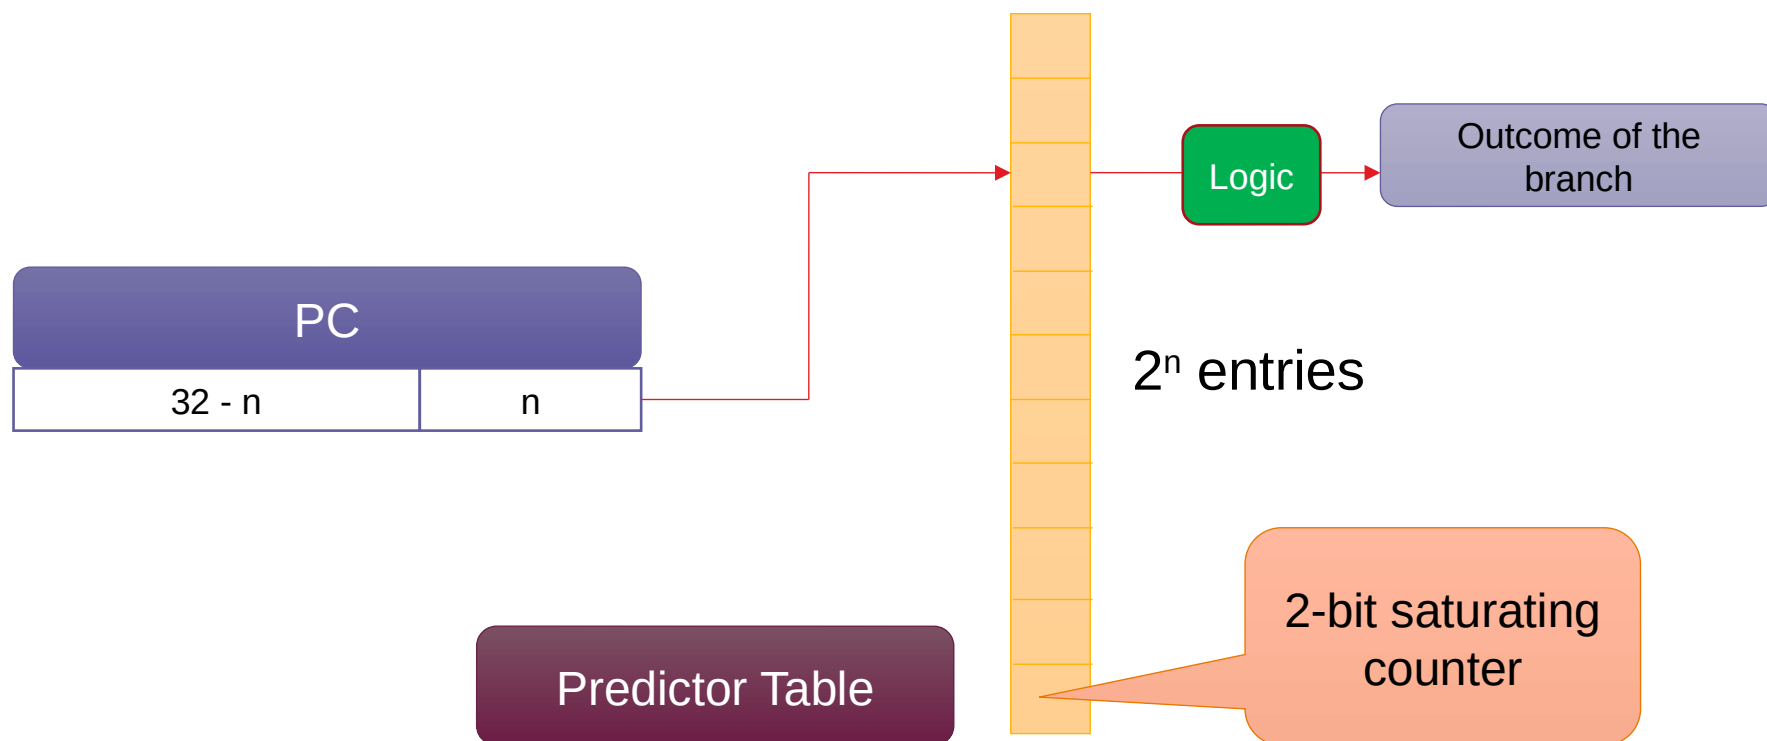
- If a branch is **taken**, increment the saturating counter
- If it is **not taken**, **decrement** the counter

For prediction:

- 00 and 01  $\Rightarrow$  **not taken**
- 10 and 11  $\Rightarrow$  **taken**

State	Name
00	Strongly not taken
01	Weakly not taken
10	Weakly taken
11	Strongly taken

# Bimodal Predictor with Saturating Counters



# Will it help?

```
.foo:
    mov    r0, 0
.loop: cmp    r0, 5
      beq    .exit
      add    r0, r0, 1
      b .loop
```

1<sup>st</sup> Time: Predict not-taken. **Correct**. Starts with 01 Moves to 00

2<sup>nd</sup> Time: Predict not-taken. **Correct**. Remains at 00

...

6<sup>th</sup> Time: Predict not-taken. **Wrong**. Starts with 00. Moves to 01



Misprediction rate: Down from 2/6 to 1/6

# Why do saturating counters work?

- In this case, we have a degree of **hysteresis**. The status of the branch (*beq .exit*) moves between the **strongly not taken** and **weakly not taken** states.
- If a branch is strongly **biased** towards one direction, but once in a while changes its direction, we should use **saturating counters**.
- They capture **stable** behavior with occasional anomalies.

# Can we do better?

What if we had?


C

```
void foo() {  
    for (i=0; i < 5; i++) {  
        ...  
        if (i == 4) {  
            foobar();  
        }  
    }  
}
```

Assembly

```
.foo:  
    mov    r0, 0  
.loop: cmp    r0, 5  
    beq    .exit  
    cmp    r0, 4  
    bne    .inc  
    call   .foobar  
  
.inc:  add    r0, r0, 1  
    b .loop
```

# Global History Register (GHR)



- Let us have a **shift register** that records the history of the last  $n$  **branches** encountered by the processor
  - We have one bit for each branch (regardless of the PC)
- We record: 1  $\Rightarrow$  taken branch, 0  $\Rightarrow$  not taken branch
- Let us consider a 2-bit shift register also known as the **GHR**
  - **GHR**  $\Rightarrow$  Global History Register
- We have two conditional branches in the running example
  - *beq .exit*
  - *bne .inc*

# Status of the GHR

Beginning of the 2<sup>nd</sup> Iteration

- 01

Beginning of the 4<sup>th</sup> Iteration

- 01



Beginning of the 5<sup>th</sup> iteration

- 01

Beginning of the 6<sup>th</sup> iteration

- 00



```
.foo:
    mov    r0, 0
.loop: cmp    r0, 5
     beq    .exit
    cmp    r0, 4
     bne    .inc
    call   .foobar

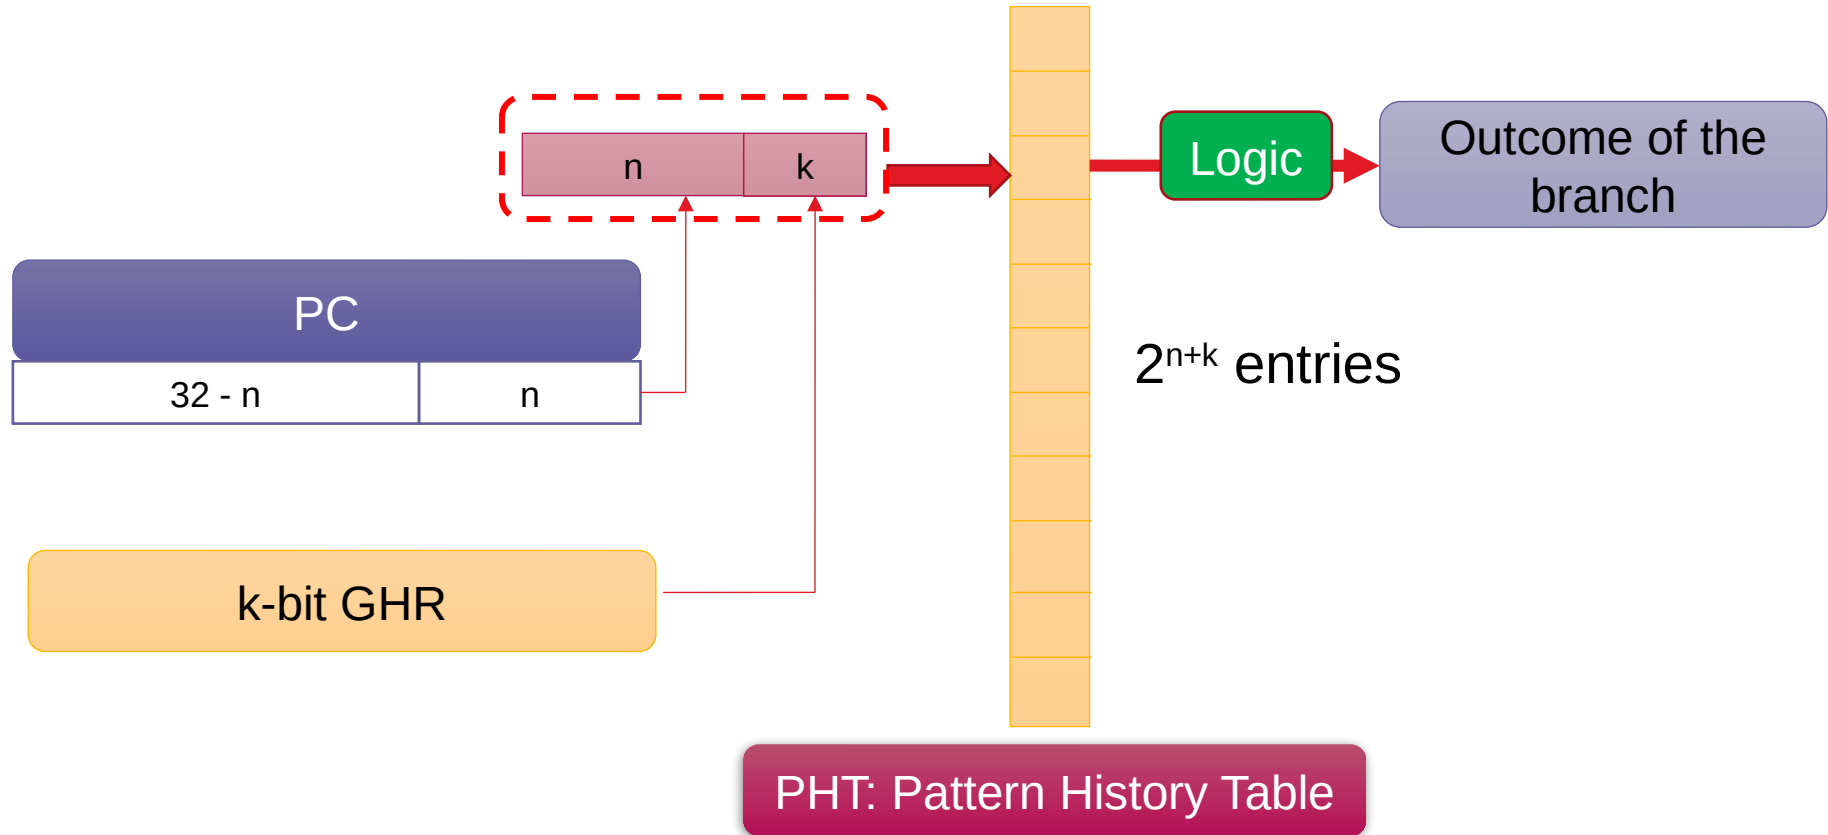
.inc:  add    r0, r0, 1

      b     .loop
```



Use the GHR information to also decide the direction of the branch.


# GAP Predictor





## GAp Predictor - II


- We create  $2^k$  times more **entries** in the predictor table
- **Capture** the global branch history
- Use both the global history as well as local (per PC) history to make the prediction.
- The accuracy is **expected** to **increase**.
  - The branch in the last iteration was **predicted correctly**
- In general we can create different **combinations** of:
  - The PC bits and the GHR's branch history



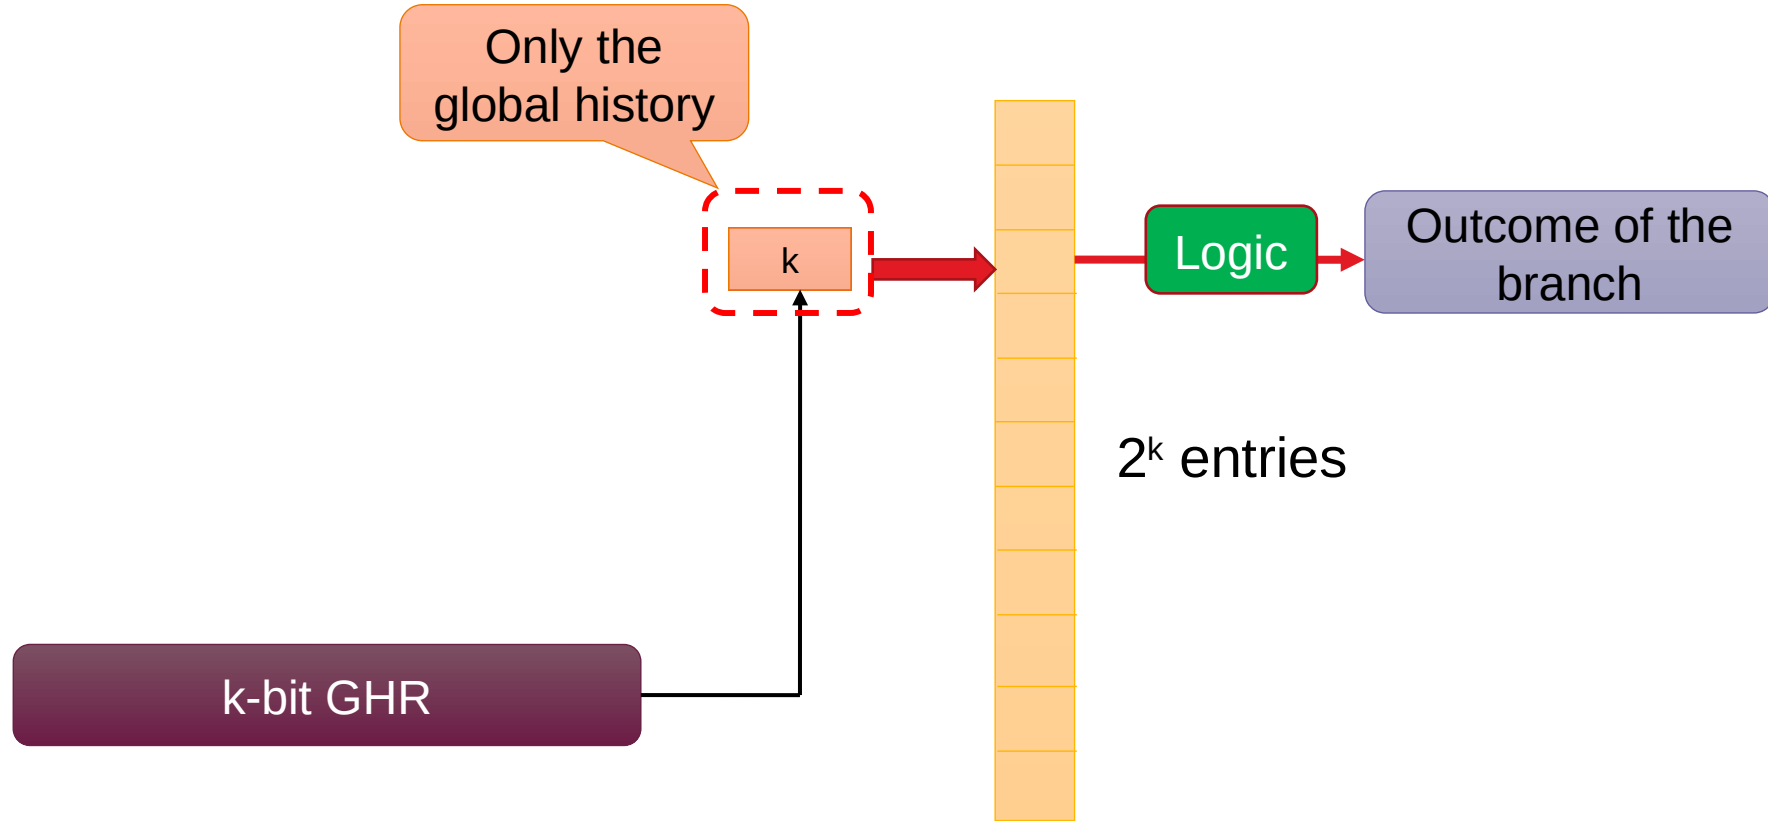
We have a branch that has an **alternating pattern**: **taken**, **not-taken**, **taken**, ... . Will a GAp predictor capture the pattern?

# Can we design a class of predictors?

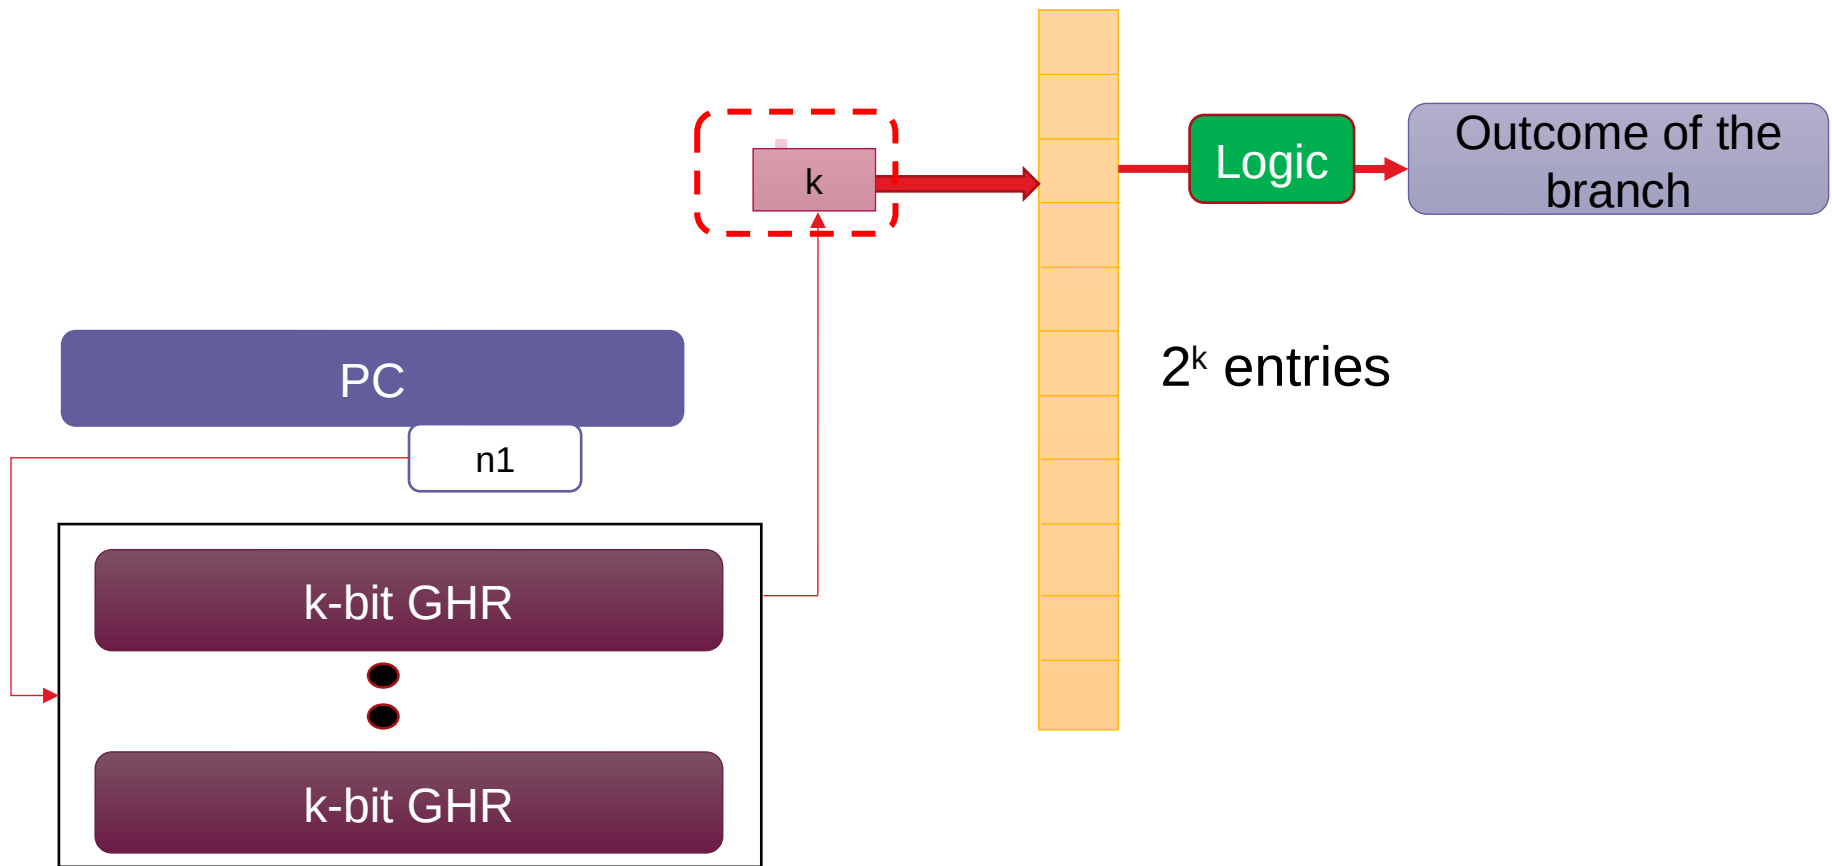
$G \sqsubseteq$  Global history,  $P \sqsubseteq$  Per PC pattern history



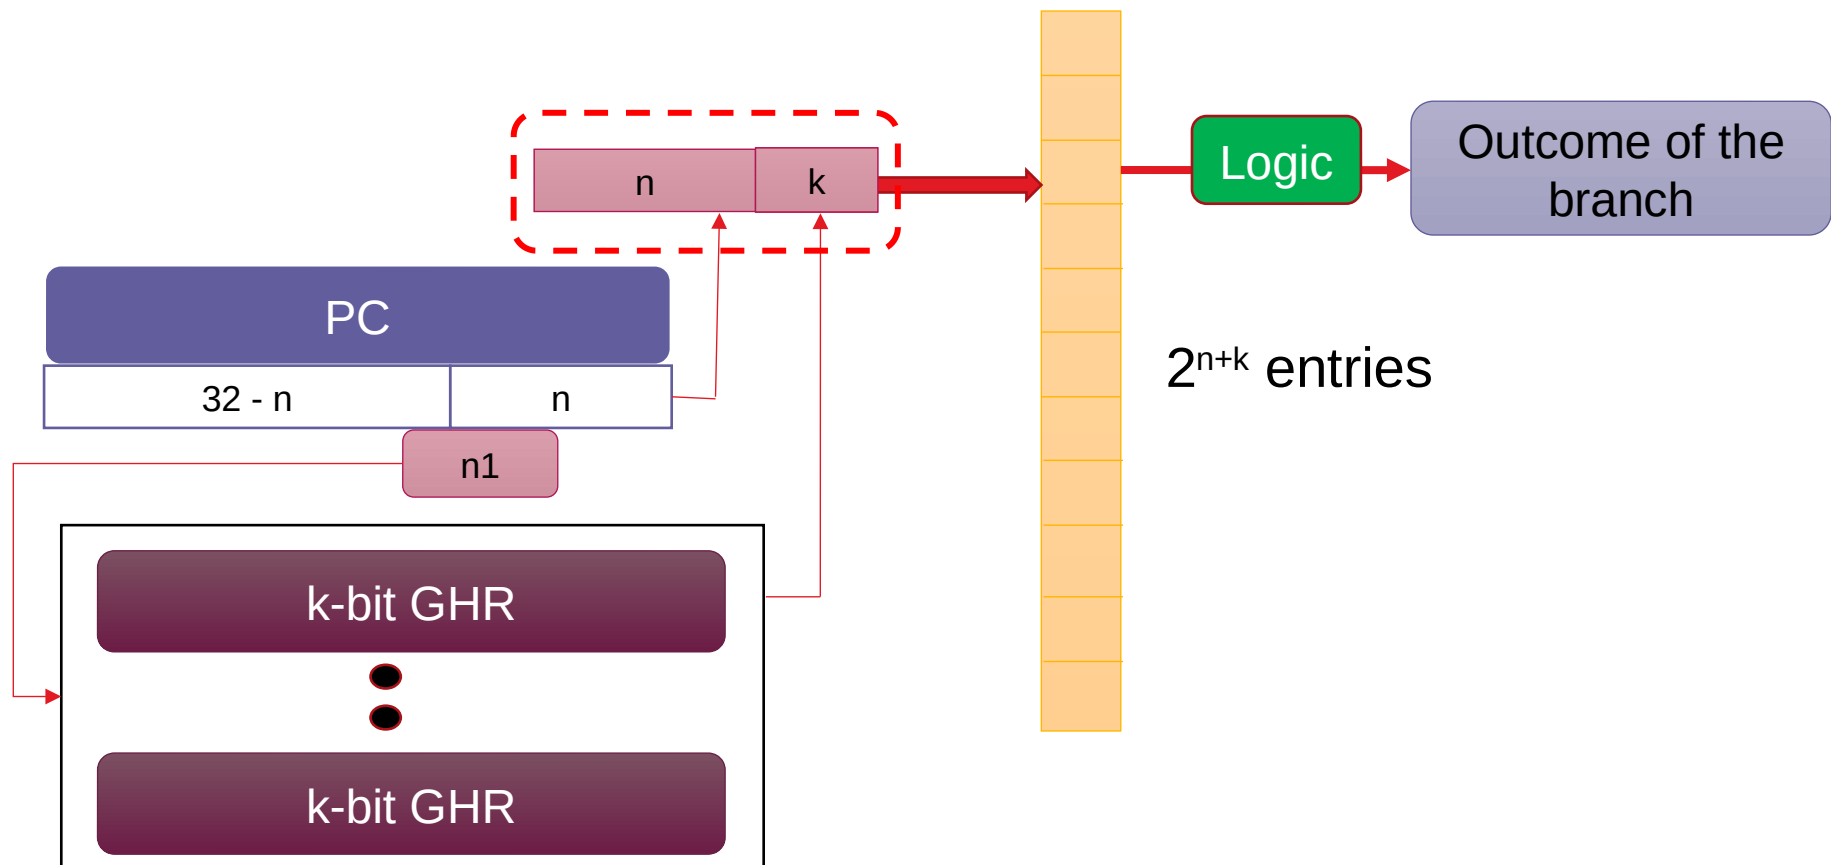
# GAg Predictor



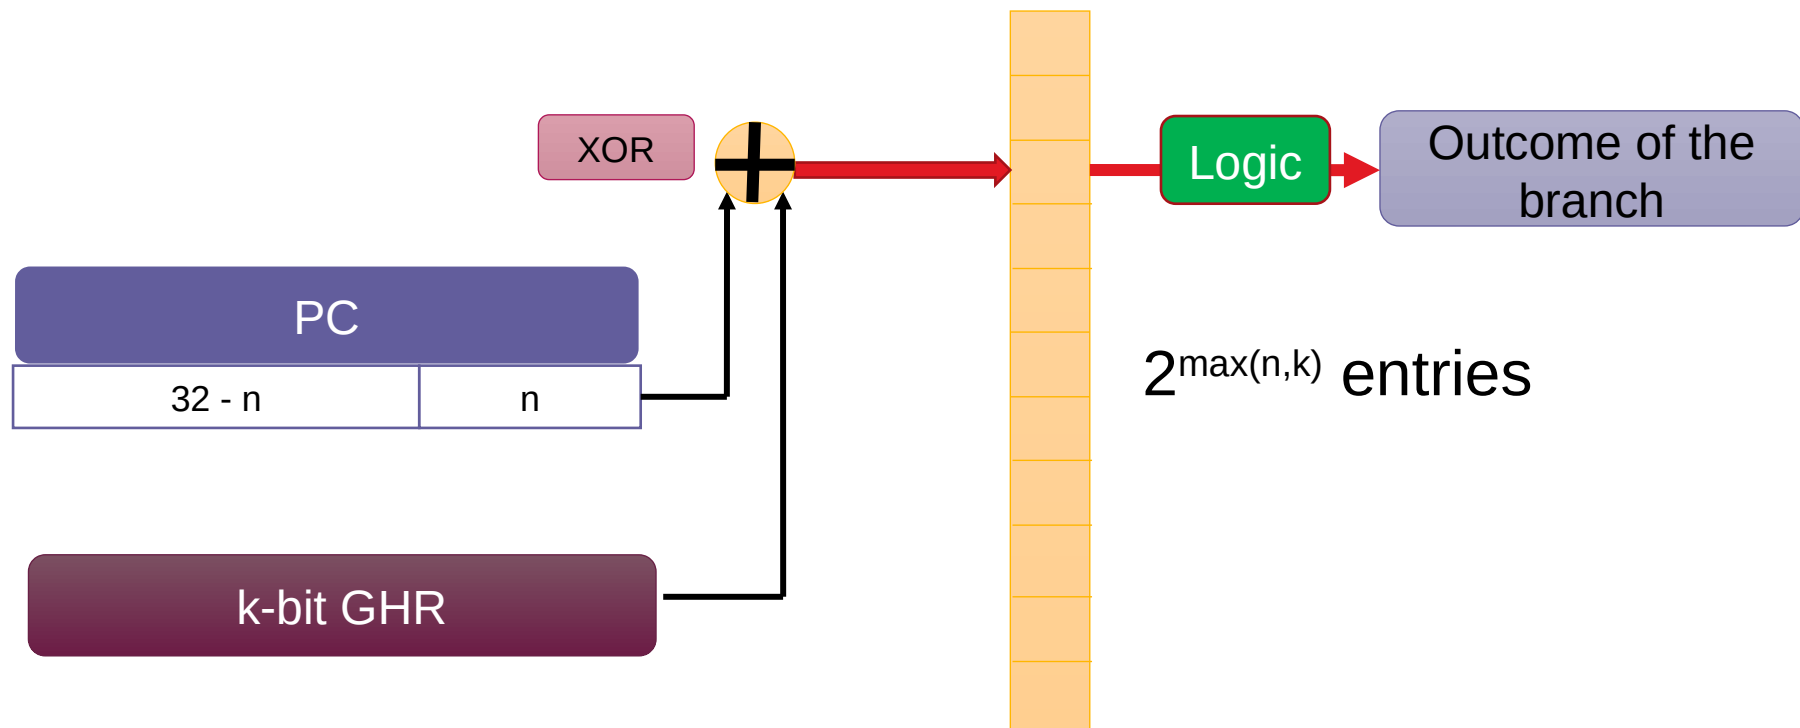
# PAG predictor



# PAP predictor



## Another way of combining information: GShare




# Tournament Predictor

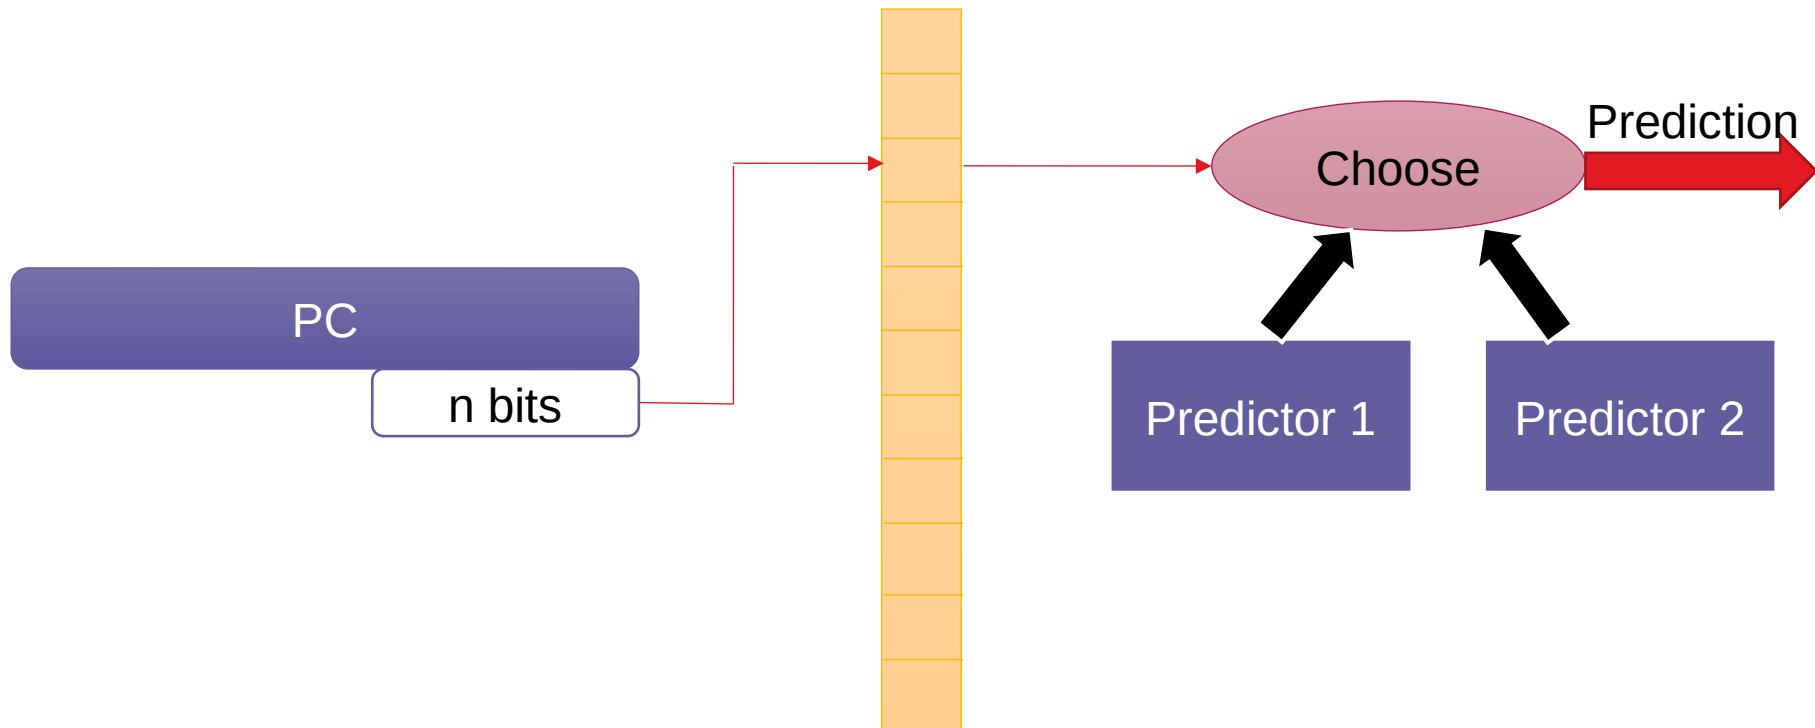
- Different predictors have different accuracies for different **program snippets**
- For a very **biased branch** a bimodal predictor with saturating counters is the best
- For an **alternating pattern**, a PAp predictor works very well
- How to know which predictor is the best for which piece of code?

What to do?

- Answer: Use two **predictors**, and **choose** between them



# Tournament Predictor





# Operation of a Tournament Predictor

## Prediction

- Find the entry in the **choice array**
- Choose the predictor based on the value of a saturating counter
- Use its prediction  $\Rightarrow$  because of the saturating counter, we automatically choose the predictor that performs the best for a given PC.

## Training

- Train both the **predictors**
- Train the entry in the **choice array** if we chose the wrong predictor
- If both the predictions are the **same**, we don't modify the choice array
- If they **differ**, we increment the counter if Predictor 1 was **correct** and decrement it if it was **incorrect**.

# Other methods of prediction

## Reduce **aliasing**

- Incorporate a few tag bits in each **entry** of the predictor
- Have multiple predictors for different subsets of branches
- Include a **bias bit** with every branch (its most likely direction), and just predict if we need to agree with the bias bit or not (agree predictor)

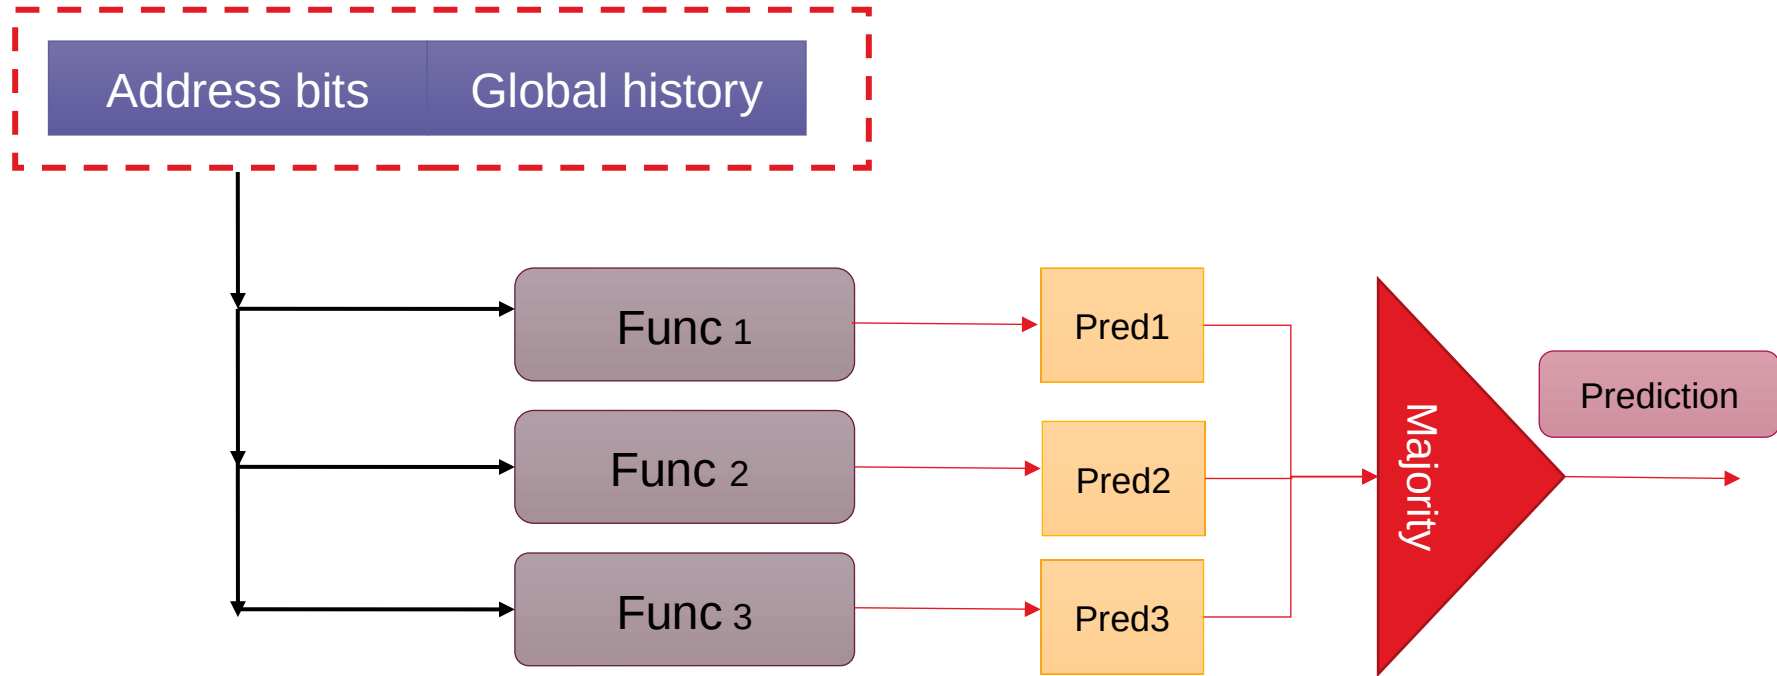
## Better **use** of bits

- Separate **high confidence** and **low confidence** branches. Dedicate more bits to low confidence branches

## Examples of other predictors:

- Bi-mode, Agree, Skew, YAGS, TAGE

# Skewed Branch Predictor




# Prediction and Compression

- Consider a sequence of  $\langle \text{PC}, \text{outcome} \rangle$  pairs
- Let's compress this sequence using a standard program: zip, rar
- Is the prediction accuracy related to the compression ratio (size of compressed file/ size of original file)?


- YES

Take a course on information theory


- Read about the Fano's inequality




# Contents

- 
1. Prediction of the Instruction Type
  2. Prediction of the Branch Outcome
  3. Prediction of the Branch Target
  4. Decode Stage

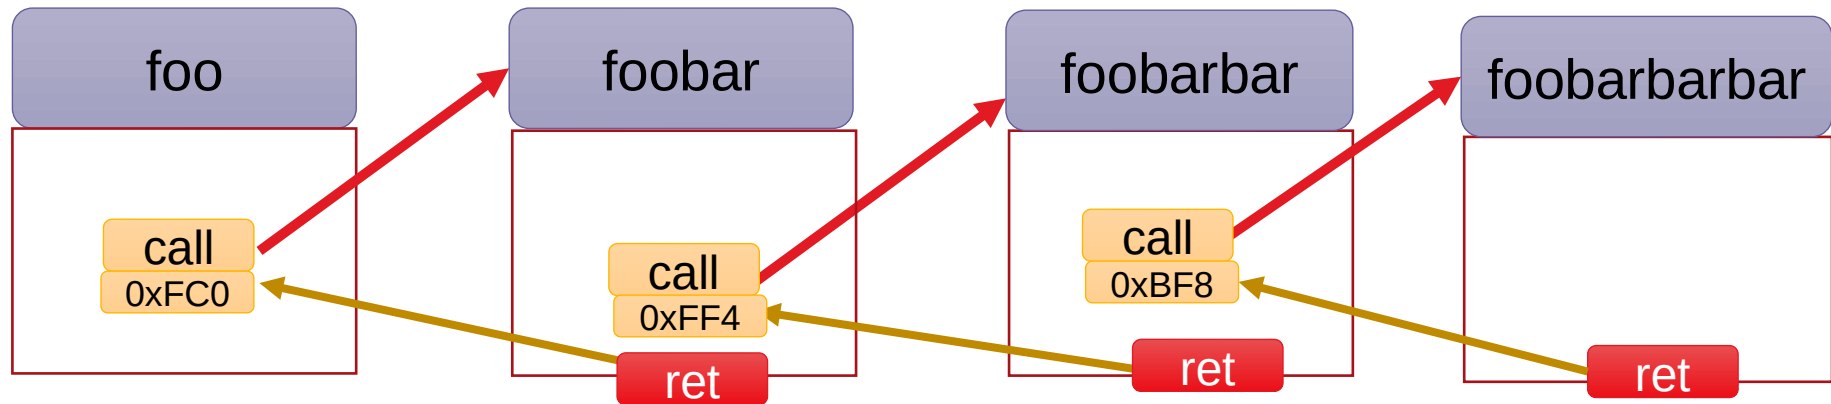
# Predict the Branch Target



# Use the IST. Let us call it the Branch Target Buffer (BTB)



# Calls and Returns

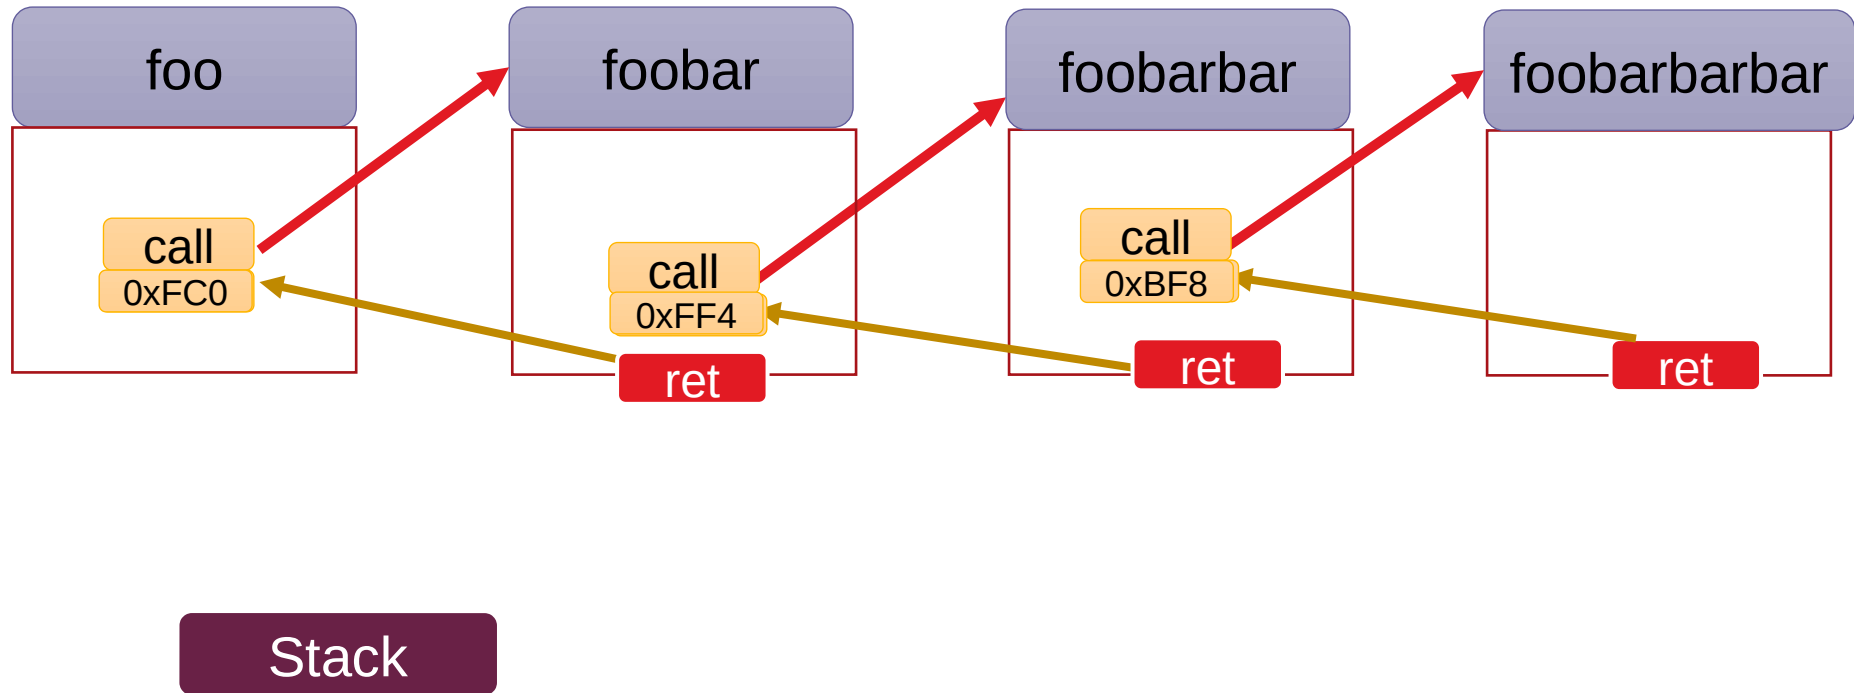


How to predict return addresses?

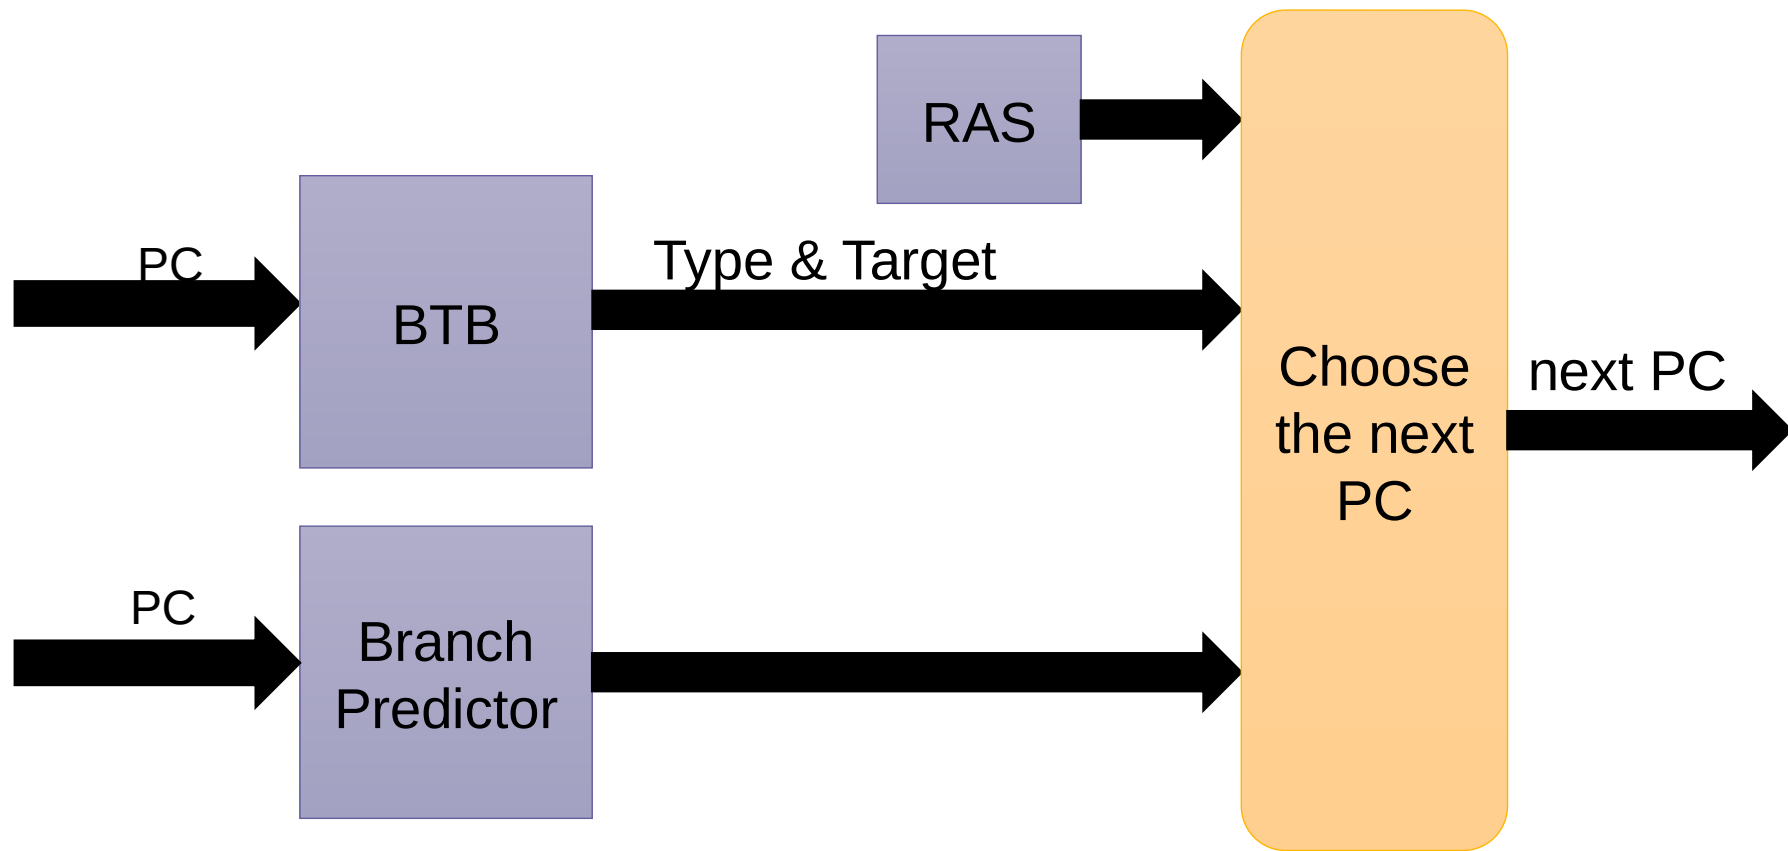
Solution: Use a stack




# Return address stack (RAS)



## Summary: The Branch Prediction System



# Contents


1. Prediction of the Instruction Type
2. Prediction of the Branch Outcome
3. Prediction of the Branch Target
-  4. Decode Stage

# The Process of Decoding


- **Expand** all the immediate values to 32 or 64-bit values
- **Extract** all the fields
- **Compute** the branch target (if branch is taken)
- **Add** all implicit sources (such as for the return instruction)
- **Create** the instruction packet

# Issues with CISC Instructions

## CISC Instructions

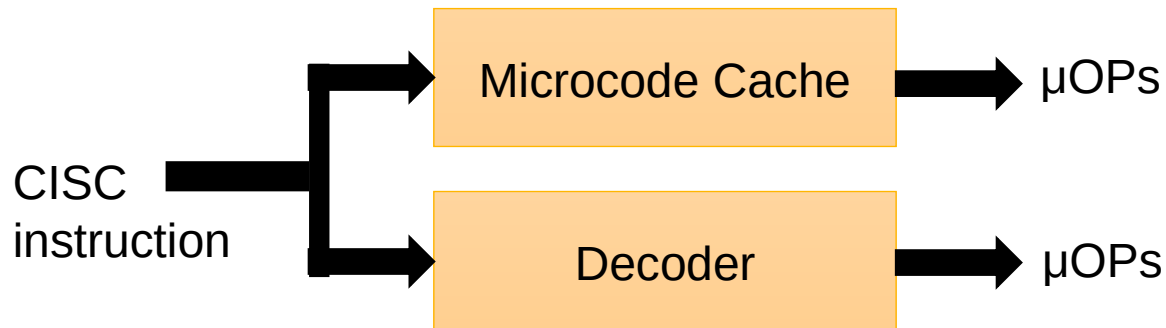


- Do not have a **fixed** length
- In **x86** the length varies from 1 to 15 bytes
- It is **hard** to fetch multiple instructions at once
  - Need to know the **boundaries of instructions**
- It is hard for OOO pipelines to **process** them
  - Most CISC processors internally convert from CISC to RISC



# Regular Decoder vs Microcode Cache

- x86 has some very **complex** instructions
- Consider the *rep movsd* instruction
  - It can be used to **copy**  $n$  elements from one location to the other
    - all in one single instruction
  - Such instructions **map** to a large number of  $\mu$ OPs
- In comparison, simple **add** and **sub** instructions with register operands map to a single  $\mu$ OP.

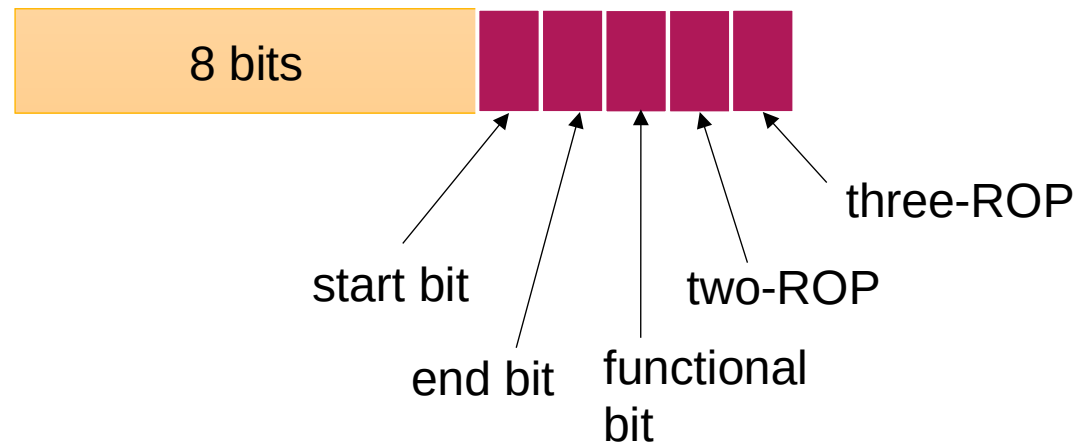


# Predecoding CISC Instructions



- When we **read** in an instruction into the i-cache
  - We **predecode** it
  - ☐ We **annotate** each byte with additional information
  - When we read the bytes from the i-cache, we can **use** this additional information to decide the instruction boundaries.  
[Narayan and Tran, 1999]

# Predecoding - II



- **Start bit**  $\Rightarrow$  Starting byte of an instruction
- **End bit**  $\Rightarrow$  Last byte of an instruction
- **Functional bit**  $\Rightarrow$  Interpretation depends on the implementation.
- **two-ROP and three-ROP bits**  $\Rightarrow$  Number of  $\mu$ OPs in an instruction



# Optimizing Operations on the Stack Pointer

## Example code

```
sub sp, sp 8
st r1, 0[sp]
st r2, 4[sp]

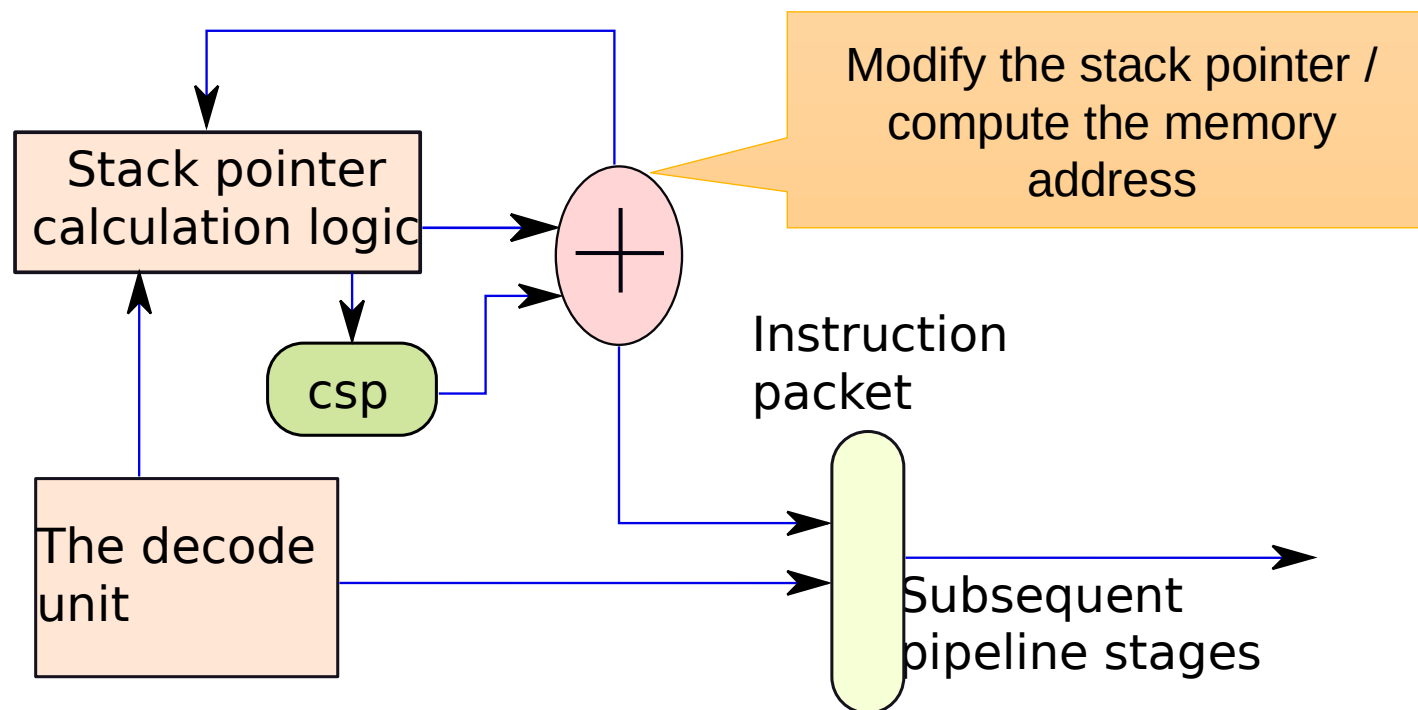
....
ld r2, 4[sp]
ld r1, 0[sp]
add sp, sp, 8
```

- There is a **pattern** here.
  - We either **modify** the stack pointer by adding or subtracting a constant
  - The **load** and **store** instructions access the stack pointer.
  - Update the stack pointer locally or get the memory address directly



Have a copy of the *sp* register in the decode stage itself.

## Add a *csp* register in the decode stage



# Corner Cases

*ld sp, 12[r1]*

- When we **encounter** such instructions, we set the value of **csp** to null.
- Treat *sp* as a regular register in the OOO pipeline.
- Accumulate the difference,  $\Delta$ , between the value returned by the load **addr** and the current value of the stack pointer.
- When the load instruction returns with its value:  
set

$$\text{csp} \equiv \text{addr} + \Delta$$


# Advantages

- **Load** instructions to the stack can be issued **early** because we compute the address at decode time.
- For **load** and **store** instructions to the stack, the address need not be computed in the pipeline.
- We can get **rid** of many instructions that update the stack pointer. This will effectively reduce the dynamic instruction count in the rest of the pipeline.

# Instruction Compression

- Instruction caches have a **limited size**: typically 32 to 64 KB
- The performance is extremely sensitive to their size
- Hence, we wish to **pack** as many instructions as possible
- **Approach 1**: Reduced-width instructions
  - **Support** a limited number of opcodes
  - **Avoid** encoding complicated flags and options
  - Reduced view of architectural registers
  - Reduce the size of the immediate fields
  - Implicit operand (accumulators):  
$$\text{add } r1, r2 \qquad (r1 \equiv r1 + r2)$$

# Instruction Compression



- The compiler or profiler **identifies** sequences of frequently executed instructions
- Each sequence is **replaced** with a code word
- The code word **uses** fewer bits
- The code word  $\Rightarrow$  sequence mapping is **stored** in a separate hardware structure: **dictionary**
- **Reduces** code size. We can even **store** decoded instructions

# Conclusion

**We can use past history to predict if a PC is a branch or not.**

**Saturated counters are used to capture steady-state behavior that can have occasional anomalies.**

**Often, the context determines the direction of branches. A global history predictor such as GAp or PAp is useful in this case.**

**The BTB records the branch type and the target.**

**Decode time optimizations are required to increase performance.**



# The End