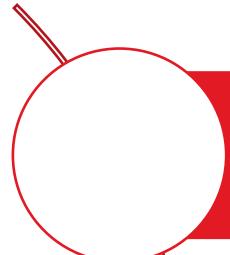
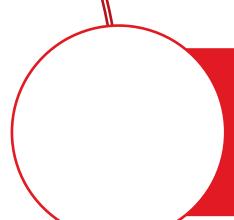


# Chapter 4: Issue, Execute, and Commit Stages

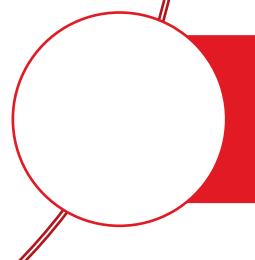
# Background Required to Understand this Chapter



In-Order Pipelines

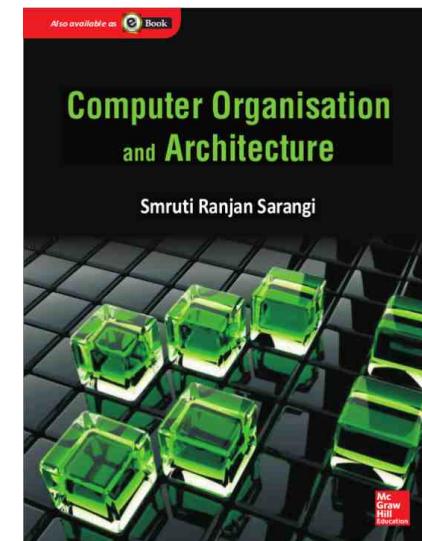
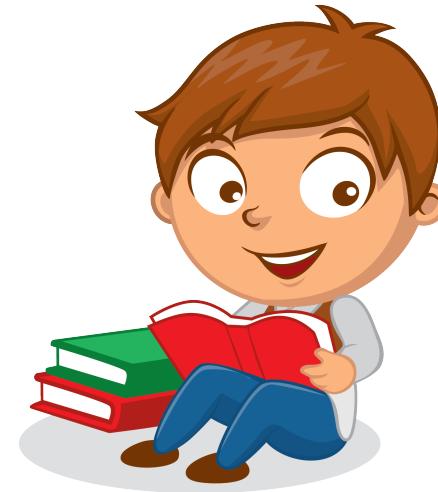


Precise Exceptions



ooo Pipelines: Basic idea

<http://www.cse.iitd.ac.in/~srsarangi/archbooksoft.html>



# Contents

- 
- 1. Instruction Renaming**
  - 2. Instruction Dispatch, Wakeup, Select**
  - 3. The Load-Store Queue (LSQ)**
  - 4. Instruction Commit**

# Renaming (architectural registers $\rightleftharpoons$ physical registers)

Consider a 4-issue processor

- 4 instructions need to be renamed each cycle
- 4 instructions have 8 read operands, and 4 write operands
- Each read operand, needs to read its value from a physical register
- Each write operand needs to be assigned a new physical register
- Dependences between the instructions need to be taken care of

There are two main approaches

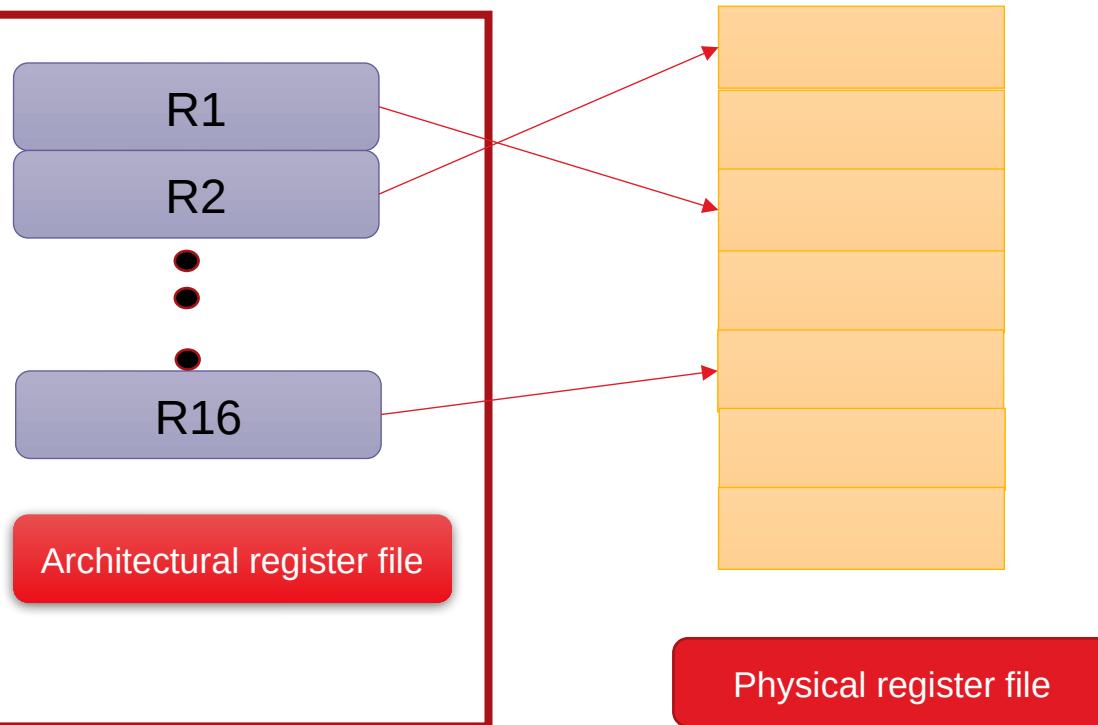
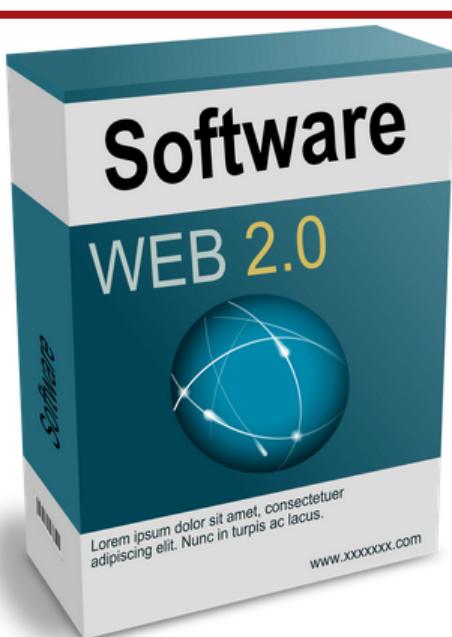
- Renaming with a physical register file
- Renaming with reservation stations

# Renaming with a physical register file

Every ISA has a set of registers that are visible to software

- These are called **architectural registers**
- x86 has 8
- x86-64 has 16
- ARM has 16
- MIPS has 32
- To ensure precise exceptions, the architectural registers should appear to be updated **in-order**
- However, inside the **processor** we need to do renaming to eliminate **output** and **anti dependences**

# Physical Register File and the Process of Renaming

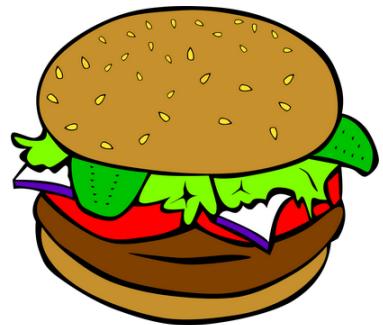


# Physical Register File

# EXAMPLE



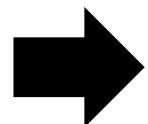
- Consider a system with 16 **architectural registers** and 128 **physical registers**
- In this case, architectural registers exist **in principle**. There is **no** separate architectural register file.
- Each architectural register is **mapped** to one and only one physical register at any point in time.
- 112 physical registers are **unmapped**. They either are empty or contain a value that will be **overwritten** by a later instruction



Why are 112 registers unmapped?

# Example

```
mov r1, 1  
add r1, r2, r3  
add r4, r1, 1  
mov r2, 5  
add r6, r2, r8  
mov r1, 8  
add r9, r1, r2
```



```
mov p11, 1  
add p12, p2, p3  
add p41, p12, 1  
mov p21, 5  
add p61, p21, p8  
mov p13, 8  
add p91, p13, p21
```

- $r$  series  $\sqsubseteq$  architecture,  $p$  series  $\sqsubseteq$  physical
- At the beginning  $ri$  is **mapped** to  $pi$
- With time, the mapping **changes**
- $p12$  contains the value of  $r1$  till the value is **overwritten**.  $p13$  contains the **new** value. At this point,  $p12$  contains an **interim value**

## Example - II

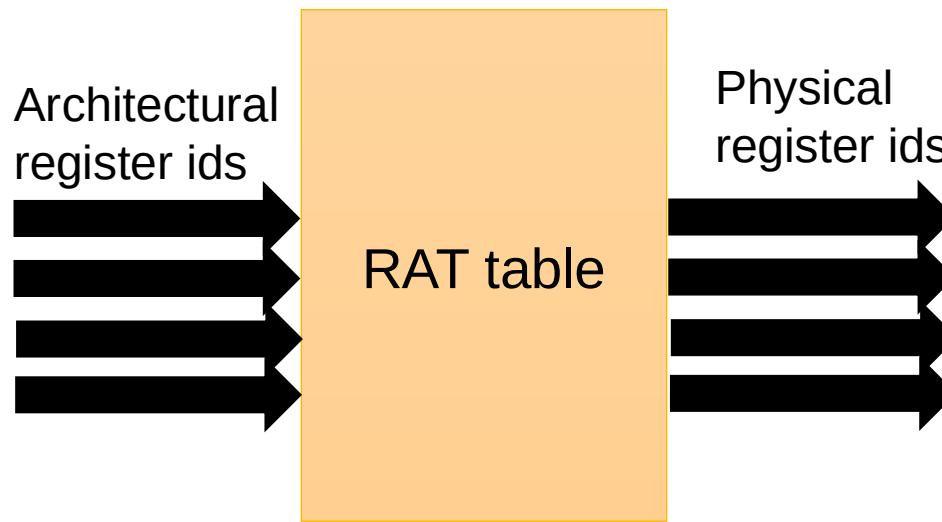
- We need 16 **physical registers** to contain the **latest** values of architectural registers
- 112 **physical registers** can possibly contain interim values that might be required by instructions in the OOO pipeline (also called **in-flight instructions** like p12)
- More is the number of physical registers  $\Rightarrow$  higher is the number of instructions that can simultaneously be **in flight**  
 $\Rightarrow$  more is the ILP
- Renaming: convert architectural register id  $\Rightarrow$  physical register id

# Rename Stage

## Three Structures

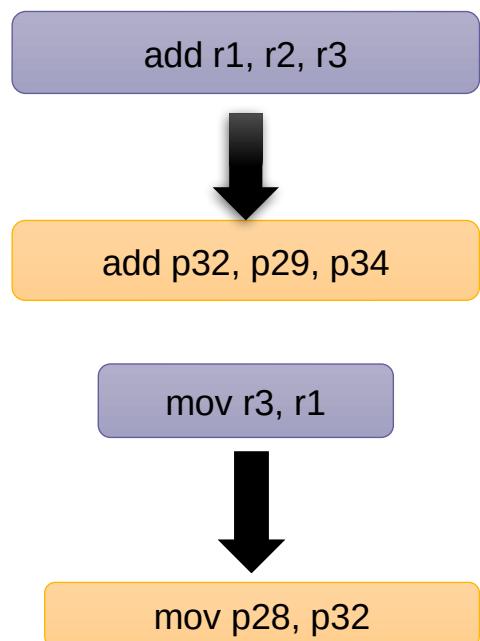
- **Register Alias Table (RAT)** □ Translate architectural register id to physical register id
- **Dependence Check Logic (DCL)** □ Take care of dependences between instructions renamed in the same cycle
- **Free List or Free Queue** □ Maintain a list of unmapped physical registers

# Register Alias Table



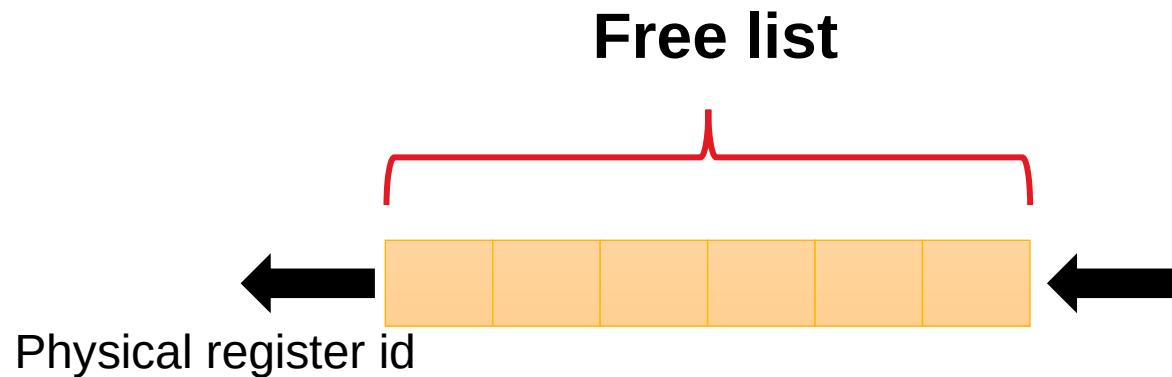
- **Rename** multiple instructions at once
- We get a **stream** of physical register ids
  - Eliminates WAW and WAR hazards

# Operation of the RAT Table



## Free List

- It is a queue of unmapped physical register file ids
- An entry is dequeued from it when we need to assign it to an architectural register id
- When is an entry added to it? **LATER**
- **Update the RAT with the new mapping**

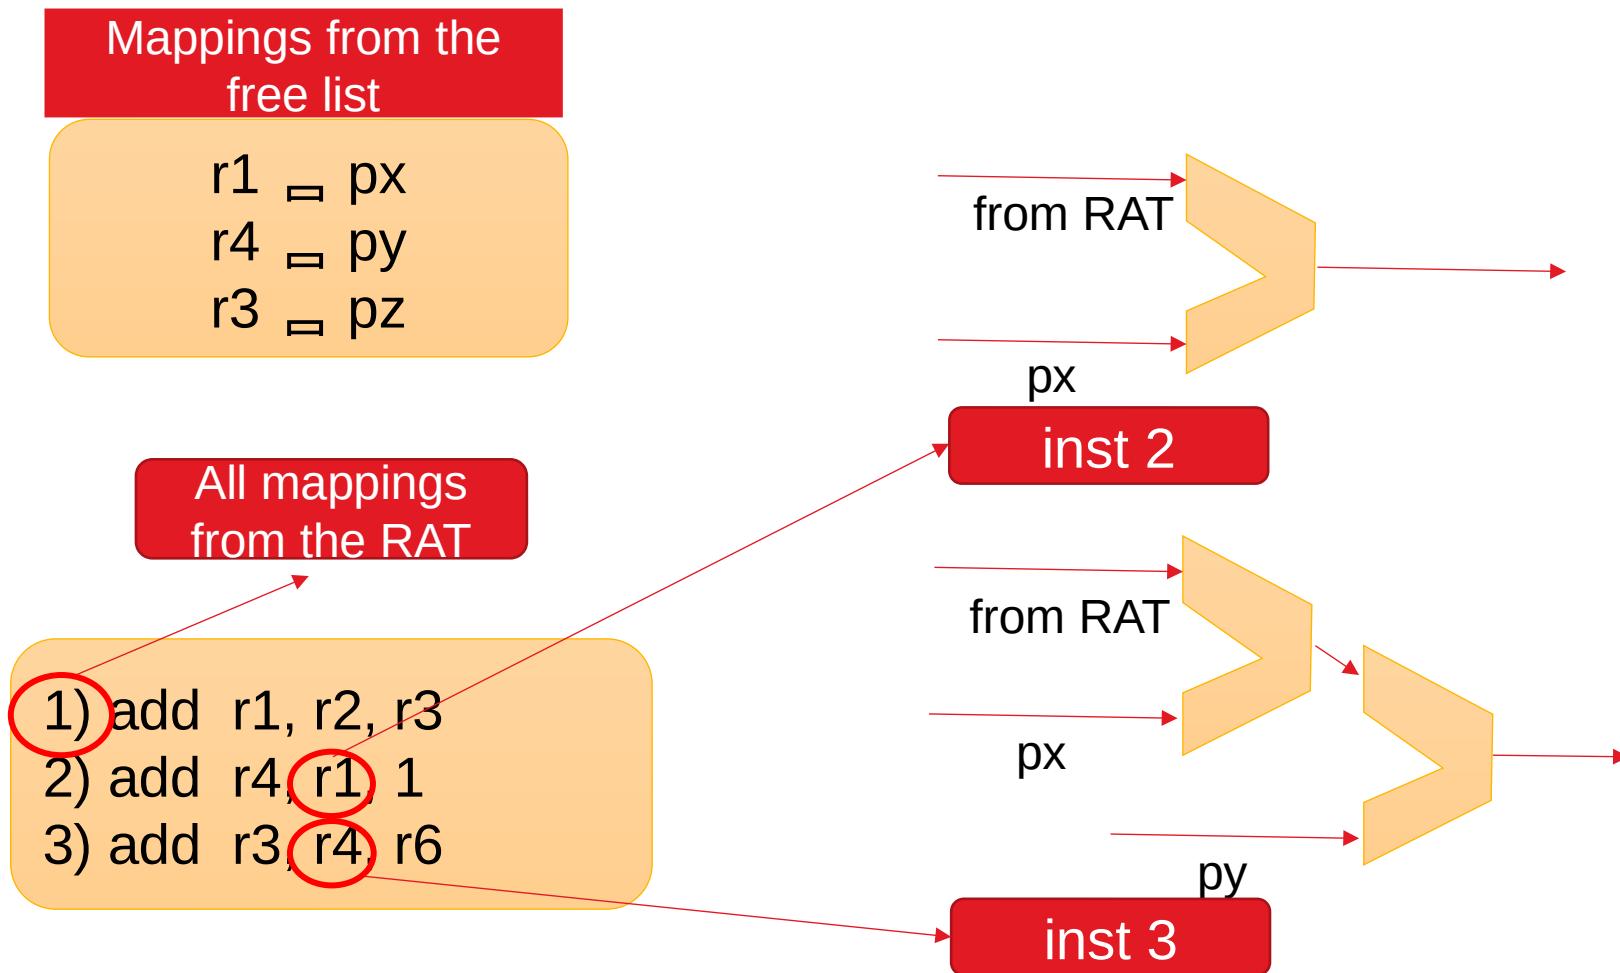


# Dependence Check Logic

```
add r1, r2, r3  
add r4, r1, 1
```

- Let us say  $r1$  is mapped to  $p11$
- The second instruction needs to read the mapping of  $r1$  as  $p11$
- What if both instructions are sent to the RAT together?

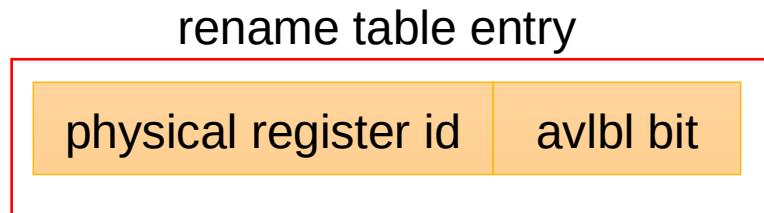
# Dependence Check Logic - II



# Available bit in the rename table entry

We need some more **information** in each entry

- An **available** bit (*avlbl*)
- It indicates whether the result can be found in the **physical register file and the rest of the pipeline** or not.
  - If not, the instruction has to **wait** in the instruction window.



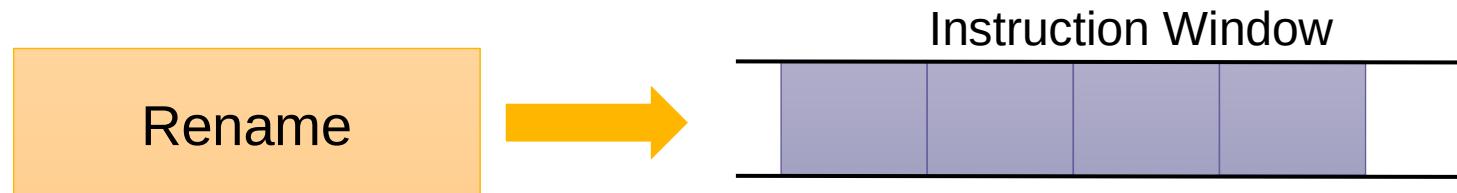
# Contents

- 
1. **Instruction Renaming**
  2. **Instruction Dispatch, Wakeup, Select**
  3. **The Load-Store Queue (LSQ)**
  4. **Instruction Commit**

# Dispatch

Till the rename stage, instructions proceed **in-order**.

**Dispatch**  $\Rightarrow$  Send the instruction to the instruction window (IW)



# Instruction Window

Status of an instruction in the window

valid	opcode	src tag 1	ready bit 1	src tag 2	ready bit 2	dest tag
		imm1		imm2		

**opcode** = type of operation: add, subtract, multiply

**tag** = physical register id

**ready bit** = does the physical register file contain the corresponding value

# At any point in the instruction window

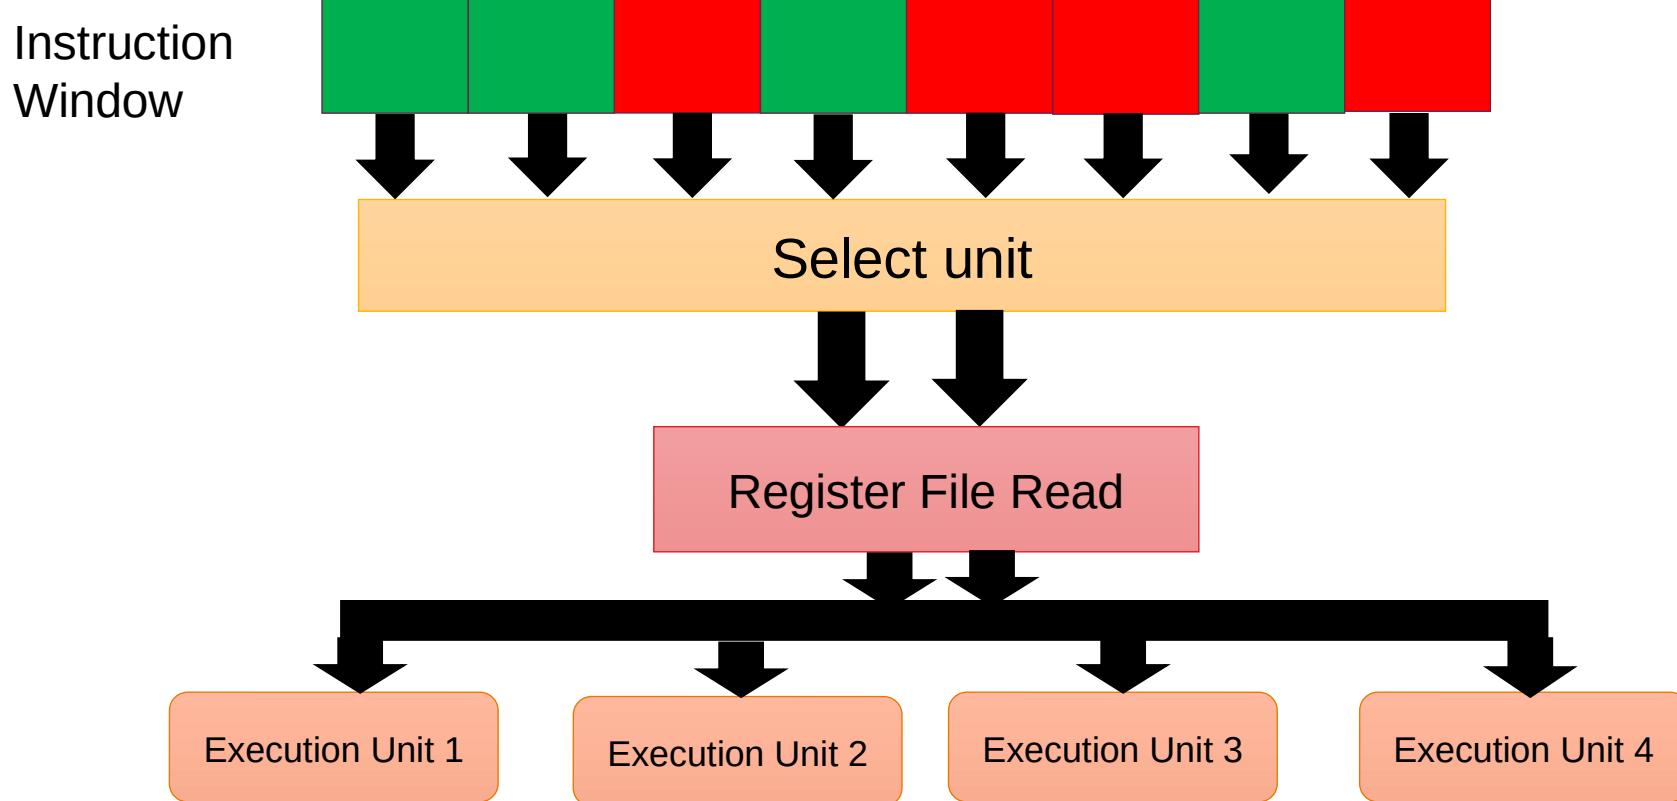


Instructions waiting for an operand to be ready



Instructions whose operands are ready

# Select and Execute



# Wakeup

- How does an instruction in the IW know that its operands are ready?

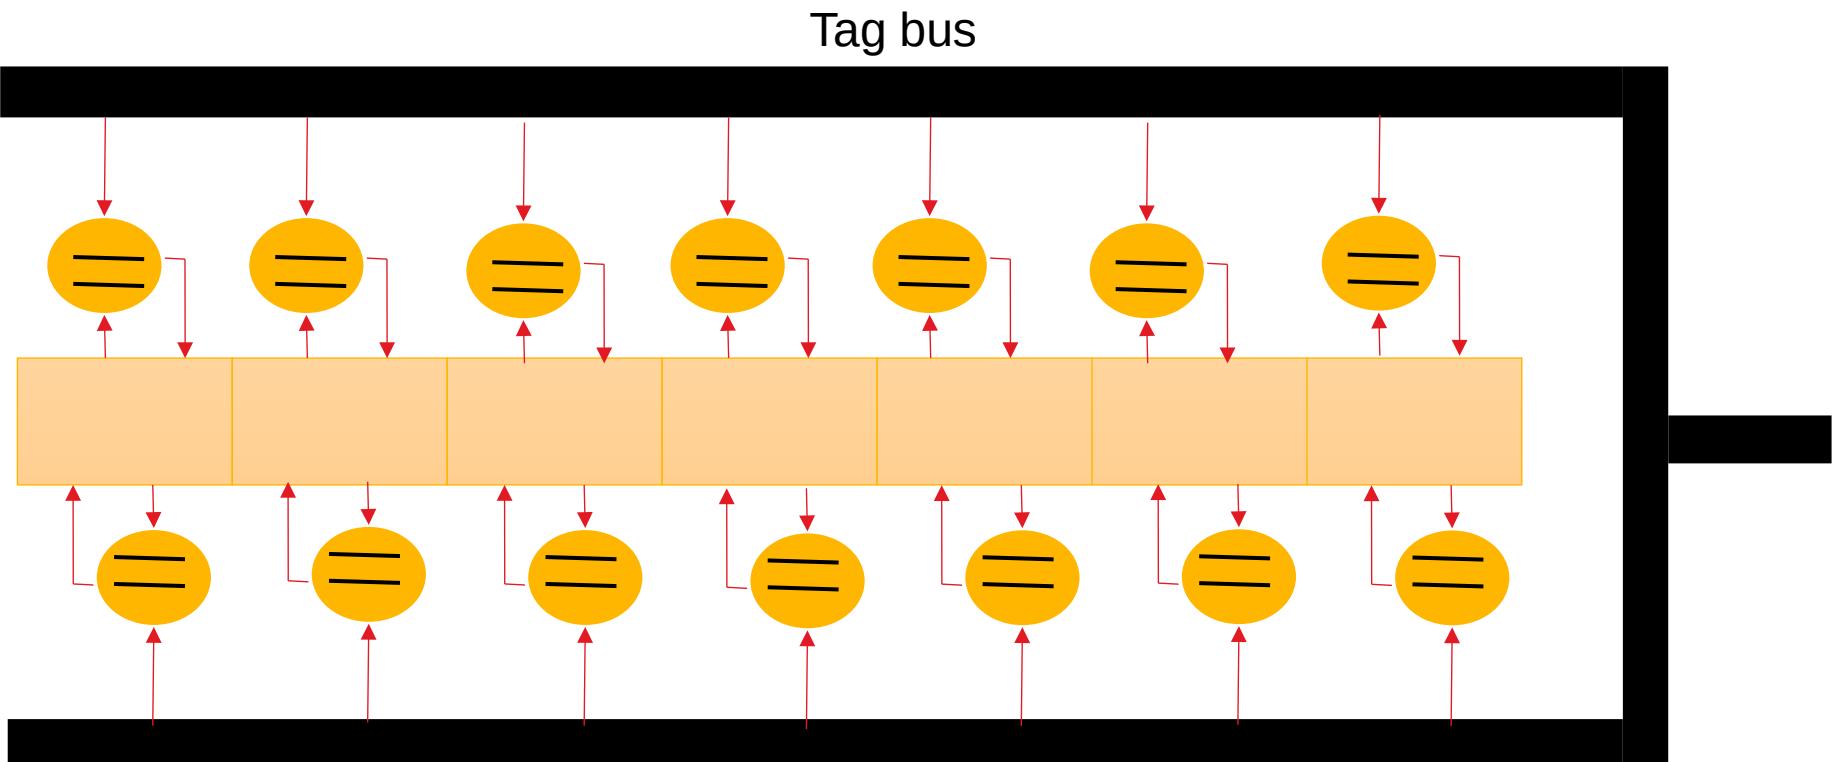


The producer instructions need to let it know

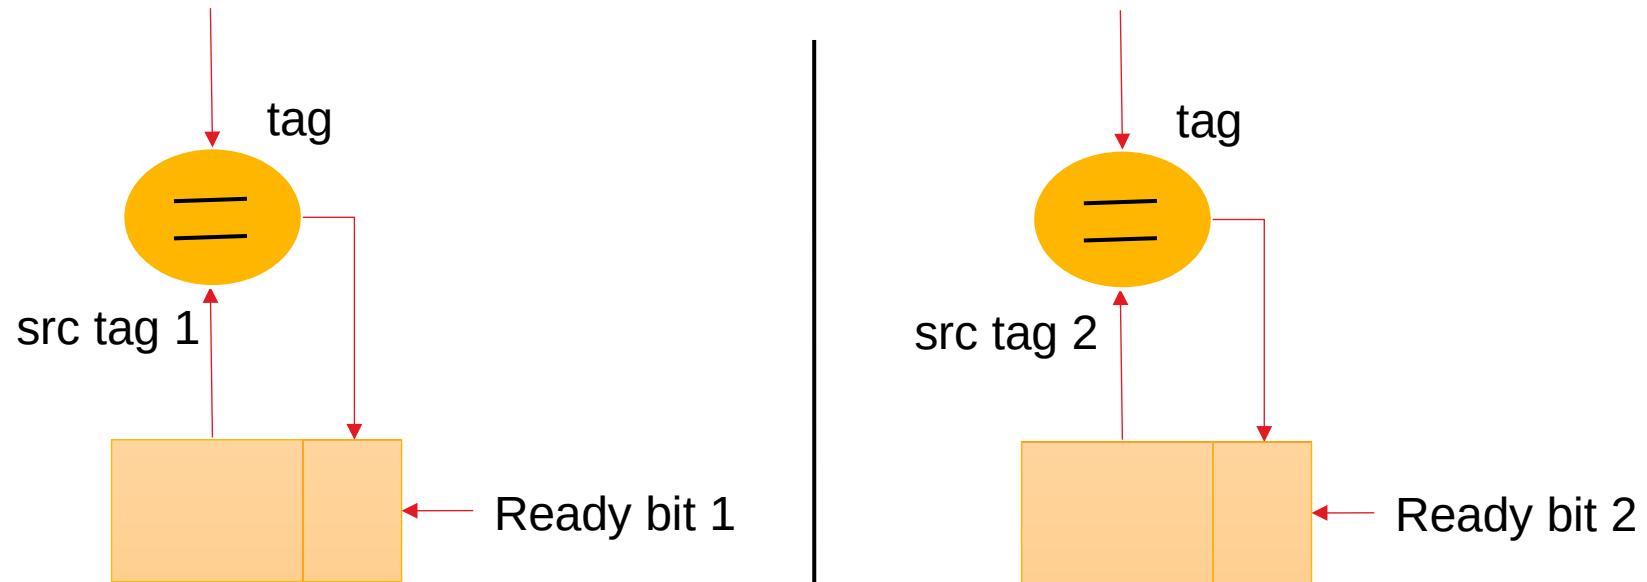
Once a **producer** finishes executing

- Broadcasts the **tag** (destination physical register id) to the entries of the IW
- Each **entry** marks its corresponding source operand as **ready** if the tag matches
- This is called **wakeup**

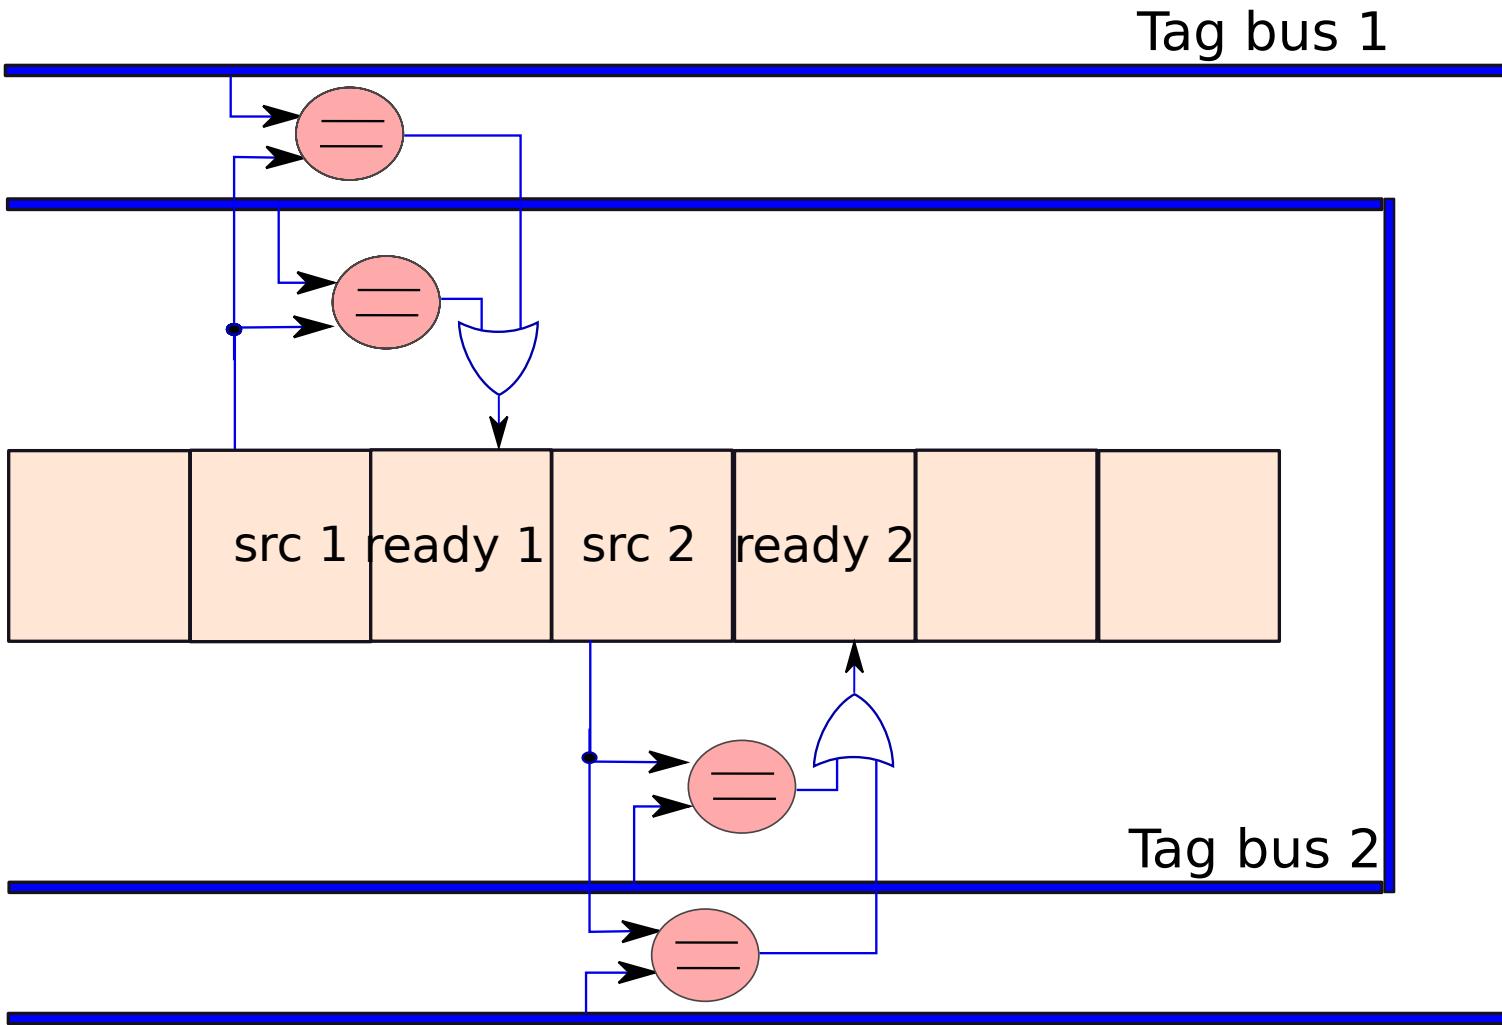
# Instruction Window



## Instruction Window - II

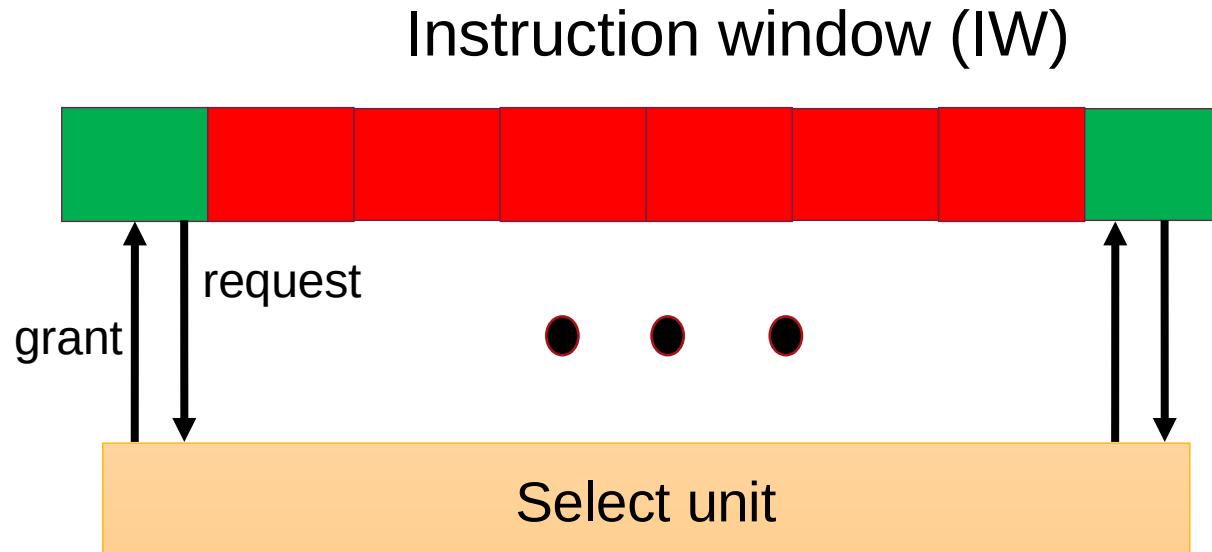


# Broadcast multiple tags at once



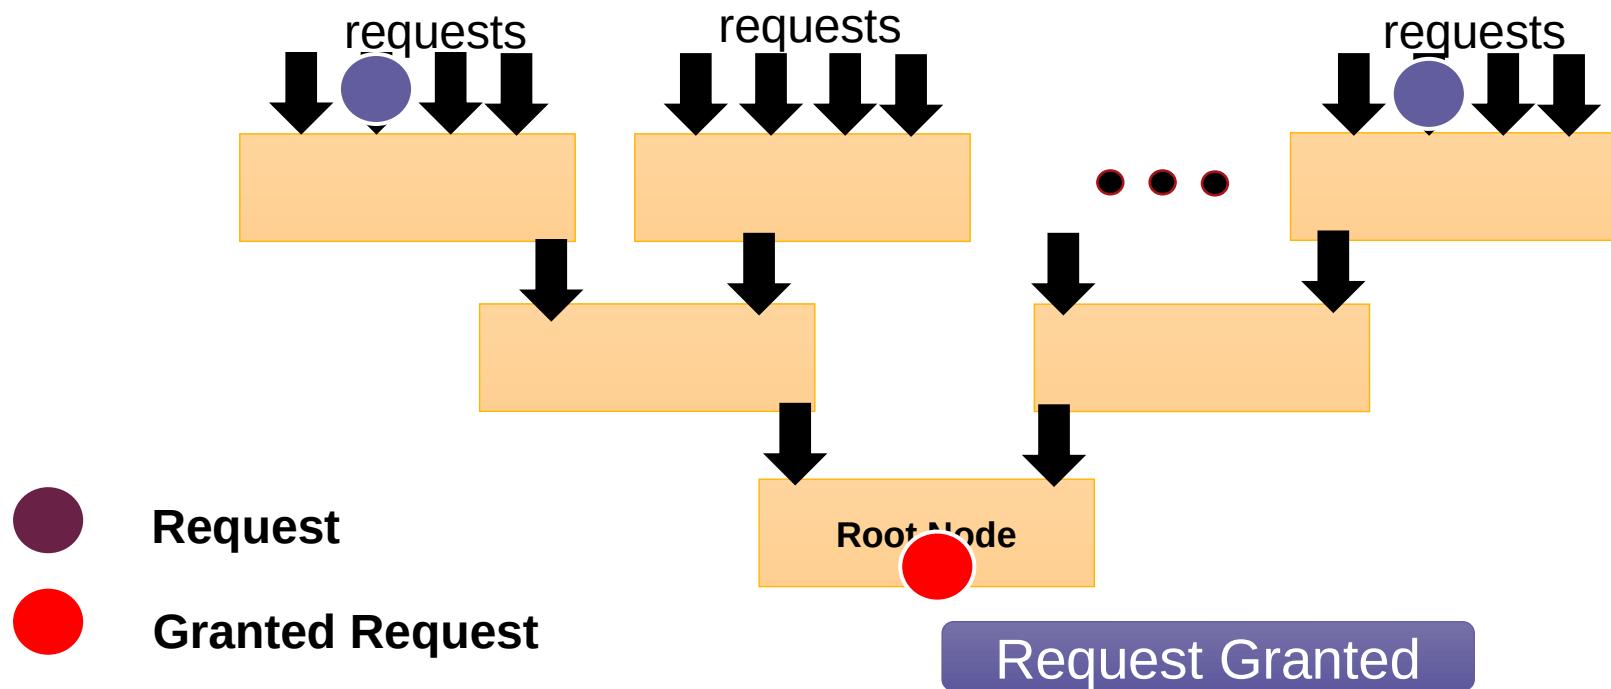
# Instruction Select

# Select Unit



- Once an IW entry has all its operands **ready**, it is marked as **ready**.
- It asserts its attached **grant line**
- The select unit **chooses** one among the ready instructions

# Tree Structured Select Unit: Each node is a choice box

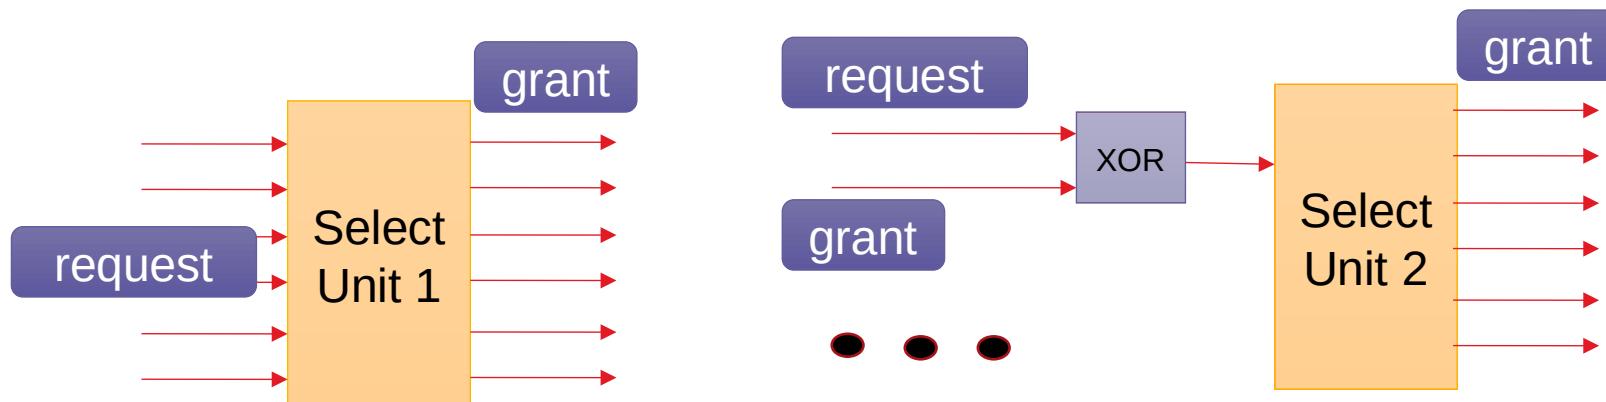


# Select Unit for Multiple Resources

What if we have 2 adders, 2 multipliers, and 1 divider?

## Solution:

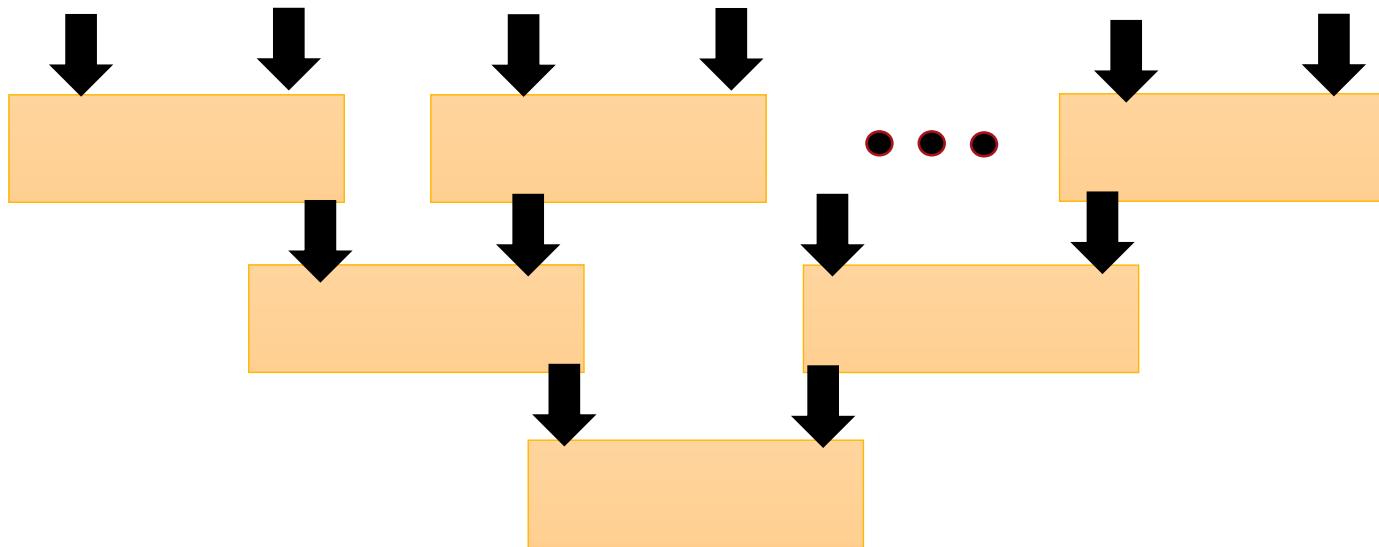
- Have a **separate** select unit for each class of instructions
- For a select unit that caters to multiple FUs (functional units)
- Three **options**: Option 1 □ **chain select units**



## Option 2: Multiple request signals

- Send **multiple** request signals (up to  $n$ )
- Every level sends at most  $n$  **requests** to the root
- The root selects at most  $n$  **requests**, and **issues** grant signals
- This circuit is more complicated than serial select circuits
  - The **choice boxes** need more elaborate logic

## Option 3: Asymmetric Select Units



- Have two select units with **different policies**.
- **Select unit 1:**
  - Each choice box gives priority to its **left input**
- **Select unit 2:**
  - Each choice box gives priority to its **right input**

# Select Policies

The select policy has an important implication on performance. We should select those instructions that are the most performance critical.

- Random
- Oldest First

## Type of opcode

- Load instructions typically have higher priority than stores.
- Other priorities can also be learnt based on execution patterns.



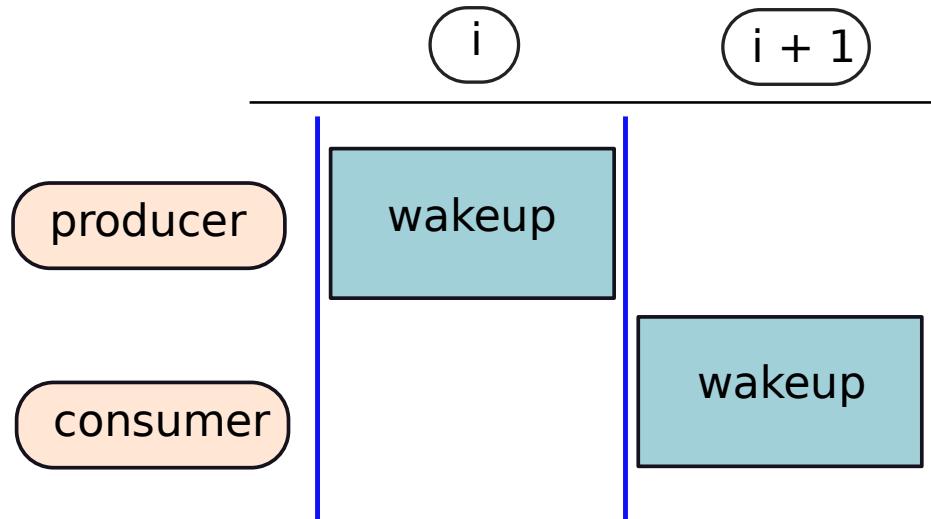
# Tag Broadcast

# RAW Dependencies in our Pipeline

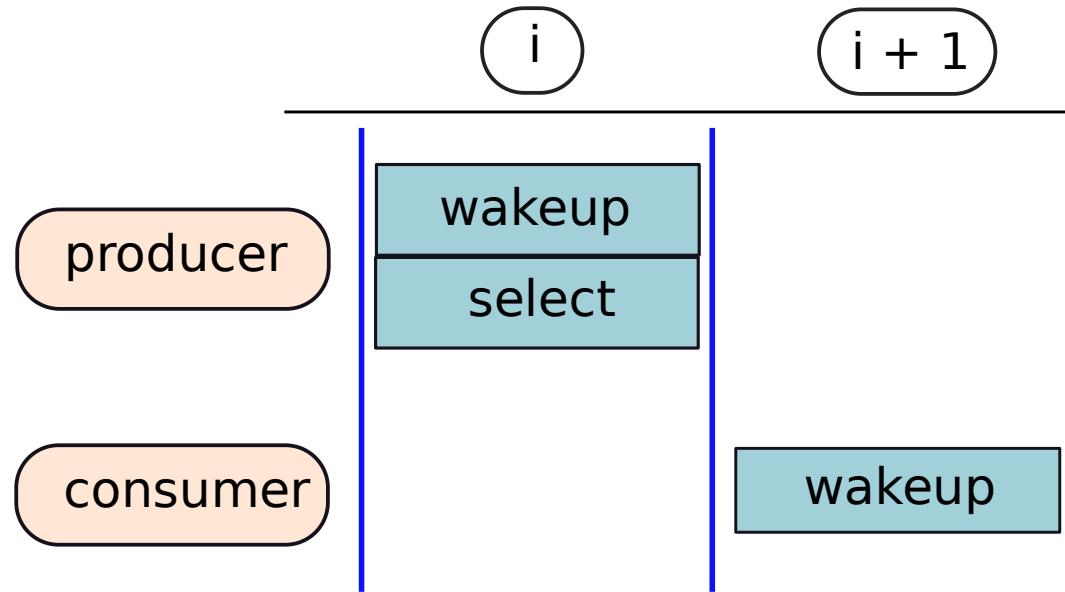
Consider consecutive instructions with RAW dependences

- 1) add r1, r2, r3
- 2) add r4, r1, r5

We want **back-to-back execution** – Execution in consecutive cycles.

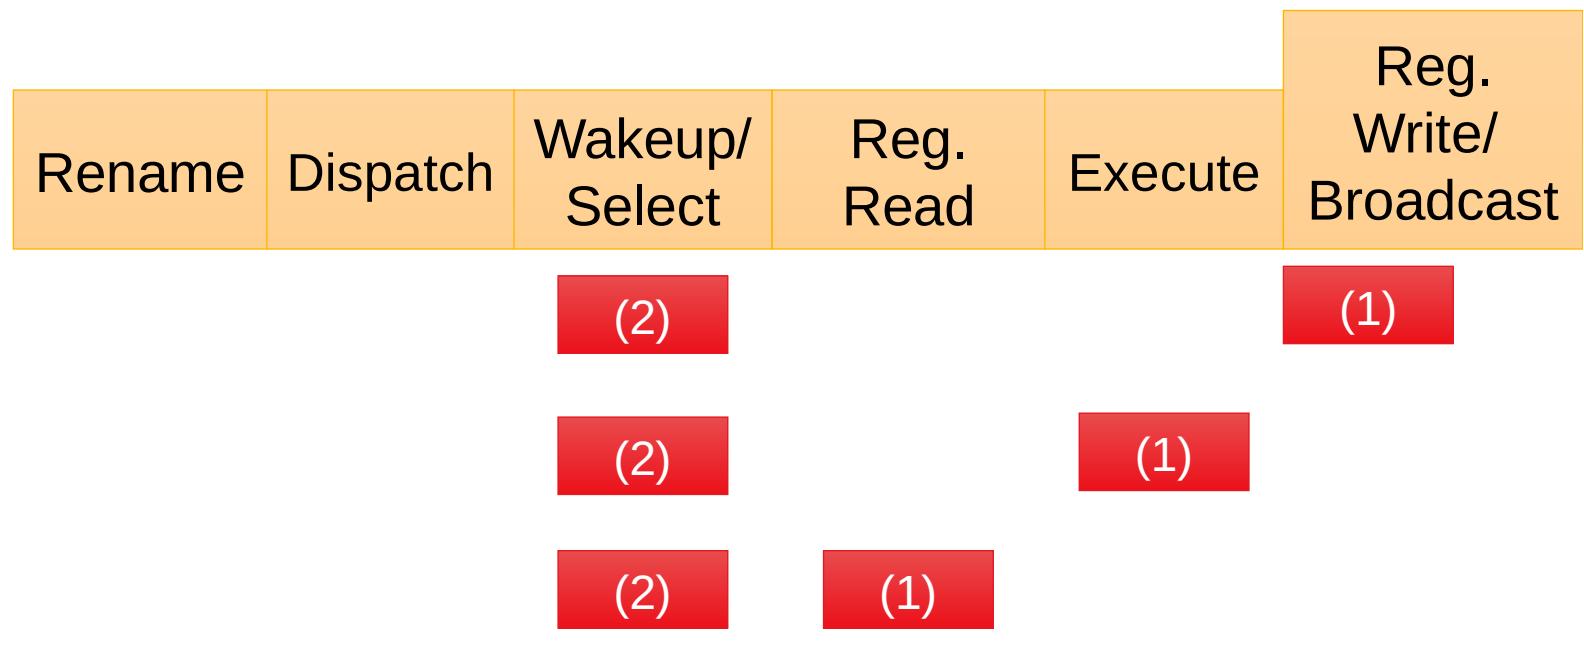


# What does this imply?



Assume we have slightly **optimized** our pipeline to have wakeup and select in the same cycle.

# When should we broadcast the tag?



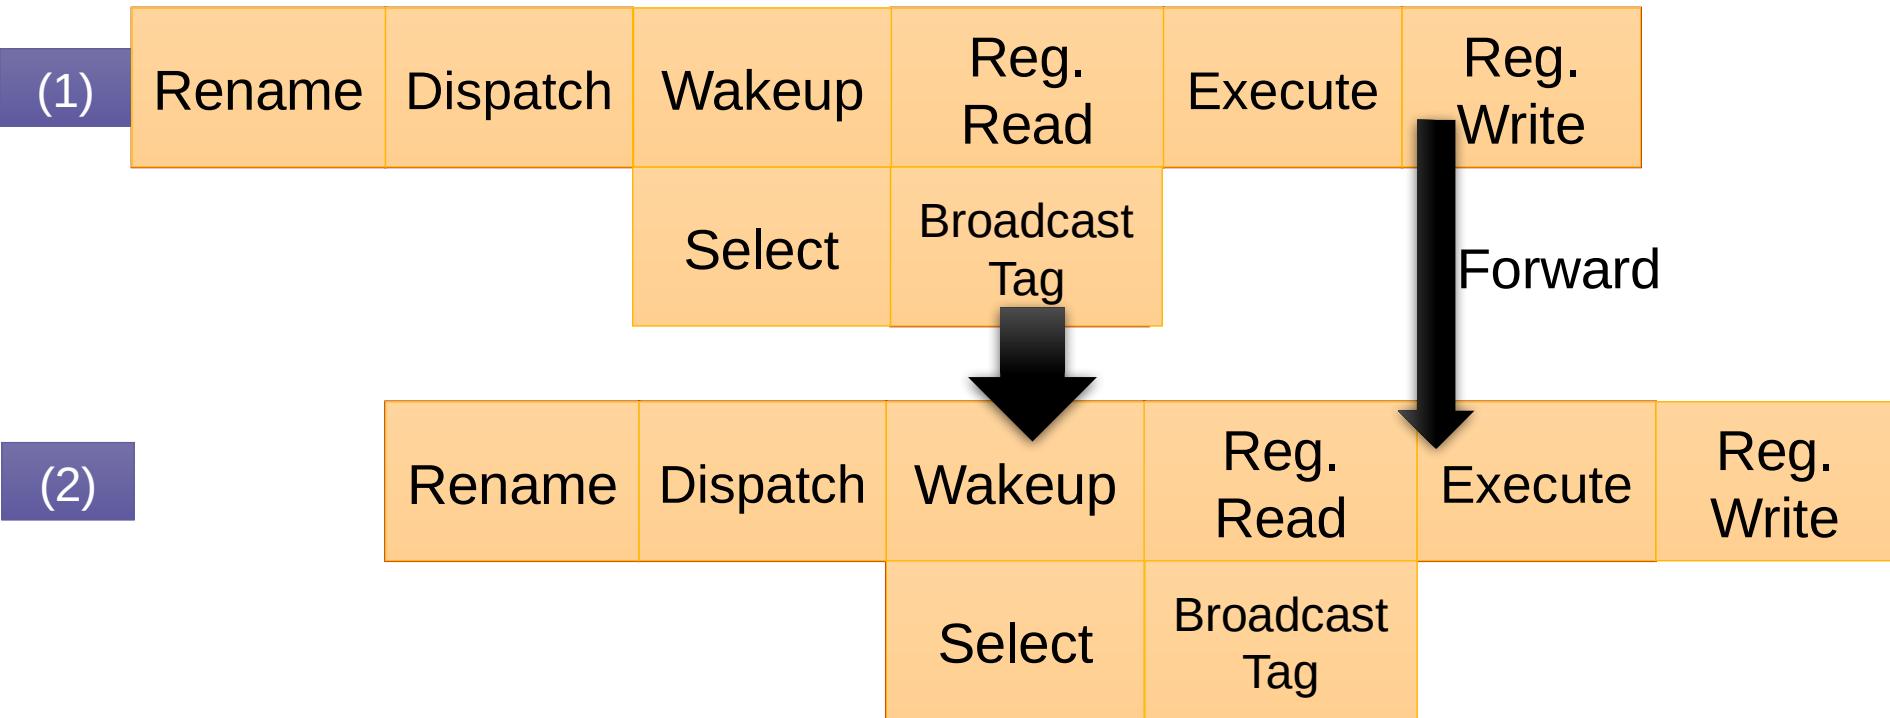
# Let us consider Option 3 – Early Broadcast

- Allows back-to-back (in consecutive cycles) execution

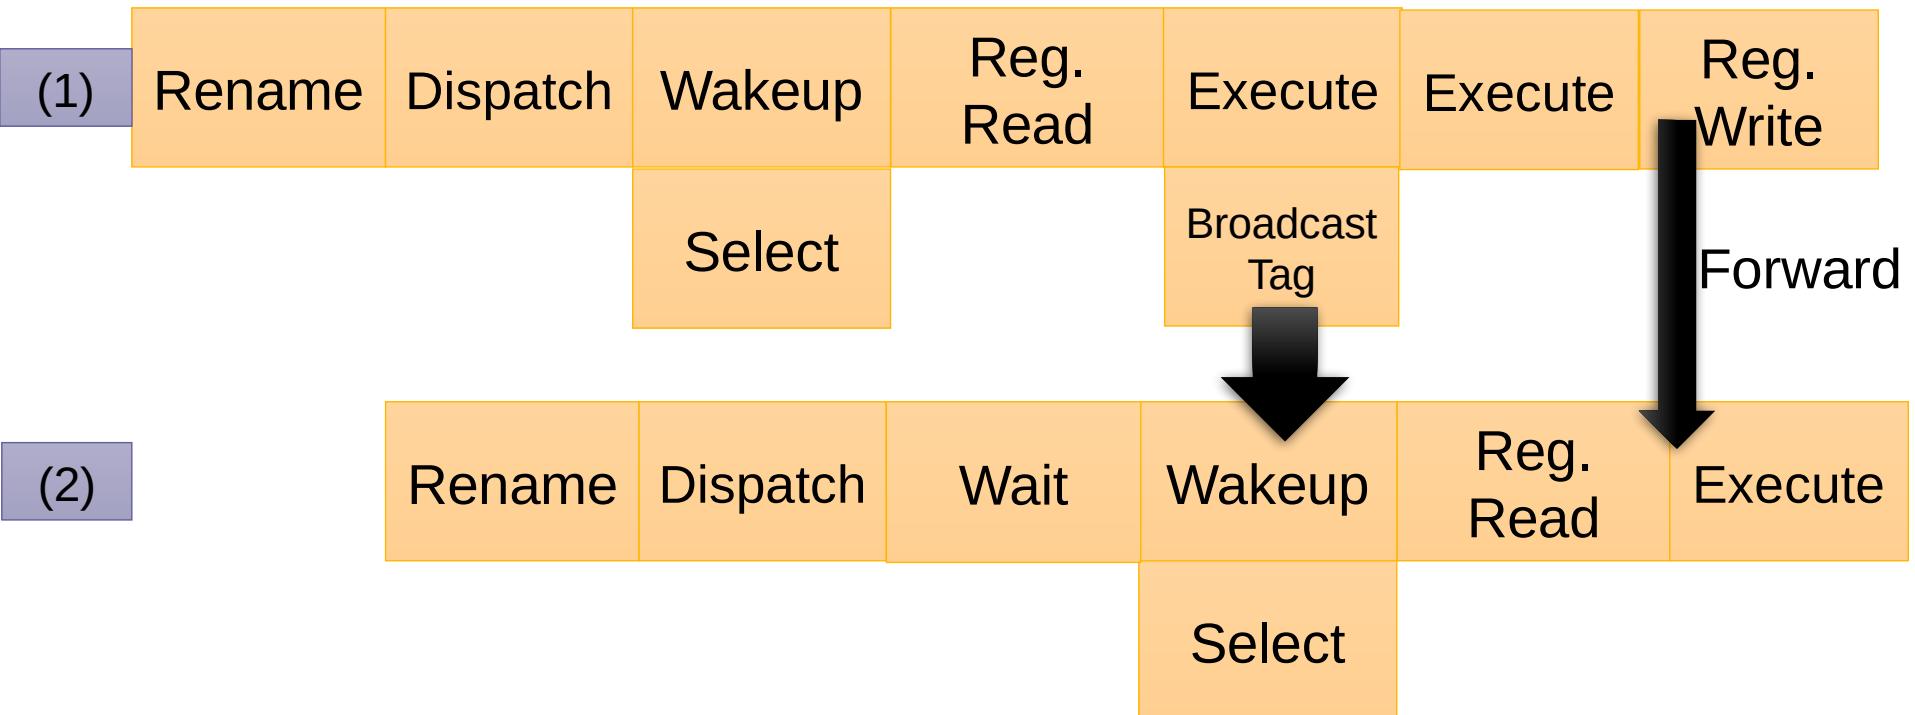
## What needs to be done:

- When (1) is in the Reg. read stage, (2) needs to be in the select stage
- This is possible if we broadcast the tag for (1) when it is in the reg-read stage
- (2) picks up the tag, wakes up, and gets selected for the next cycle

# Classic Example of Forwarding in an OOO Pipeline



## 2-cycle latency for producer: Load-use Hazard



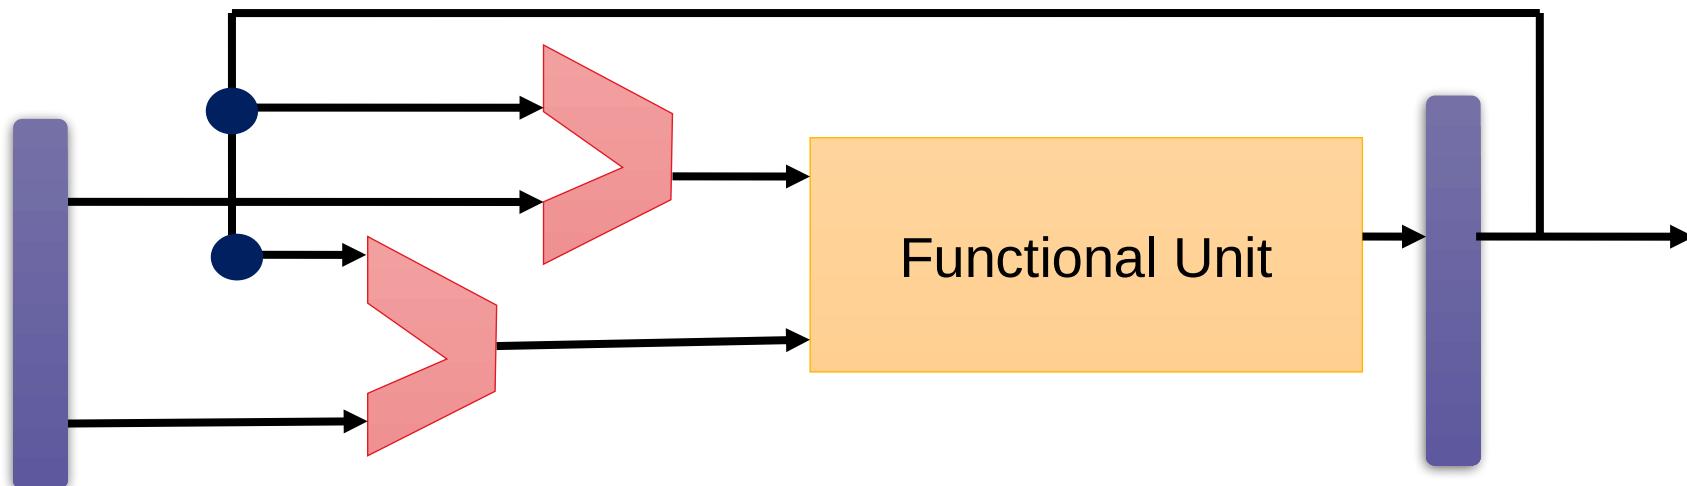
# Structure of the Pipeline



- For **different** instructions, we broadcast the tags at different points of time.
- If the **execution** requires  $k$  cycles, we **broadcast** the tag  $k$  cycles after getting selected.

# Forwarding Logic in the Execution Unit

Also called **bypass** paths

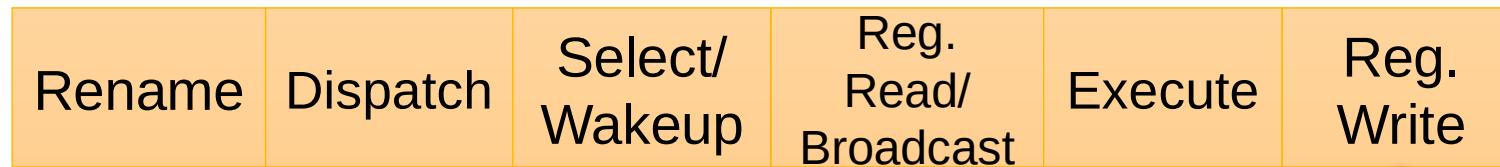


from the  
Register File

# When should we mark the *avbl* bit?



**Option 1:** Along with the register write

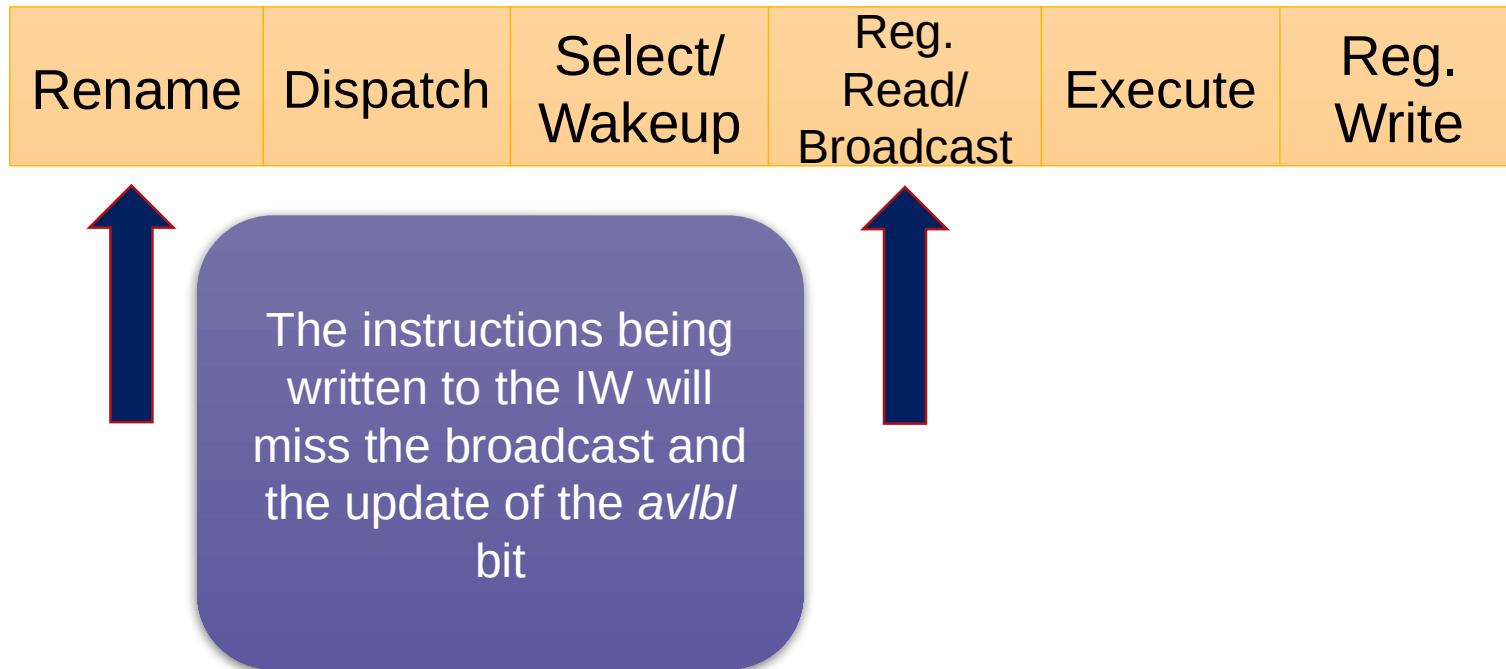


Some instructions that have arrived after the broadcast might wait forever

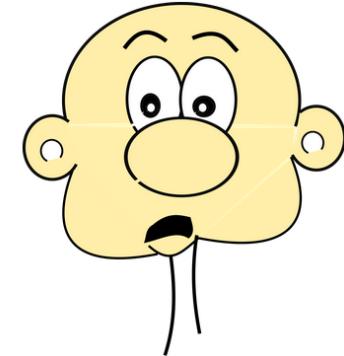


# When should we mark the available bit?

**Option 2:** Along with the broadcast



# What to do?



Let us **fix** solutions 1 and 2

Basic **features** of a new solution:

- Realization: **Instructions** being written to the IW will **miss** the broadcast
- These instructions would have also read the **avlbl** bit as 0
- As a **result**: they will **wait** forever

What does the teacher do if some students in the class are sleeping?

- **ANSWER**: Take one more class

What is the solution here:

- **Double broadcast**

# Double Broadcast

Both **solutions 1** and **2** can be **fixed**, if **two** broadcast messages are sent

Let us look at one solution

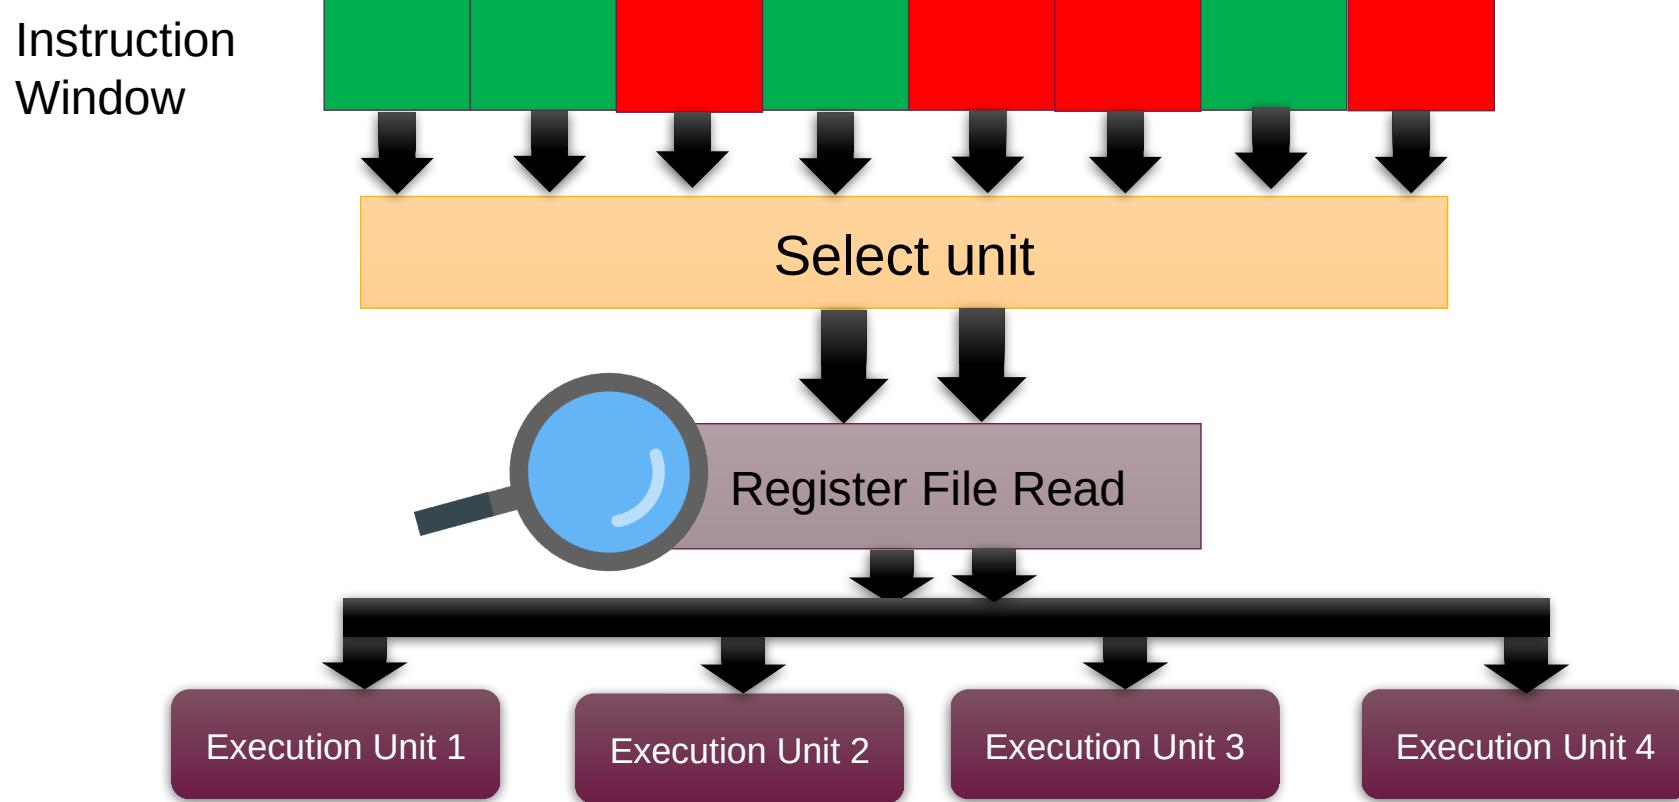
- Set the **avlbl** bit to 1 along with sending the **Broadcast** message
- There might be some **instructions** that miss the **broadcast**
- **Log** the **Broadcasted** tags in a temporary structure
- Also **record** all the instructions that are being dispatched in a separate structure called the *dispatch buffer*

In the **next** cycle

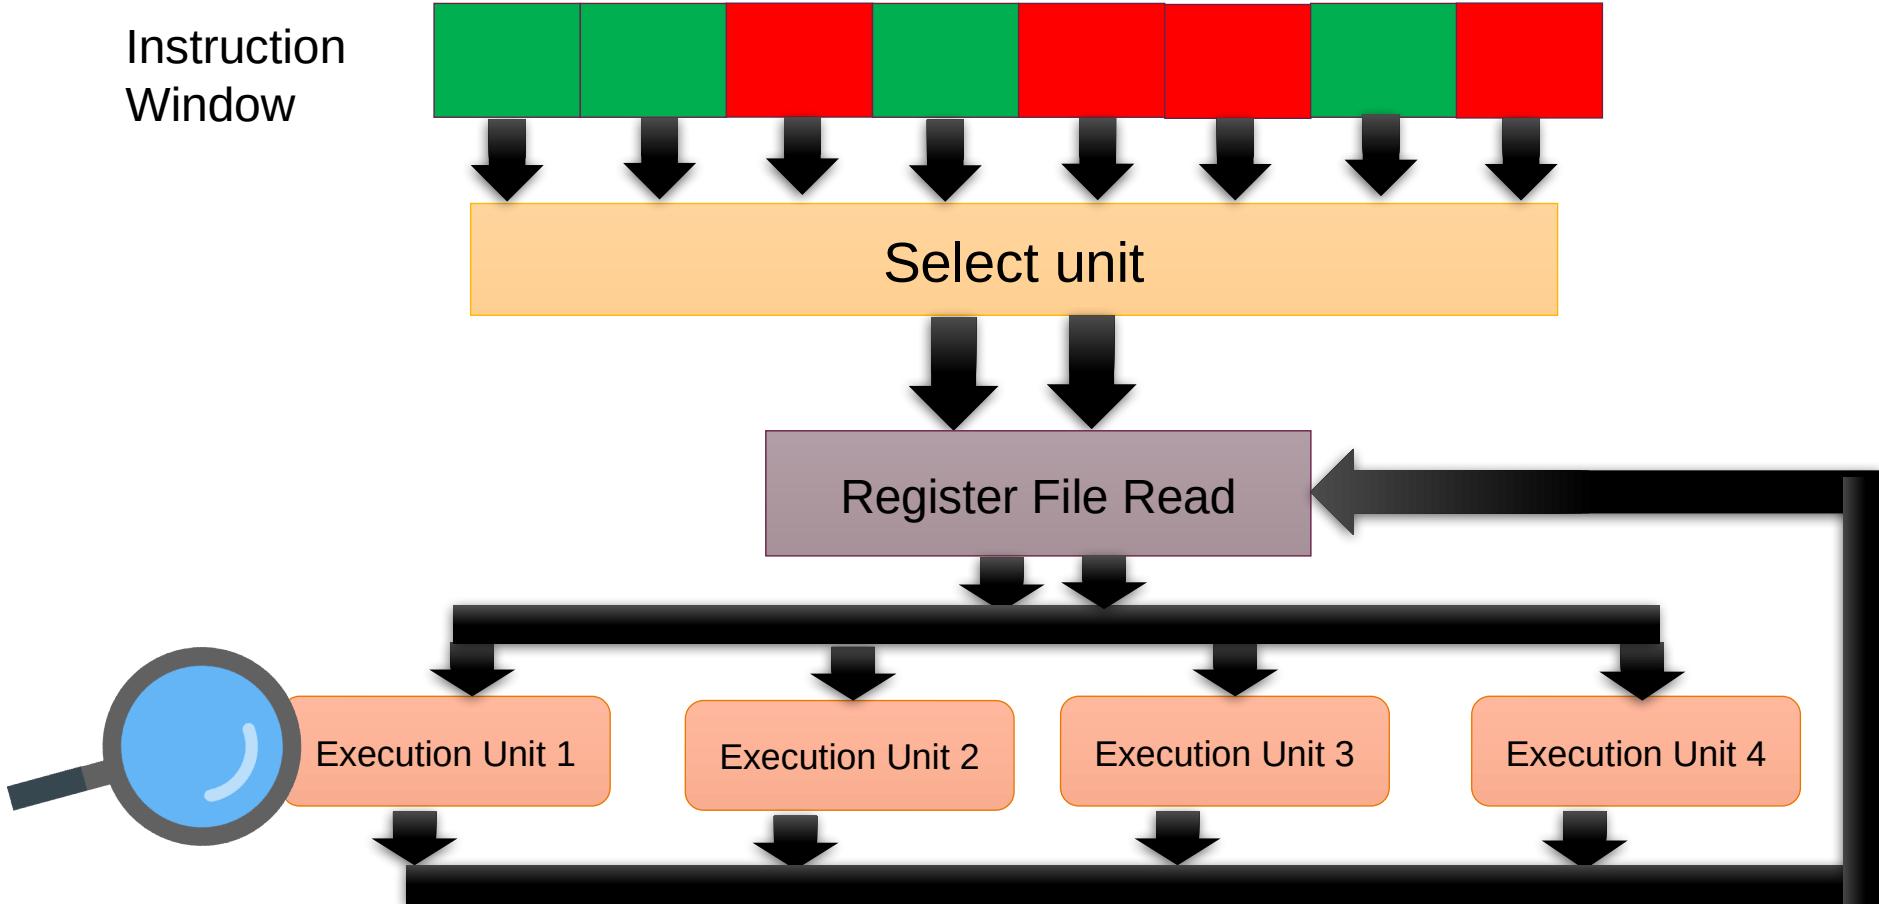
- Match the tags with entries in the dispatch buffer.
- If there is a **match**, update the corresponding IW entry
- **Clear** the entries for the last cycle in the *dispatch buffer*

# Register Read and Execute

# Read operands from the register file and Issue to the execution units



# Execute



# Contents

- 
- 1. Instruction Renaming
  - 2. Instruction Dispatch, Wakeup, Select
  - 3. The Load-Store Queue (LSQ)
  - 4. Instruction Commit

# Case of Load and Store Instructions

```
ld r1, 4[r3]  
st r2, 10[r5]
```

First, the instructions are sent to the adder to compute the effective address. In this case:  $4 + (\text{contents of } r3)$ ,  $10 + (\text{contents of } r5)$

Next, they are sent to a load-store unit

- Sends the loads and stores to the data cache
- Loads can execute immediately
- Stores update the processor's state (remember precise exceptions)
- We cannot afford to write the value of stores on the wrong path
  - Hence, for stores we wait.



# Requirements: Loads and Stores

Stores wait

They need some storage space

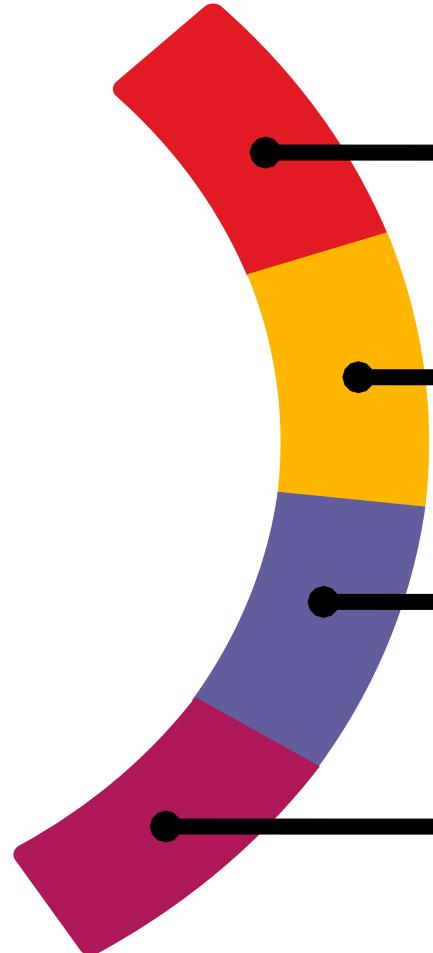
Loads first check this dedicated space

Loads cannot directly be sent to the cache

There might be stores before them that are unresolved

Unresolved stores might be to the same address

# Key requirements



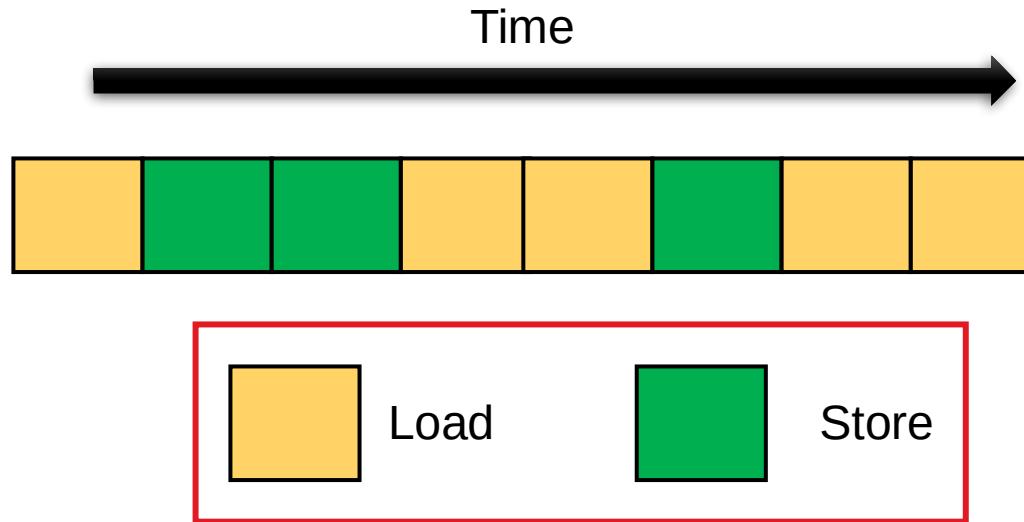
Stores have to be sent to the memory system in **program order**. This is needed for precise exceptions. We need dedicated storage space  $\sqsubset$  store queue

Stores can only be sent when they are guaranteed to be on the **correct path**. A store needs to be the earliest instruction in the pipeline when it is sent.

Loads first **check** the store queue. The store queue is like a cache.

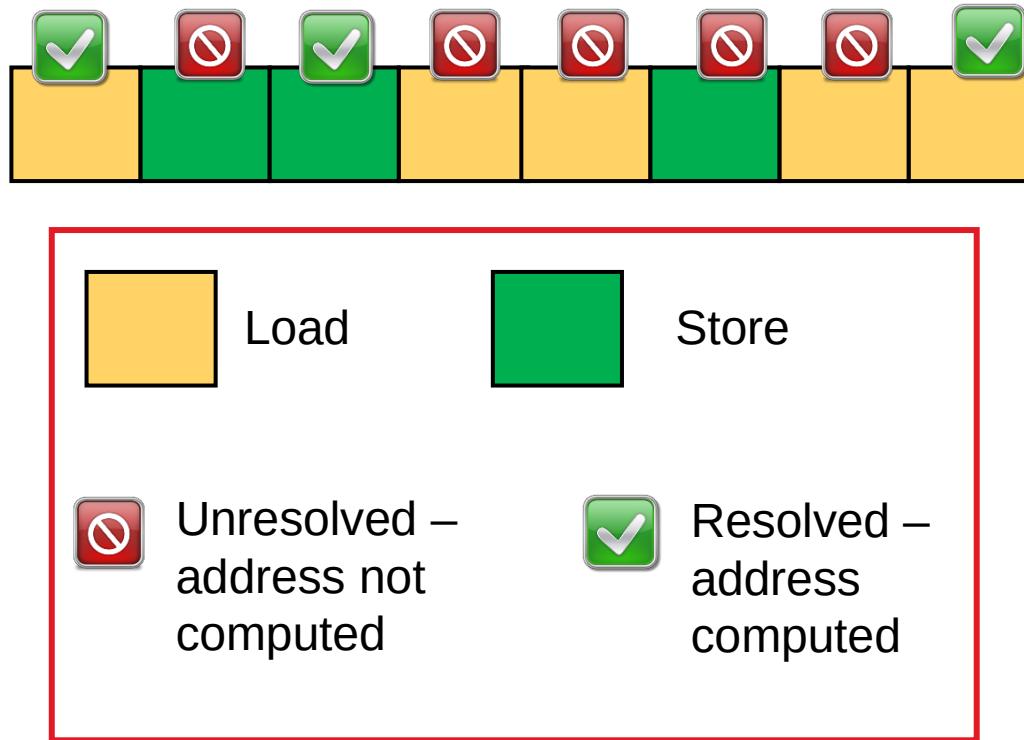
Loads cannot be sent to the memory system until all the stores before it are **resolved** and do not write to the same address.

# Load-Store Queue (LSQ)

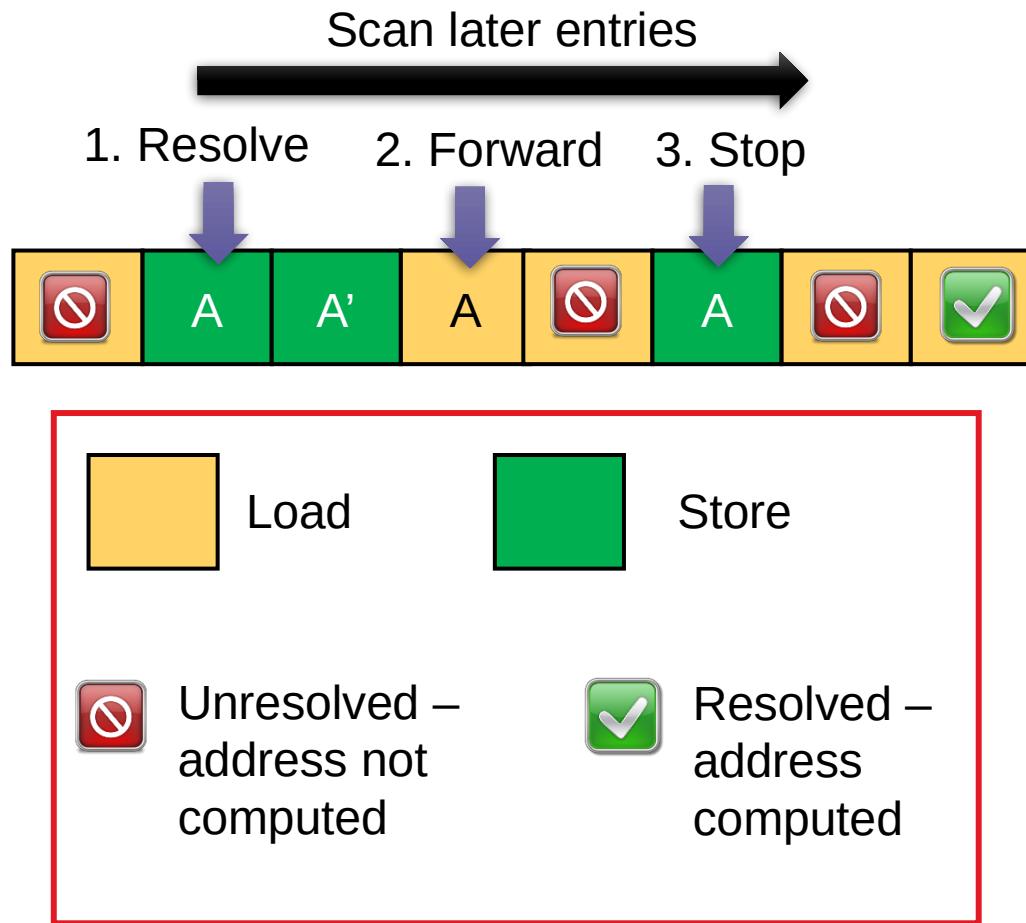


- Allocate an entry at decode time (allocated in order)
- Don't ask when to deallocate an entry (ask later)
- When the effective address is computed, update the entry

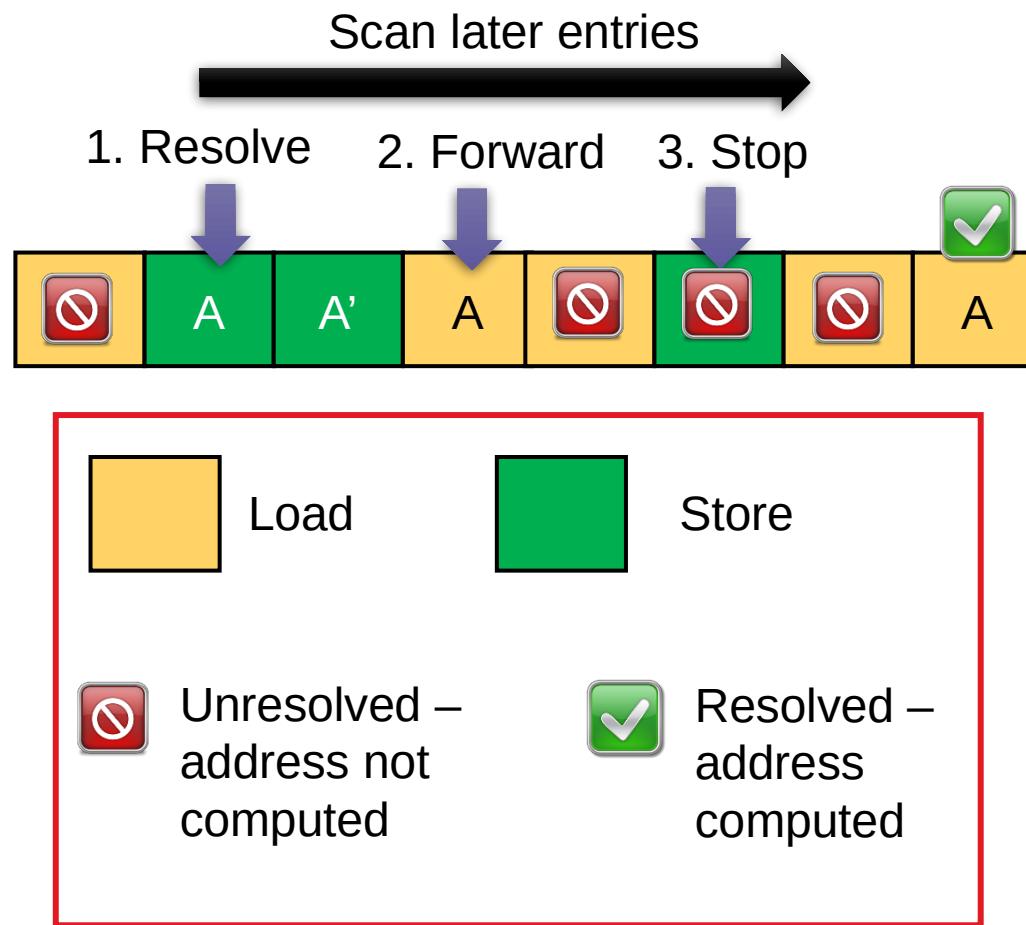
# Resolved and Unresolved Entries



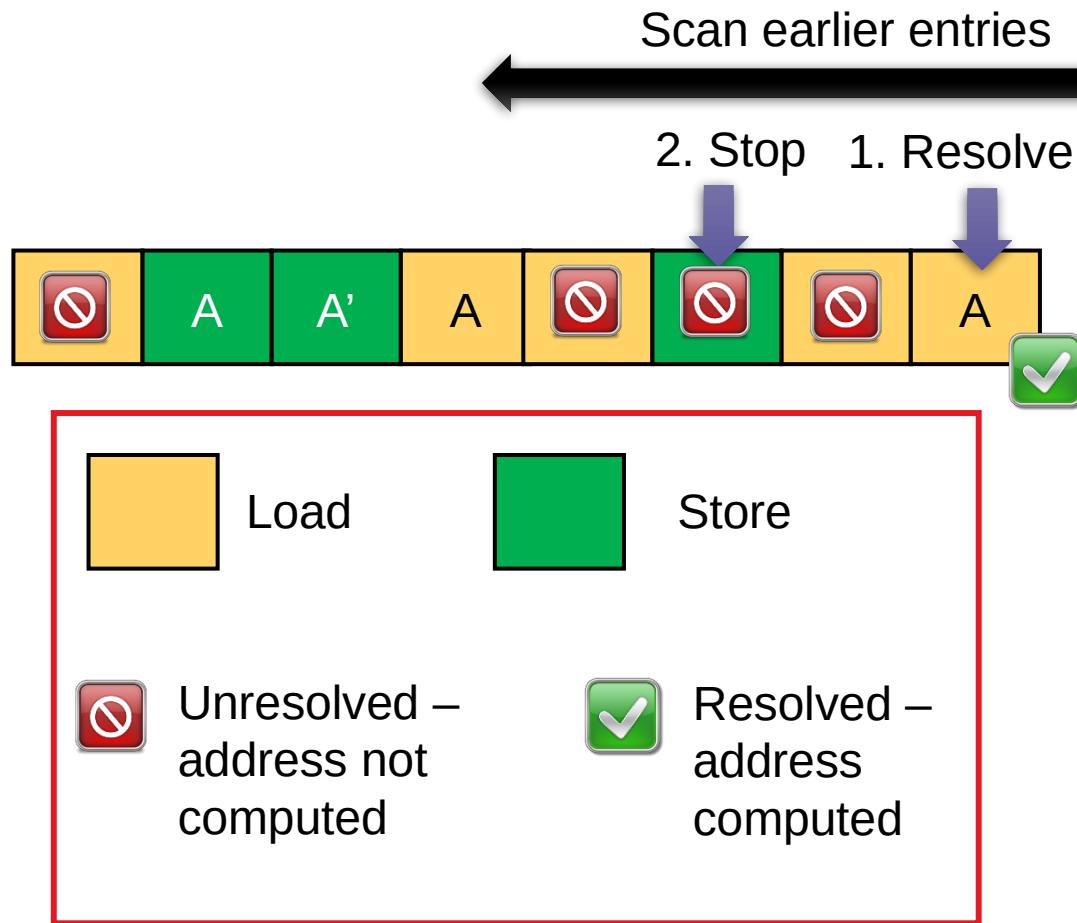
# Compute Address of a Store – Case I



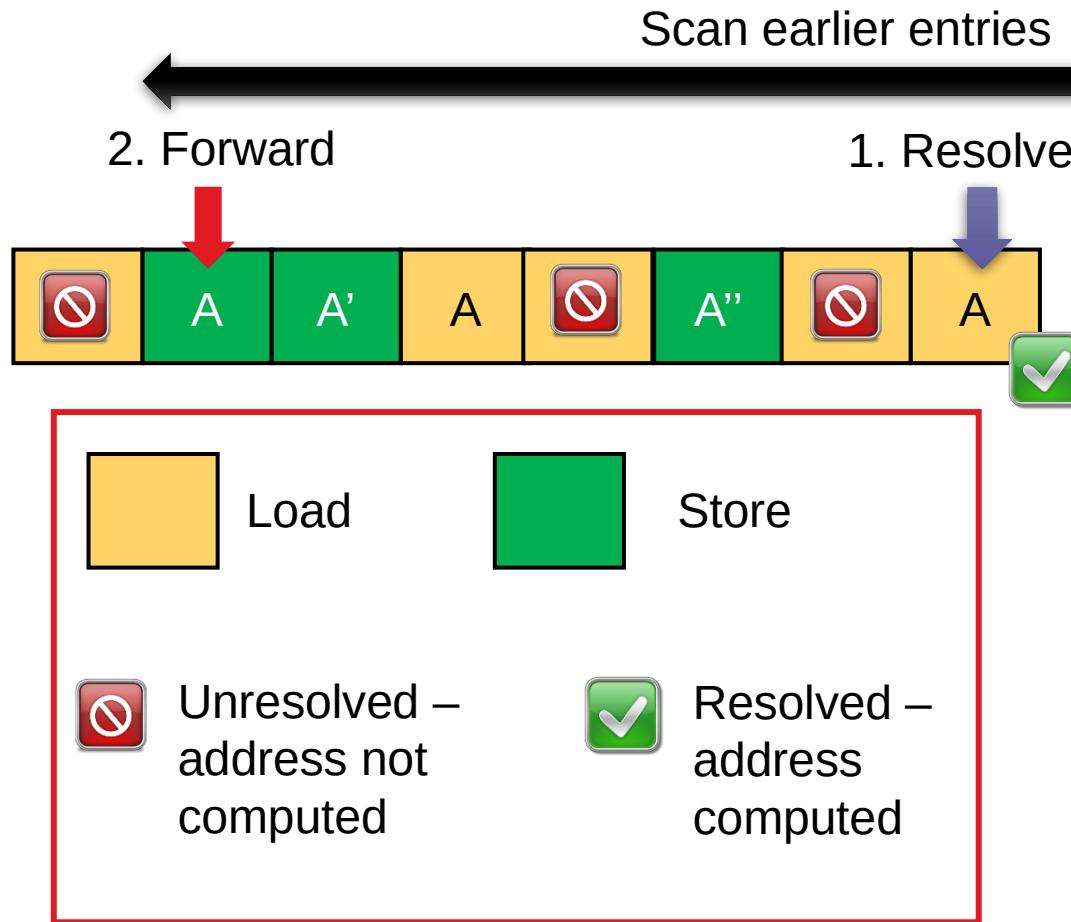
# Compute Address of a Store – Case II



# Compute Address of a Load – Case I



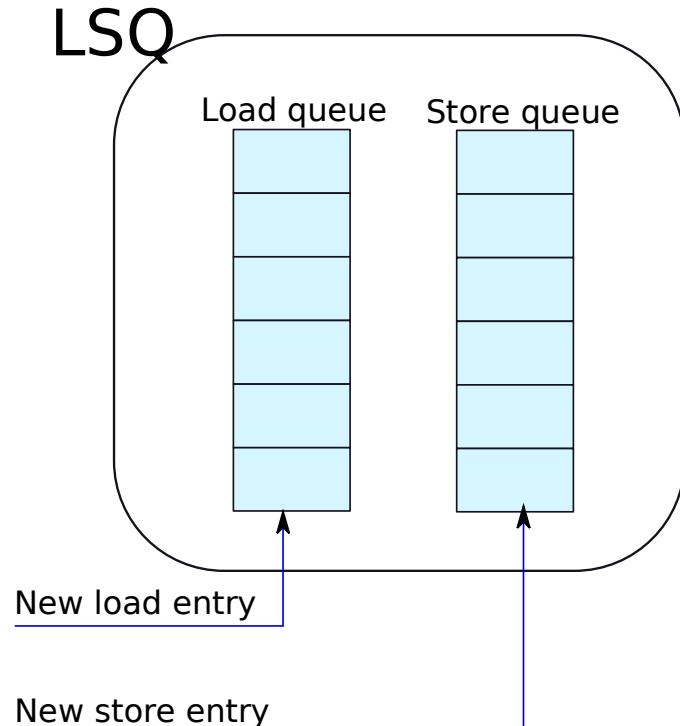
# Compute Address of a Load – Case II



# Summary of the Discussion

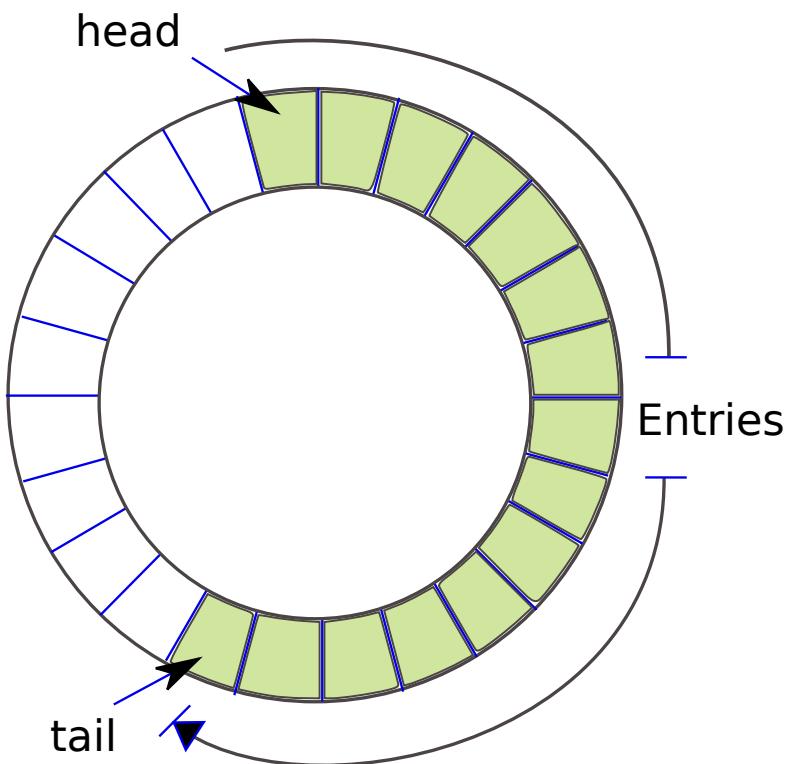
Type	Search direction	Action: Condition
Loads	Search all the stores before it	<ol style="list-style-type: none"><li>1. Terminate and forward value: Store to the same address</li><li>2. Terminate: Store with an unresolved address</li><li>3. Go to d-cache: Conditions (1) and (2) not met</li></ol>
Stores	Search all the loads and stores after it	<ol style="list-style-type: none"><li>1. Terminate: Store to the same address</li><li>2. Terminate: Store with an unresolved address</li><li>3. Forward value: Load from the same address</li></ol>

# Actual Implementation

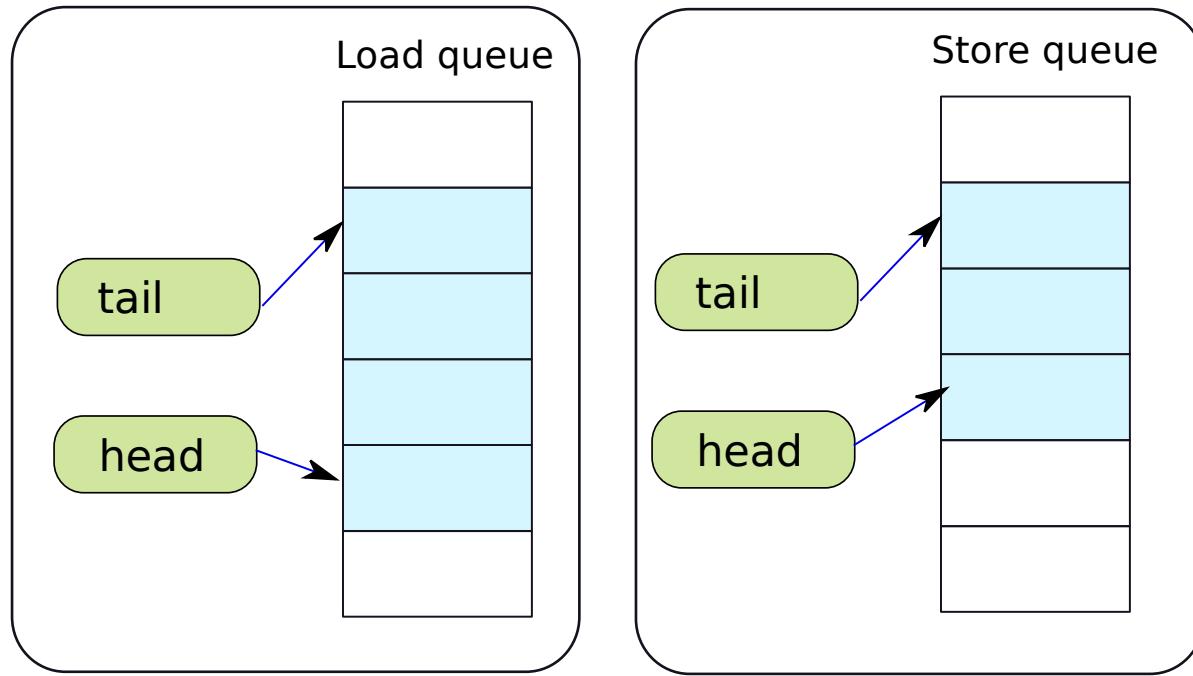


- For the purpose of **efficiency**, have two queues:
  - One load queue
  - One store queue

# Both are arranged as a circular queue



# Organization



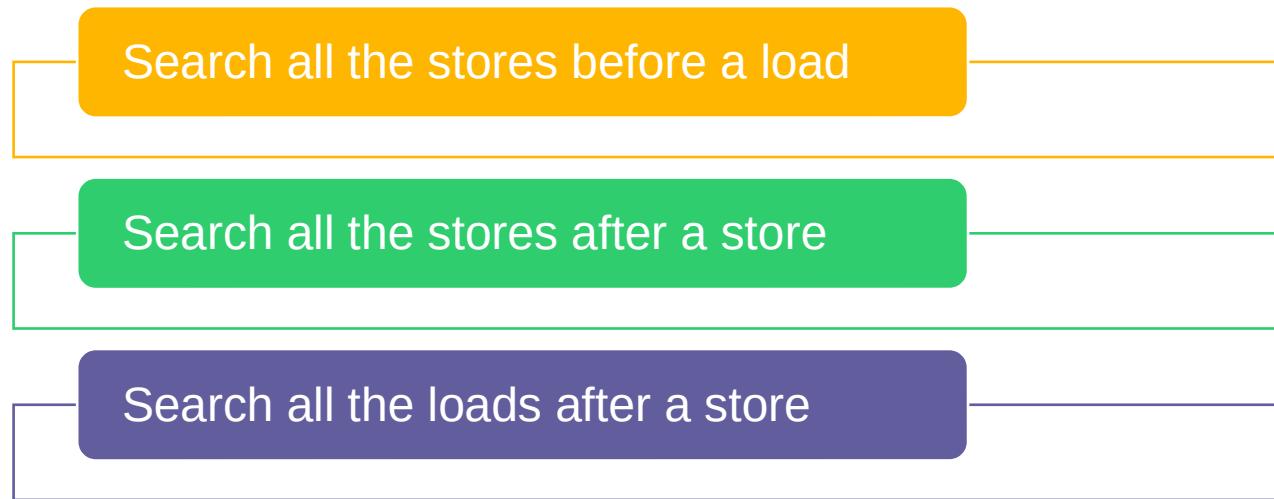
## Load queue entry

1. Load address
2. Index of the *tail* ptr in the store queue when the entry was added

## Store queue entry

1. Memory address and value
2. Index of the *tail* ptr in the load queue when the entry was added

# Basic Search Operations



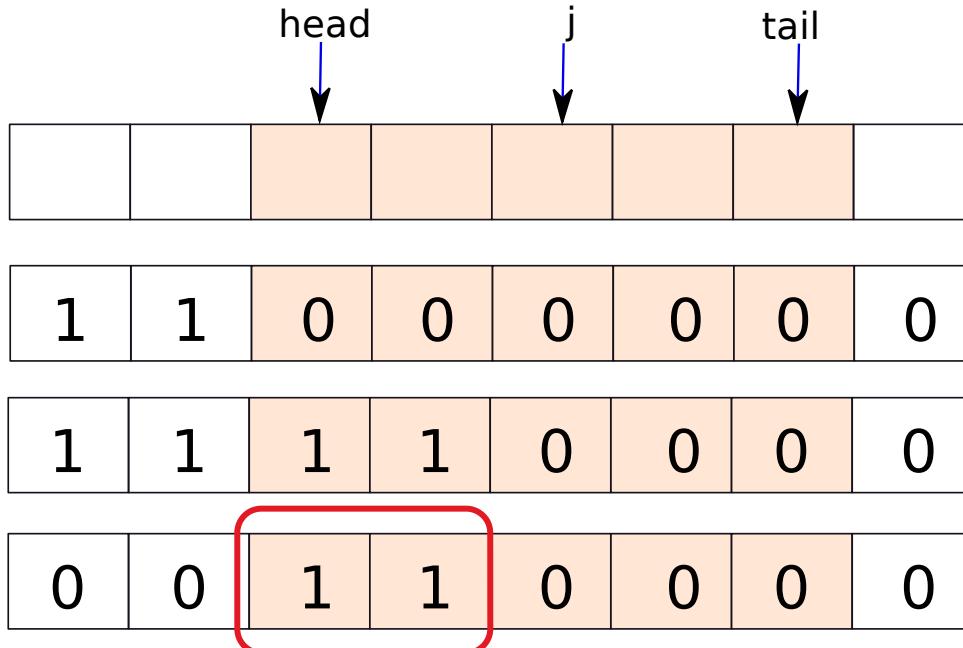
- Represent locations as a bit vector. If a queue has  $N$  entries, we have  $N$  bits: one bit per each entry.
- Create a bit vector called *prec* – all the locations before a given location after set to 1.

# Implementation of functions: *before*

$j \geq head$

$before(j) = \overline{prec(head)} \wedge prec(j)$  be

Compute quickly in HW



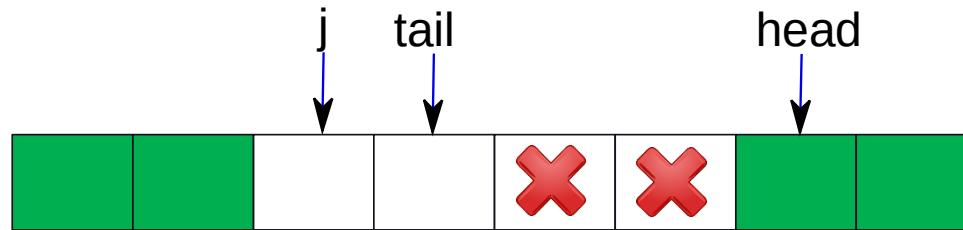
Entries before *j*  
in the queue

# Expressions for other cases

$j < head$

Wraparound

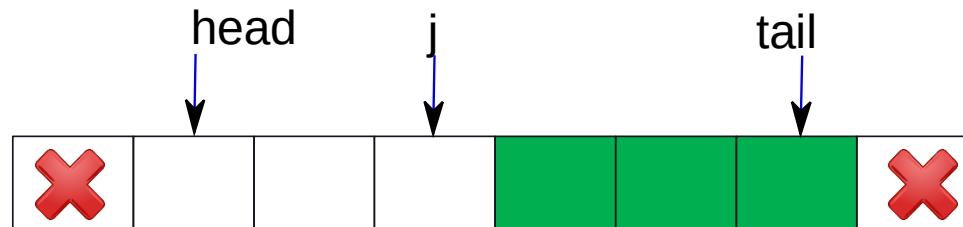
$$before(j) = \overline{prec(head)} \vee prec(j)$$



*after* function

$j \leq tail$

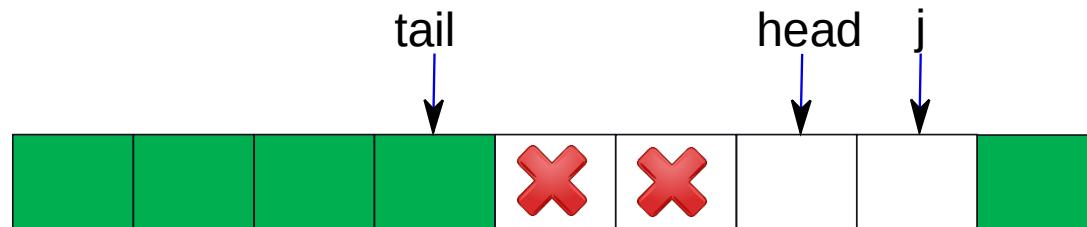
$$after(j) = \overline{prec(j)} \wedge \overline{map(j)} \wedge (prec(tail) \vee map(tail))$$



$j > tail$

Wraparound

$$after(j) = (\overline{prec(j)} \vee prec(tail) \vee map(tail)) \wedge \overline{map(j)}$$



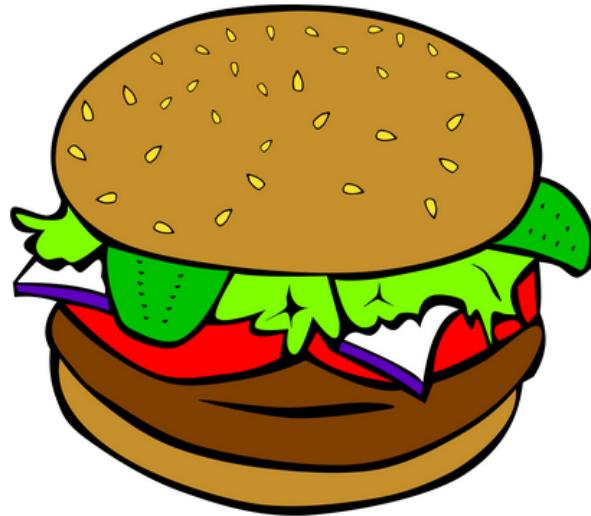
# Operation of the LSQ

Type	Search direction	Action: Condition
Loads	Search all the stores before it	<ol style="list-style-type: none"><li>1. Terminate and forward value: Store to the same address</li><li>2. Terminate: Store with an unresolved address</li><li>3. Go to d-cache: Conditions (1) and (2) not met</li></ol>
Stores	Search all the loads and stores after it	<ol style="list-style-type: none"><li>1. Terminate: Store to the same address</li><li>2. Terminate: Store with an unresolved address</li><li>3. Forward value: Load from the same address</li></ol>

- We need to maintain a few more bit-vectors
  - *resvd*: bit-vector to store whether an entry is resolved or unresolved
  - *match*: all the entries that match a given address (we can use a CAM array for this purpose, Chapter 7)
- *Example*: All matching/unresolved stores before a load:

# Wrap-up

- We can compute various functions and find a set of locations of interest.
- We need to find either the first or last location.
- How do we do this quickly in hardware?



Can we use a select unit?

# Contents

- 
- 1. Instruction Renaming
  - 2. Instruction Dispatch, Wakeup, Select
  - 3. The Load-Store Queue (LSQ)
  - 4. Instruction Commit

# Precise Exceptions



For an external observer

- Instructions need to appear to **execute** in order
- At any point of time, the instruction flow can be **paused**, and **resumed** later.

# New Structure: Reorder Buffer (ROB)

## Reorder Buffer (entry)

- Contains 1 entry for every instruction that has been fetched
- After decoding an instruction, we enter it in the ROB
- If there are no free entries, the OOO pipeline stalls
- Instructions are entered into the ROB in program order
- The ROB is essentially an in-order queue

## Updating an entry

- Whenever an instruction finishes its execution, we update the entry of the ROB
- Mark it ready

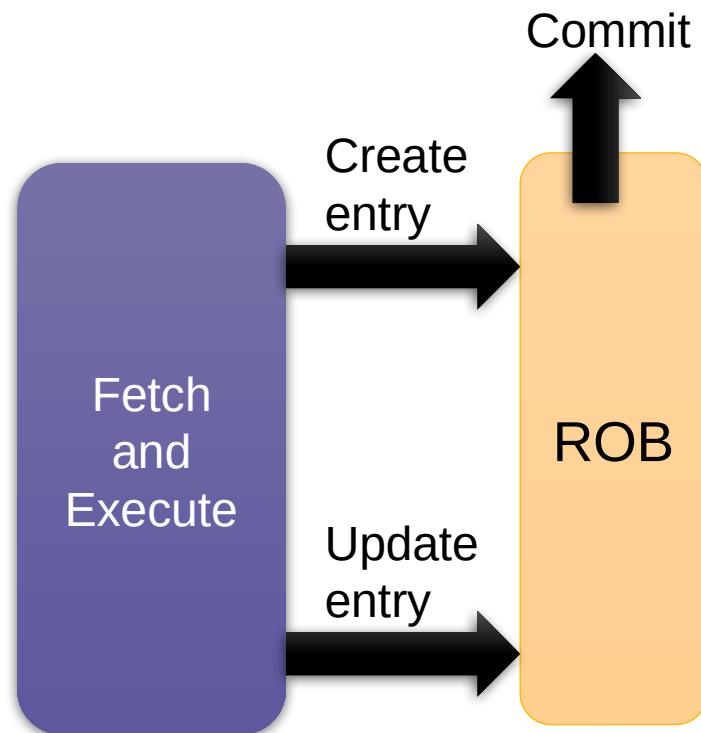
## Removing an entry

- It is called retiring or committing
- This happens in program order (keep reading ... )

# Removing Entries from the ROB (Commit/Retire)

Every processor has a **commit width**

- Number of instructions to **commit** per cycle
- Let the commit width be **W**
- 2, 4, or 6 (typical **values**)
- Start from the top of the **ROB** and access the next W entries
- Stop at the first entry, which is not ready or continue till we reach the W<sup>th</sup> entry
- Commit all the **entries** that we find to be **ready**



# What does it mean to commit an entry?

Remove it from the ROB

AND



Depends on the type of instruction

Consider instructions with a register destination

- Inst J:  $r1 \leftarrow r2 + r3$
- Assume that  $r1$  is mapped to  $py$
- Before this instruction,  $r1$  was mapped to  $px$

When can we free  $px$ ?

- After J **commits**, there are no instructions earlier than J in the **pipeline**
- There is no **instruction** that requires the value in  $px$
- $px$  can be reclaimed, and added to the **free list**

# Updates to renaming

Whenever we see an **instruction** of the form:

- Instruction J:  $r1 \equiv \dots$
- Remember the previous mapping of  $r1$  by reading the **rename table**
- Assume that  $r1$  was mapped to  $px$
- Save this mapping in the **ROB** entry for instruction J
- When J is committed, **unmap**  $px$ , and return it to the **free list**

This can be done for all **instructions** that have a **destination** operand as a register: ALU instructions, load instruction

# Load and Store Instructions

When should a store be written to **memory**?

- **ANS:** When we are sure that it is not in the **wrong path** of a branch, or no **interrupt** will come before it
  - Only possible at **commit time**.
  - Solution: Send stores to memory at **commit time**.
  - At the same time remove them from the store queue
- 

When should a load be removed from the LSQ?

- At commit time ^
-

# What do we in the case of an interrupt?

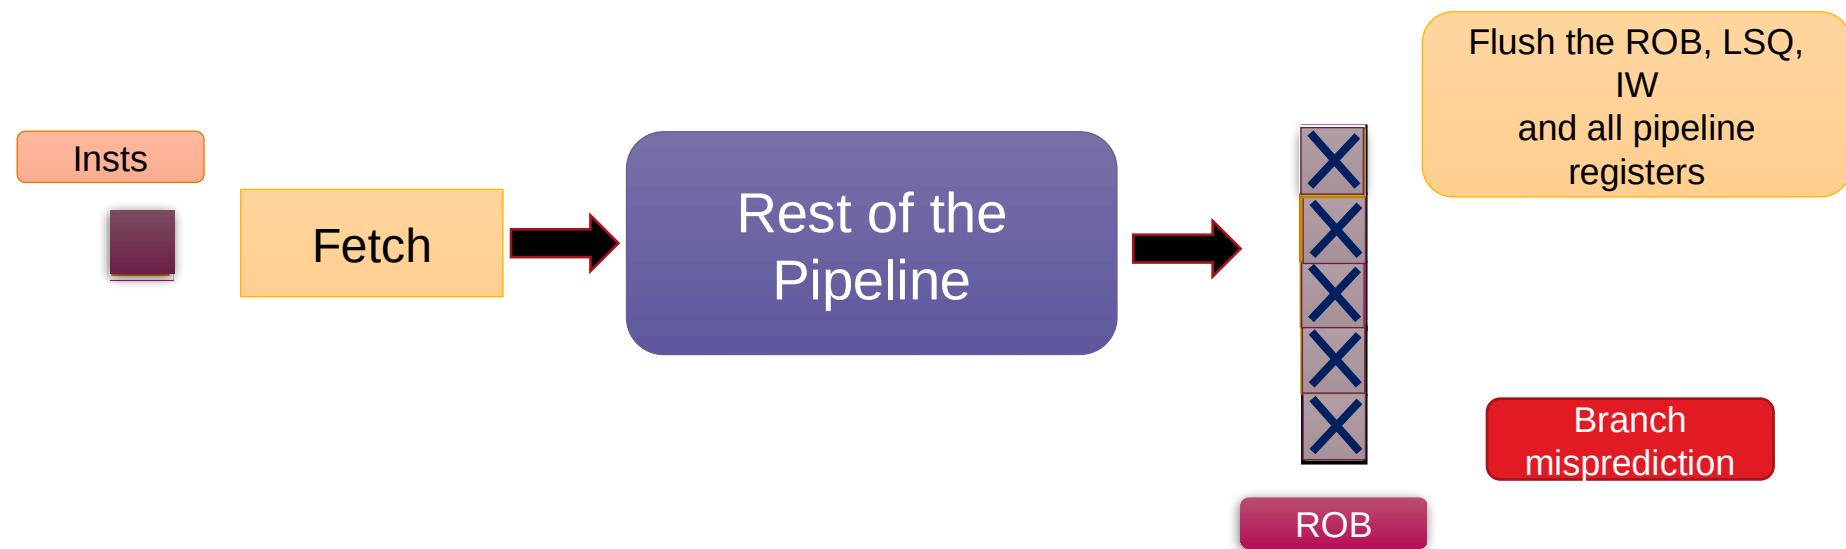
Let us treat interrupts, exceptions, and branch mispredictions in the same category

We will clearly have some instructions in the pipeline that are on the wrong path. They should not be committed

Use the ROB:

- Mark the instruction that has had an exception, or suffered from a misprediction, in the ROB
- In the case of an interrupt mark the topmost entry in the ROB
- Wait till the marked instruction retires
- Initiate the recovery sequence
- More on the next slide 

# Flushing the Pipeline



# Recovery Sequence

- Flush the pipeline
- Remove the entries from the IW, LSQ, and ROB
- Remove entries from all pipeline registers and other temporary structures (not the rename table or register file though)

## Restore the state

- Architectural register state for the last committed instruction and the PC
- The contents in the rest of the physical registers does not matter

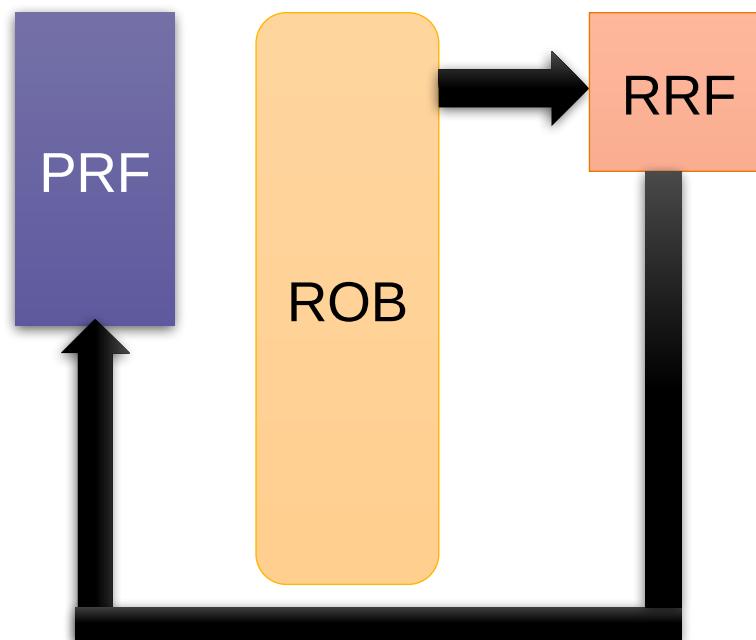
# Approach 1: Retirement Register File (RRF)

RRF  $\sqsubseteq$  Maintain the state of all architectural **registers** at commit time

Each entry of the ROB gets augmented with the value of the **destination register**

Write the value of the **destination register** at commit time.

The **RRF** maintains a checkpoint at commit time



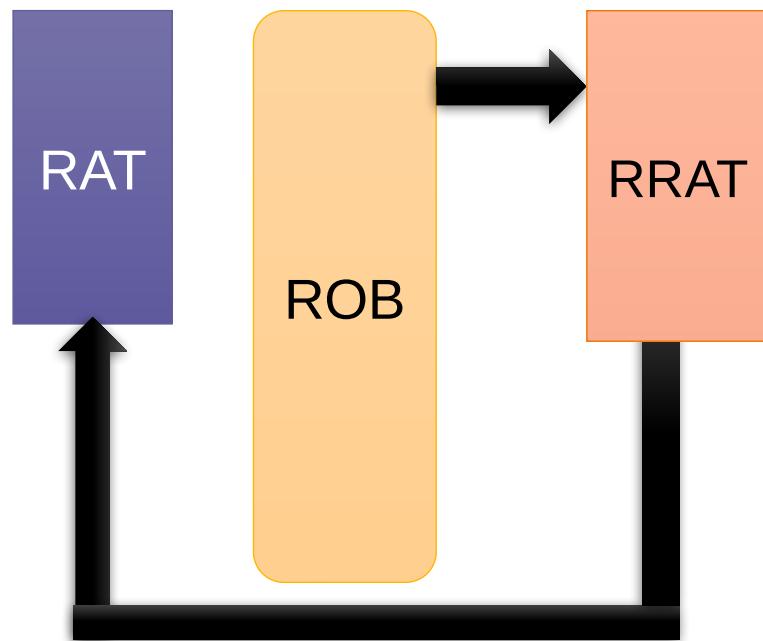
## Approach 2: Retirement RAT (RRAT)

Why do we need to write the **values** into the RRF?

Aren't just the **mappings** sufficient.

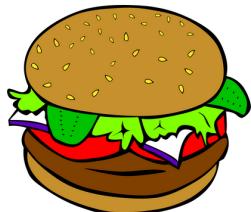
Maintain a copy of the mapping of the destination register in the **ROB** entry

- Eg:  $r1 = r2 + r3$ ,  $r1$  gets mapped to  $py$
- Maintain  $r1 = py$  in the **ROB** entry
- Update the RRAT with this mapping



At the time of **recovery**

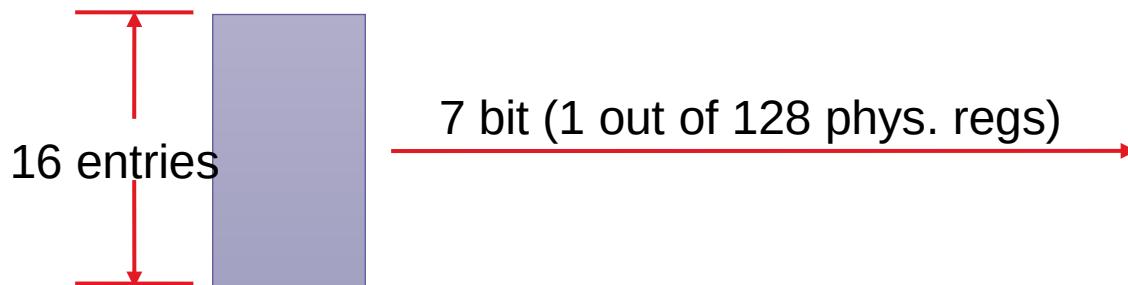
- Transfer the **mappings** in the RRAT to the RAT



Prove that this can be done.

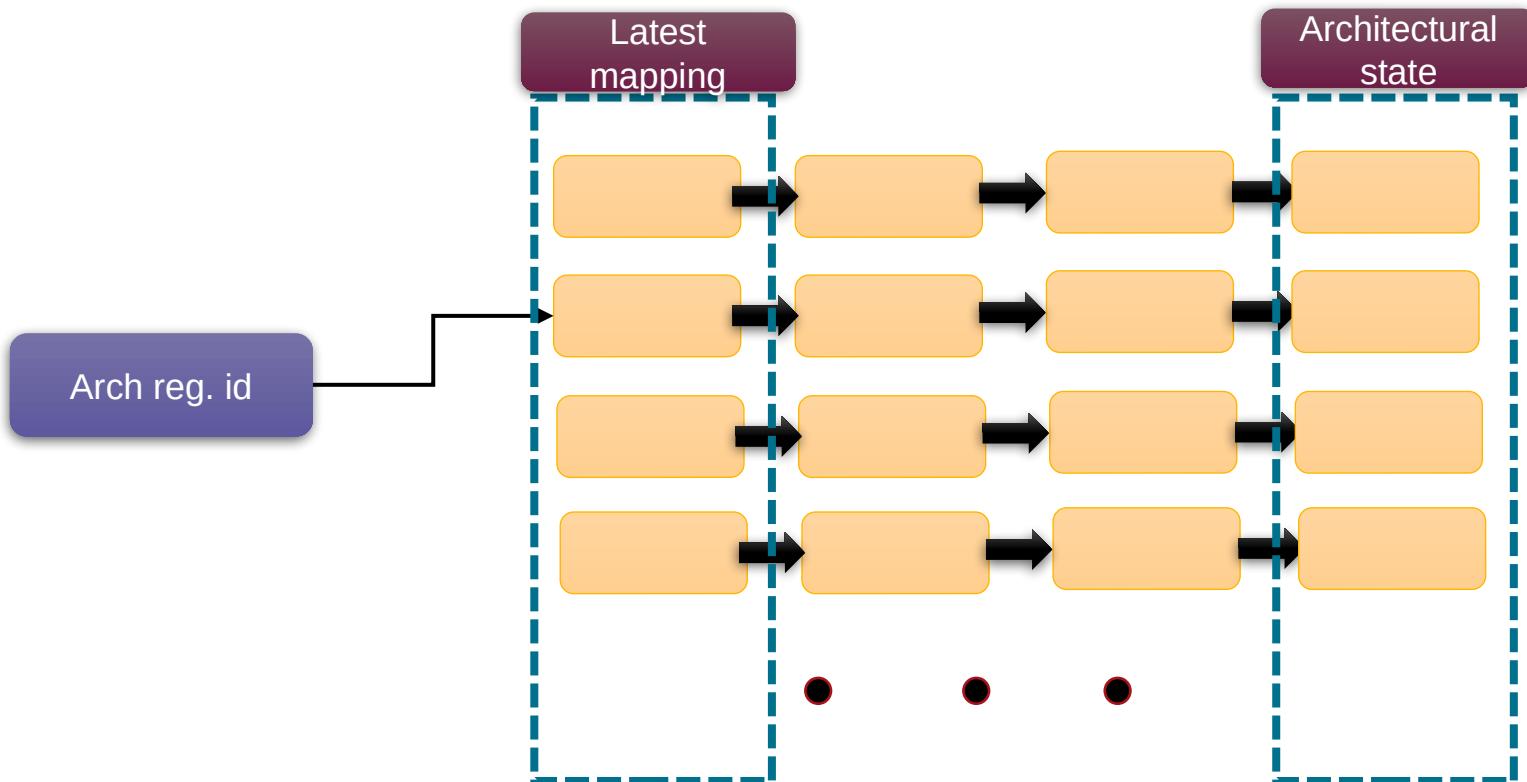
## Approach 3: RAT Checkpointing (SRAM based)

Consider the following design



- At every branch let us maintain a **checkpoint**
- Set of 16, 7-bit ids
- In each **entry** of the **RAT**, have a **shift register**, where each entry of the shift register is 7 bits wide
- To create a **checkpoint**, shift the entries in each row.
- The last **entries** of each **shift register** represent the current **architectural state**
- Remove the last entries once the corresponding instructions **commit**

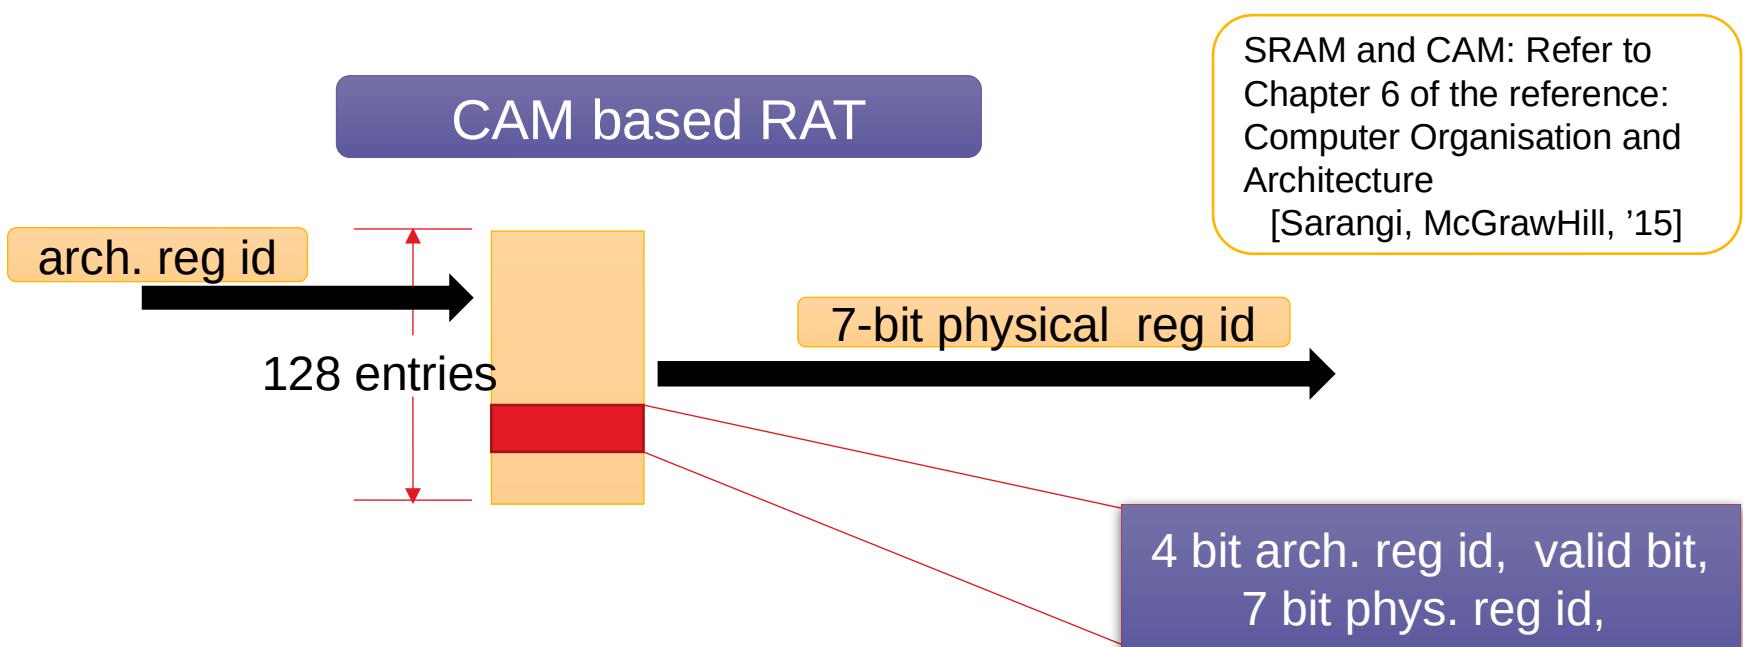
# RAT



# Approach 4: RAT Checkpointing (CAM based)

A **CAM** (content addressable memory) is like a **hash table**

- You address each element not by an index, but by a subset of its **content**
- It is **slower** than a structure that is addressed by an **index**

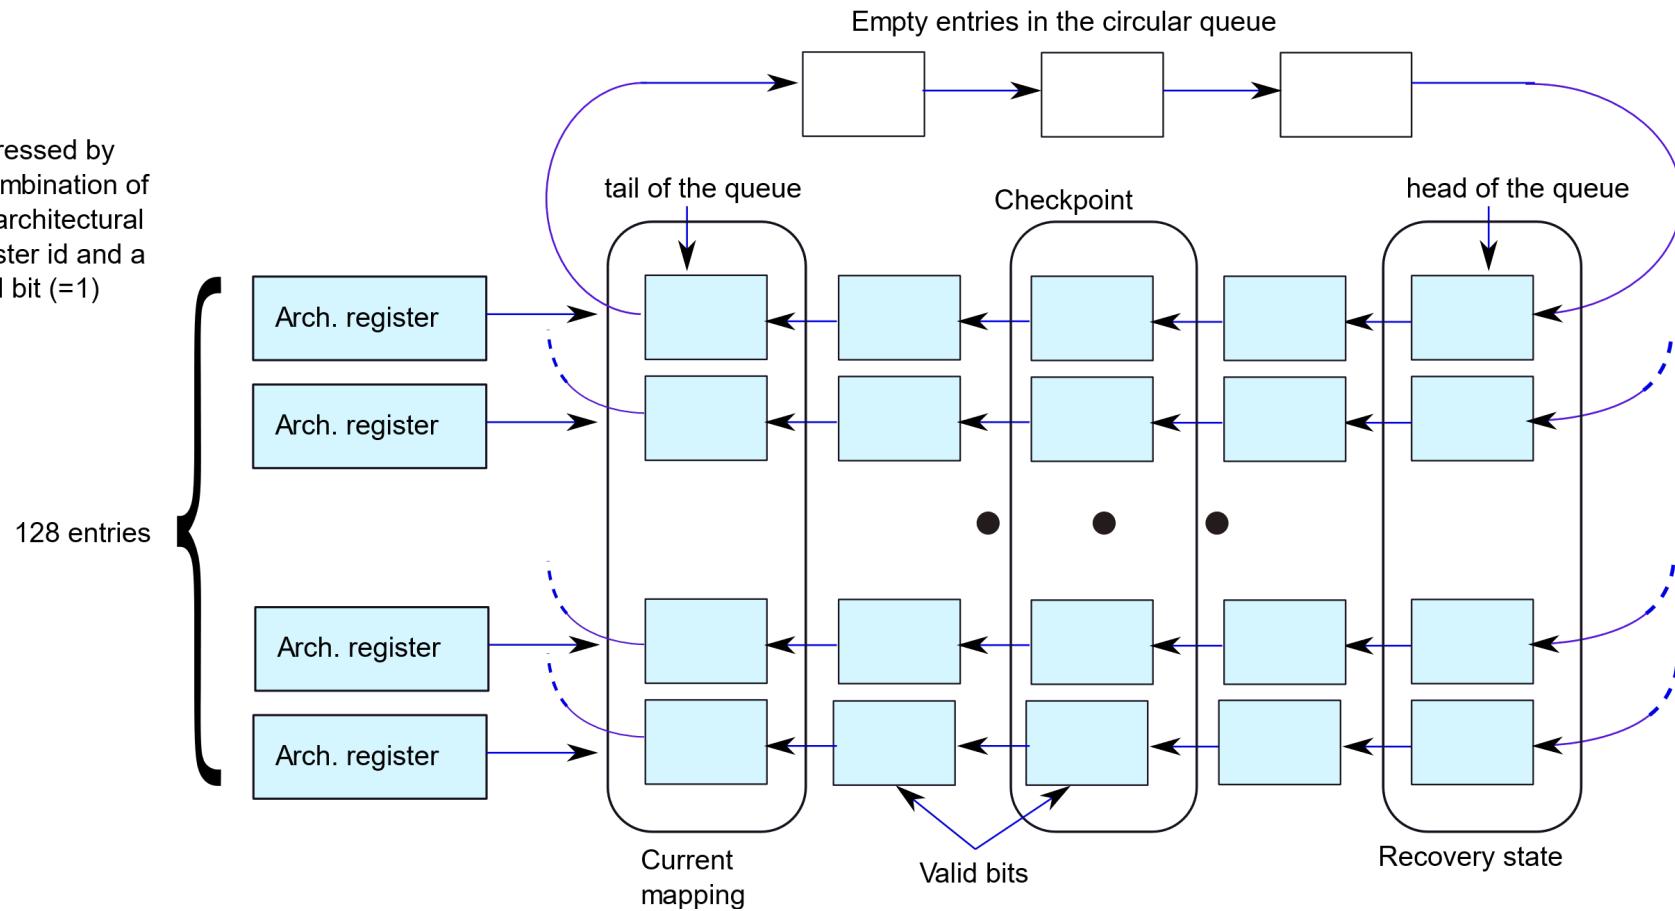


# CAM based RAT

- At any point of time, 16 bits are set to 1. They contain the current mapping (latest state)
- The **architectural state** is also a set of 16 entries in the **RAT** table.
- We can save a **checkpoint** by saving the set of valid bits in a 128 bit **vector**.
- For every **branch**, we can take a **snapshot** of the **valid** bits.
- The best way to achieve this is to have a shift register in each row of the RAT table, where each entry is 1 bit wide
- Use the same **idea** as the SRAM based RAT

# CAM based checkpointing, one 1 bit per row

Addressed by  
a combination of  
the architectural  
register id and a  
valid bit (=1)



# Pros and Cons

## RRF

- + point □ Simple to implement. Transferring the checkpoint is easy
- point □ Extra register writes every cycle. More power
- point □ Need to save the value in the ROB. More space.

## RRAT

- + point □ Requires less space in the ROB than the RRF
- point □ Has activity every cycle.

## SRAM based

- + point □ Activity only on a branch. Shift operation.
- point □ Each row of the RAT is wider. Cannot resume from non-branch points.

## CAM

- + point □ Shift operations are easy. Shift only 1 bit.
- point □ The CAM per se is a slower structure than an SRAM.

# Conclusion

Renaming is done using the RAT table that stores a mapping between an arch. register and a physical register.

We additionally need a free list and dependence check logic.

The scheduler comprises the wakeup, select, and broadcast logic. We often perform an early broadcast to avoid stalls.

The LSQ manages the dependences between loads and stores. It forwards data between in-flight loads and stores.

The ROB is used to keep track of in-flight instructions. Instructions are committed when they complete.



The End

The word "The" is positioned at the top in a dark blue serif font. The word "End" is positioned below it, partially overlapping, in a larger dark blue serif font. The background features a large red triangle pointing towards the bottom-left, a smaller maroon triangle at the base, and a light beige triangle pointing towards the top-right. The text is overlaid on these geometric shapes.